

The Metric Blind Spot in Deep Learning-Based Brain MRI Reconstruction: A Systematic Review

Dat Tat Mai^{a,*}, Thai Viet Pham^b, Thu Nguyen Thi Dang^c, James Jin Kang^a

^a*School of Science, Engineering & Technology, RMIT University Vietnam, Ho Chi Minh City, Vietnam*

^b*School of Computing Technologies, RMIT University, Melbourne, VIC, 3000, Australia*

^c*School of Health and Biomedical Science, RMIT University, Melbourne, VIC, 3000, Australia*

Abstract

Objective: Deep learning accelerates brain MRI four- to tenfold, but learned models can erase a lesion or invent tissue, and pixel-averaged scores such as PSNR and SSIM miss both. We review whether evaluation practice could detect this blind spot. **Methods:** Following PRISMA 2020 we searched seven databases without date limit, included 263 studies (1995–2026), appraised them with QUADAS-2 and matched instruments, and synthesised narratively. Categories came from titles, abstracts and controlled vocabulary; every prevalence figure is a floor. Duplicate screening achieved substantial agreement (Fleiss’ $\kappa = 0.877$); appraisal agreement was substantial (0.788; 0.390 where observable); extraction is unaudited. **Results:** Only 18 of 263 studies (6.8%) record a fidelity metric and a reader assessment on the same data, leaving the field’s central proxy unmeasured. Reader studies mostly measure reader-versus-reader agreement, which is weak: fastMRI 2020 concordance was 0.457 and 0.386 (Kendall’s W), improving only where SSIM had diverged. Erasing a 100 mm³ lacunar infarct moves global PSNR about 0.03 dB under the stated error model. The corpus grew fivefold; reader assessment fell from 32% to 18%, recovering to 21%. Among the four largest families the most hallucination-associated (generative, 39%) is among the least reader-evaluated (11.3%); self-supervised is higher still (47%) with no readers. One study in twenty both releases code and runs readers; none evaluated a model observer; no named dataset covers acute stroke or

*Corresponding author.

Email addresses: `dat.mai2@rmit.edu.vn` (Dat Tat Mai), `s4229249@student.rmit.edu.au` (Thai Viet Pham), `s4205224@student.rmit.edu.au` (Thu Nguyen Thi Dang), `james.kang@rmit.edu.vn` (James Jin Kang)

25 haemorrhage. **Conclusions:** On these floors, evaluation practice cannot certify
26 diagnostic safety. We derive five requirements safety-oriented evaluation must
27 meet.

28 *Keywords:* Deep learning, Brain MRI reconstruction, Reconstruction safety,
29 Hallucination, Image quality assessment, Systematic review

30 1. Introduction

31 Brain MRI is widely used because it can clearly distinguish different types
32 of brain tissue and reveal subtle abnormalities that are difficult to detect with
33 other imaging modalities. However, this image quality comes at the cost of long
34 acquisition times. Unlike conventional imaging methods that capture an image
35 more directly, MRI builds an image from a sequence of raw signal measure-
36 ments (samples of spatial frequencies, known as k -space) collected over time.
37 This prolonged acquisition reduces patient throughput and increases sensitivity
38 to patient motion, which can introduce artifacts and degrade image quality (Za-
39 itsev et al., 2015). To address these limitations, acceleration techniques such as
40 parallel imaging, including Sensitivity Encoding (SENSE) (Pruessmann et al.,
41 1999), and compressed sensing (Lustig et al., 2007) were developed to reduce
42 the amount of data that must be collected. More recently, deep learning has
43 enabled further acceleration by reconstructing high-quality images from incom-
44 plete MRI measurements, through cascaded convolutional networks (Schlemper
45 et al., 2017), model-based unrolled networks (Hammernik et al., 2018; Aggar-
46 wal et al., 2018), generative adversarial networks (Yang et al., 2018), and score-
47 based diffusion models (Song et al., 2021), at accelerations from fourfold to, in
48 exploratory work, an order of magnitude (Radmanesh et al., 2022).

49 The reconstructed images look convincing, and that is the problem. Small
50 changes in the input can produce large errors in the output (Antun et al., 2020;
51 Gottschling et al., 2025). Four failure types recur across the studies we re-
52 viewed: strong priors smooth away faint structure; motion causes ghosting the
53 model never saw in training (Sommer et al., 2020); a change of scanner or sam-

54 pling degrades results in ways that in-distribution testing misses (Güngör et al.,
 55 2023); and generative models can invent plausible tissue (Zhou et al., 2024).
 56 The scores used to certify these methods, peak signal-to-noise ratio (PSNR)
 57 and structural similarity (SSIM) (Wang et al., 2004), agree only partially with
 58 radiologist quality judgement even on conventional image degradations (Mason
 59 et al., 2020); and SNR, though reproducible across observers, tracks no-reference
 60 quality models only moderately (Yu et al., 2018), so an image can score well
 61 while the error that matters goes unrecorded. We call this the **metric blind**
 62 **spot**: standard image-quality metrics cannot catch these small but clinically se-
 63 rious errors, because a lesion of a few cubic millimetres barely moves an average
 64 taken over millions of voxels.

65 Parts of the field have responded, training quality metrics on radiologist
 66 rankings (Tang et al., 2025), attaching uncertainty maps to reconstructions
 67 (Khawaled and Freiman, 2024), or running reader studies that test an acceler-
 68 ated protocol against a reference one directly (Lang et al., 2024; Radmanesh
 69 et al., 2022). Causal and counterfactual methods can ask the most direct ques-
 70 tion, whether a lesion survives a controlled change of acquisition, but have
 71 mostly been applied to classification rather than reconstruction (Pawlowski
 72 et al., 2020; Castro et al., 2020). These responses share a root cause. Ac-
 73 celeration works because a learned prior fills in information the measurement
 74 does not contain, and that same filled-in information can hide a real finding or
 75 create a false one, while the field certifies the result with averaged pixel scores
 76 that cannot see either event.

77 To our knowledge no systematic review has assembled the evidence on this
 78 problem. Existing surveys cover architectures and benchmark performance
 79 (Liang et al., 2020; Montalt-Tordera et al., 2021); none asks how often con-
 80 ventional metrics miss clinically important errors, why they miss them, or what
 81 evaluation would have to look like to catch them. We address that through
 82 three questions:

83 **RQ1** *How reliably do conventional fidelity metrics track radiologist diagnostic*

84 *ratings in deep learning-based accelerated brain MRI, and how wide is the*
85 *evaluation reporting gap in the literature?*

86 **RQ2** *What underlying mathematical mechanisms cause standard metrics to*
87 *overlook localized pathology erasure and generative hallucinations in deep*
88 *learning-based reconstruction?*

89 **RQ3** *What can each evaluation paradigm in current use establish about prior-*
90 *induced diagnostic error in deep learning-based MRI reconstruction, and*
91 *what requirements for safety-oriented evaluation follow from those limits?*

92 We contribute: (C1) a synthesis of what 263 primary studies report about
93 agreement between pixel-wise metrics and radiologist judgement; (C2) four
94 mathematical mechanisms that explain the divergence, with an exact identity
95 bounding the PSNR cost of erasing a lesion; (C3) an analysis of *where* the
96 evaluation gap comes from, showing that it co-occurs with a division of evalu-
97 ation practice between research communities and with a benchmark infrastruc-
98 ture lacking acute pathology; and (C4) five requirements that safety-oriented
99 evaluation must meet. The auditing framework these requirements motivate is
100 developed and evaluated in a companion methods paper, not here.

101 Three boundaries should be stated at the outset, because they determine
102 what the findings can and cannot be asked to support. First, this is a review of
103 *evaluation practice*, not of reconstruction performance: we count what studies
104 measured, not how well their methods performed, and nothing here estimates
105 how often a deployed reconstruction actually erases a lesion. Second, the mech-
106 anisms in RQ2 are mathematical properties of the metrics, derived under a
107 stated error model; they establish that a focal error *can* pass unrecorded, not
108 that it commonly does. Third, prevalence counts rest on labels extracted from
109 what studies report, so every one of them is a lower bound on practice rather
110 than an estimate of it, and we treat the direction of that bias explicitly in Sec-
111 tion 5.2. The contribution is a diagnosis of what the field’s current instruments
112 are capable of detecting, which is logically prior to, and considerably cheaper
113 to establish than, a measurement of how often they fail.

114 2. Related Work

115 Four strands bear on this review, and the claim that none of them has
116 assembled the evidence we assemble here is a claim about each in turn.

117 *Surveys of reconstruction methods.* Reviews of deep learning MRI recon-
118 struction catalogue architectures, acceleration factors and benchmark accuracy
119 (Liang et al., 2020; Montalt-Tordera et al., 2021; Knoll et al., 2020a; Kazer-
120 ouni et al., 2023). They are organised by method, and their outcome variable
121 is reconstruction fidelity against a fully sampled reference. Safety and failure
122 modes appear as caveats in discussion sections rather than as an object of syn-
123 thesis, and none of them tabulates what evaluation each primary study actually
124 performed, which is the quantity this review extracts.

125 *Task-based image quality assessment.* The tradition that ought to have sup-
126 plied the answer established, decades before deep learning, that image quality is
127 meaningful only relative to a diagnostic task. Barrett and colleagues formalised
128 the observer as part of the measurement: an image is good if an observer per-
129 forming a specified task on it achieves a high figure of merit, typically the
130 area under an ROC curve for a detection task, and the mathematical (ideal or
131 channelised Hotelling) observer supplies a computable surrogate for the human
132 when the task and image statistics are specified (Barrett et al., 1993; Barrett
133 and Myers, 2004). ICRU Report 54 made the same argument the international
134 reference framework for assessing medical imaging quality (International Com-
135 mission on Radiation Units and Measurements, 1996). That framework already
136 contains the concept the reconstruction literature lacks: a detectability index
137 for a specified lesion, which is by construction sensitive to exactly the focal
138 change a whole-volume L_2 score averages away. It has been carried into the
139 dose- and protocol-optimisation literatures of other imaging modalities, but it
140 has not crossed into deep learning MRI reconstruction: no study in this cor-
141 pus evaluated a model observer, and our first requirement (R1, Section 4.7) is
142 in substance a restatement of the task-based position for a field that has not
143 adopted it. The reason it has not is worth naming: model observers require a

144 specified signal and a statistical description of the background, and the null-
145 space structure that makes learned reconstruction dangerous is precisely what
146 makes those two ingredients hard to specify.

147 *Failure-mode-specific work.* Artifact detection, motion correction and ro-
148 bustness under distribution shift are each addressed in depth, but in isolation,
149 by literatures that do not share an evaluation vocabulary (Esteban et al., 2017;
150 Sommer et al., 2020; Güngör et al., 2023). Analyses of hallucination in tomo-
151 graphic reconstruction establish that prior-induced structure lives in the null
152 space of the forward operator, so it is invisible to any measurement-space con-
153 sistency check (Bhadra et al., 2021; Gottschling et al., 2025), and causal and
154 counterfactual methods for medical imaging supply the machinery to intervene
155 on image content and ask whether a finding survives (Pawlowski et al., 2020;
156 Castro et al., 2020; Sanchez et al., 2022; Peng et al., 2026). Neither line has
157 been connected to review-level evidence on how reconstruction is evaluated in
158 practice.

159 *What is new in the four mechanisms of RQ2, and what is not.* Because
160 Section 4 states four mechanisms formally, their provenance should be set out
161 plainly rather than left for a reader to reconstruct. Two are elementary prop-
162 erties of the objective rather than findings: spatial error averaging follows from
163 the definition of a mean over the image volume, and loss symmetry from the
164 fact that squaring discards the sign of the residual. Two are established results
165 we restate: reference dependence is standard in the image-quality literature,
166 and the null-space account underlying the realism–reward trade-off is due to
167 Bhadra et al. (2021) and Gottschling et al. (2025). What this review adds is
168 not the mechanisms but their quantification and their assembly — the identity
169 at Equation (2), which converts any lesion volume into the PSNR movement its
170 erasure would cause, and the setting of all four against what a corpus of 263
171 studies is recorded as measuring. We claim the assembly and the bound, not
172 the underlying results.

173 *Benchmark and annotation infrastructure.* The raw k -space resource the
174 field benchmarks on is fastMRI (Knoll et al., 2020b), whose brain release fol-

lowed the knee data described in that report, and an annotation layer is emerging on top of it: fastMRI+ adds expert pathology labels to the fastMRI brain and knee data (Zhao et al., 2022), and curated stroke-lesion datasets exist for segmentation (ISLES (Hernandez Petzsche et al., 2022), ATLAS (Liew et al., 2022)). None of these supplies raw k -space paired with acute vascular pathology, which is the combination reconstruction evaluation needs, and Section 4.5 shows that this is not an incidental gap but the binding constraint on the whole evaluation problem.

What no strand supplies is the join: a systematic account of how often the field measures metric–reader agreement, why the metrics fail when it does not, and what each evaluation paradigm can therefore establish. The four strands are not merely separate; they are held apart by what each can access. The survey and hallucination-analysis literatures work on datasets with raw k -space and no clinically verified pathology; the task-based tradition needs a specified signal and a labelled task, which those datasets do not carry; and the clinical reader literature has the pathology and the readers but no access to the raw data on which the mechanisms operate. Section 4.5 shows this is not an abstraction but a measurable division in the corpus itself. Supplementary Table S4 sets out each secondary study we identified against the questions asked here.

3. Methodology

This systematic review was conducted in accordance with the PRISMA 2020 statement (Page et al., 2021); the completed checklist, search vocabulary, extraction schema, and analysis scripts are available in the supplementary material and on the Open Science Framework (<https://doi.org/10.17605/OSF.IO/EUXA8>).

Protocol and registration status. The protocol, covering the research questions, eligibility criteria, search strategy, screening procedure, classification rubric and synthesis approach, was fixed in the project repository before screening began. In the terms of PRISMA 2020 item 24a, **this review was not prospectively registered and carries no registration number:**

PROSPERO’s scope excludes reviews without direct clinical health outcomes, and no equivalent prospective registry covers methodological reviews of this kind. The Open Science Framework deposit cited above therefore timestamps the record but is retrospective, and we make no claim of prospective registration on it. Two consequences follow. An external record does not exclude undisclosed post-hoc analytic choices; and the two analyses we ourselves identify as post-hoc — the acceleration-factor extraction of Table 6 and the venue classification of Section 4.5 — are labelled as such wherever they appear.

Search strategy. Seven bibliographic databases (PubMed/MEDLINE, IEEE Xplore, Scopus, Web of Science, SpringerLink, OpenAlex, and Semantic Scholar) were searched between 23 and 24 July 2026, with no publication-date limit applied; the 1995–2026 span quoted elsewhere in this article is the publication range of the studies that were included, not a restriction imposed on the search. Queries were executed through each database’s API rather than its web interface, so the complete Boolean match set was retrieved for every term combination rather than a ranked sample, explaining why the initial identification count is large. The strategy combined three required concept groups: neuroimaging (Block B1), reconstruction and acceleration (Block B2), and safety, failure modes, and evaluation outcomes (Block B3), with MeSH headings and wildcard truncation adapted per database (Supplementary Table S5). Table 1 details the executed query structure and hit breakdown across databases.

The breadth of this search is by design. Evidence of reconstruction failure is rarely a paper’s stated subject: artifacts, hallucinations and metric–reader disagreements are most often reported in passing inside studies whose declared contribution is a new architecture, so a query tuned narrowly to safety vocabulary would miss much of this evidence. Clinical and engineering literatures also use distinct vocabularies — sequences and artifacts against inverse problems, unrolled optimization and generative priors — and neither database family indexes the other completely. We therefore tuned for sensitivity rather than precision and accepted the screening burden. The funnel from 27,327 records to

Table 1: Query structure executed in each database and record hit breakdown.

Database	Executed Query Structure	Controlled Vocabulary	Fields Searched	Hits (N)
PubMed / MEDLINE	(B1 _{ft} [TIAB] OR B1 _{mesh}) AND (B2 _{ft} [TIAB] OR B2 _{mesh}) AND (B3 _{ft} [TIAB] OR B3 _{mesh}) NOT (X[TIAB])	Yes: 9 MeSH headings OR-ed into the three blocks	Title/Abstract + MeSH	24,693
SpringerLink	(B1) AND (B2) AND (B3) NOT (X); quoted phrases, no truncation operators	No	Full record	1,598
IEEE Xplore	(B1) AND (B2) AND (B3) NOT (X); wildcard truncation retained	No	Metadata	778
Scopus	TITLE-ABS-KEY(B1) AND TITLE-ABS-KEY(B2) AND TITLE-ABS-KEY(B3) AND NOT TITLE-ABS-KEY(X)	No	TITLE-ABS-KEY	127
Semantic Scholar	(B1) + (B2) + (B3), where denotes OR and + denotes AND	No	Title/Abstract	66
OpenAlex	title_and_-abstract.search:B1, title_and_-abstract.search:B2, title_and_abstract.search:B3	No	Title/Abstract	34*
Web of Science	TS=(B1) AND TS=(B2) AND TS=(B3) NOT TS=(X)	No	Topic (TS)	31
Total identified records				27,327
Concept blocks combine as (B1) AND (B2) AND (B3) NOT (X). B1 neuroimaging: "brain MRI", "cerebral MRI", neuroimag*, DWI, FLAIR, T1w, T2w, SWI; MeSH <i>Brain, Magnetic Resonance Imaging, Neuroimaging</i>. B2 AI and accelerated reconstruction: "MRI reconstruction", "accelerated MRI", "undersampled reconstruction", "k-space reconstruction", "unrolled network", GAN, "diffusion prior"; MeSH <i>Deep Learning, Machine Learning</i>. B3 safety, artifacts and evaluation: hallucinat*, "lesion suppression", reconstruction artifact*, PSNR, SSIM, "reader study"; MeSH <i>Artifacts, Diagnostic Errors, Reproducibility of Results</i>. X exclusion: "animal model", "case report", "in vitro". * denotes truncation. Complete term lists, MeSH headings and per-database syntax are given in Supplementary Table S5. Searches were executed between 2026-07-23 and 2026-07-24, with no publication-date limit. *OpenAlex returned 50 raw hits; 34 passed record-completeness filtering.				

263 studies is the intended signature of that choice: the burden falls on initial screening, which reviewers can absorb, rather than on coverage, which no downstream step could recover.

Eligibility and selection. Records were managed in the review's own Python

239 pipeline rather than a commercial screening platform; the scripts, the per-record
 240 decision files and the per-database query logs are released, so each step below
 241 is re-executable. Deduplication by exact identifier matching (PMID/DOI) and
 242 string/TF-IDF similarity (Jaro-Winkler ≥ 0.97 , cosine ≥ 0.95) removed 403
 243 duplicates (1.47%), leaving 26,924 unique records. That rate is low enough to
 244 suggest a matching failure but is structural: PubMed’s MeSH arm returned
 245 24,693 records and the six engineering arms 2,634 between them, and sets that
 246 barely overlap leave little to deduplicate. This is not a claim of completeness,
 247 and one failure is on record: peer review identified an arXiv preprint and its
 248 journal version that both survived to the included set, carrying different DOIs,
 249 no shared PMID, and titles differing by terminal punctuation — below the
 250 thresholds above. It was removed at revision, taking the corpus from 264 reports
 251 to 263 studies (Section 4). Automated matching will not catch a preprint–
 252 journal pair whose metadata diverges in this way, and the honest inference is
 253 that a residual duplicate rate of the same order cannot be excluded.

254 Three reviewers (D.T.M., T.V.P., T.N.T.D.) screened records against the
 255 pre-specified PICOS criteria (Table 2) at both stages, under a primary-screen-
 256 plus- adjudication design: each record received one primary label of *include*,
 257 *exclude*, *flagged* or *background only*, and every record labelled *flagged* went to
 258 a second reviewer for adjudication. Labelling was deliberately conservative,
 259 so any record with an unclear PICOS dimension, or whose text we could not
 260 obtain, was flagged rather than excluded. Title and abstract screening advanced
 261 1,149 records to full-text assessment, of which 109 were labelled *include* on the
 262 abstract alone and 1,040 *flagged*; both labels carried the record forward, and
 263 the distinction records only whether the abstract was sufficient to decide. A
 264 further 25,573 were excluded and 202 set aside as background. Full texts were
 265 pursued through institutional and title-level library subscriptions, open-access
 266 repositories, preprint servers, PubMed Central and inter-library request until
 267 no route remained: 1,068 of the 1,149 (93.0%) were obtained and assessed, and
 268 the 81 that could not be obtained are broken down by recorded reason with the
 269 selection results (Section 4).

270 The three reviewers divided the 1,068 obtained reports between them un-
 271 der the same primary-screen-plus-adjudication design used at title and abstract:
 272 each report received one primary full-text label, and a report was escalated to
 273 a second reviewer wherever the primary label was *flagged*, wherever an inclu-
 274 sion decision rested on a borderline PICOS dimension, or wherever the reviewer
 275 requested a second reading. Escalated reports were resolved by consensus dis-
 276 cussion between the two assessors, with the third reviewer consulted where the
 277 two did not converge. The escalation volumes this produced are reported below,
 278 and the per-report labels, escalation marks and reasons are released in the de-
 279 cision ledger. We state the design in this form because it determines what can
 280 be computed from the ledger: one primary label per report means no chance-
 281 corrected inter-rater coefficient is recoverable from full-text screening, which is
 282 why the separate criterion-reproducibility exercise below was run. In total, 263
 283 primary studies met all inclusion criteria upon full-text verification (two ad-
 284 ditional reports were excluded as duplicate publications). Full-text exclusions
 285 ($n = 785$) were categorized by pre-specified PICOS criteria: wrong intervention
 286 ($n = 326$), out of scope ($n = 256$), wrong outcome ($n = 98$), wrong population
 287 ($n = 93$), wrong study design ($n = 4$), non-English ($n = 1$), and non-primary
 288 articles ($n = 7$). A further 18 reports were retained as background material
 289 rather than as included studies, so the 1,068 assessed reports resolve as 263
 290 included, 785 excluded on PICOS grounds, 2 excluded as duplicate publications
 291 and 18 retained as background. The complete decision ledger is released as
 292 `fulltext_screening_decisions.csv`.

293 *Publication year.* Every analysis partitioning the corpus by time uses the
 294 year of the version of record as indexed by the publisher — the issue year for
 295 journal articles, the proceedings year for conference papers — rather than first
 296 online posting or preprint deposit. The rule matters at era boundaries, and we
 297 adopt the indexed year because it is what a reader can verify from the reference
 298 list, Crossref and PubMed without online-first metadata, which is unevenly
 299 available across the 87 venues represented. The same value is recorded in the
 300 `year` field of every bibliography entry and in the `publication_year` column of

301 `included_characteristics.csv`, so the era counts in Table 9 are recomputable
 302 from either. One consequence should not be read as an inconsistency: the
 303 searches ran on 23–24 July 2026, yet 41 of the 263 included studies (15.6%)
 304 carry a 2026 issue year, being articles published online ahead of print before the
 305 searches ran but assigned by the publisher to a 2026 issue.

306 *Screening agreement, and what it does and does not cover.* The primary
 307 screen described above records one label per record, so no chance-corrected co-
 308 efficient is recomputable from those decisions alone. We therefore ran a separate
 309 duplicate-screening exercise. A random sample of 214 records already assessed
 310 at title and abstract was re-screened independently by all three reviewers, blind
 311 both to the recorded decision and to one another. Reviewers judged the four
 312 PICOS dimensions from title, journal and abstract, opening the DOI link only
 313 where the abstract was insufficient, and returned one of three labels per record:
 314 include, exclude or unsure. The sample was drawn without stratification; the
 315 drawn records, the three label sets and the agreement computation are released,
 316 so the sample is auditable record by record. All 642 cells were completed and
 317 no record was dropped for missingness.

318 Across all 214 records Fleiss’ $\kappa = 0.877$ (95% bootstrap CI 0.823–0.925;
 319 10,000 resamples, seed 20260821), with observed agreement $\bar{P} = 0.933$ against
 320 $\bar{P}_e = 0.455$ expected by chance. The three reviewers agreed unanimously on
 321 194 of 214 records (90.7%), and pairwise Cohen’s κ ran from 0.845 to 0.931.
 322 One sensitivity analysis is worth reporting because it is the one a reader is most
 323 likely to run. Restricting to the 193 records on which no reviewer recorded
 324 *unsure* raises Fleiss’ κ to 0.943 (95% CI 0.900–0.979). That complete-case figure
 325 discards 21 records, and it discards precisely the records at least one reviewer
 326 could not decide, so it overstates agreement on the population the coefficient
 327 is meant to describe. We report 0.877 as the estimate and 0.943 only as the
 328 value obtained when the hard cases are removed. Two qualifications travel with
 329 that figure and should not be dropped. First, it measures agreement on the
 330 screening decision, not on the extraction labels from which every prevalence
 331 count in this article is derived; those remain unaudited for reliability, and the

only error rate we can quote for them is the measured 4/61 (6.6%) on the design field (Section 5.2). Second, the conservative labelling policy routed almost the whole full-text workload to adjudication in any case — 1,040 of the 1,149 records advanced at title and abstract were flagged rather than included outright — so a high coefficient is evidence that the criteria were applied consistently, not that the borderline records were easy to classify.

What the primary decision records themselves support, and what we therefore also report, are the escalation volumes that policy produced. At full text, 84 reports were escalated to a second reviewer for adjudication and 29 of those were included; a further 77 reports carry a mandatory human-re-inspection mark in the released files, of which 40 were included after that review. Screening recall is bounded by the search strategy and by the criteria as applied rather than estimated against an independent reference set of known-eligible records, since no such set exists for this question; Section 5.2 treats both points as limitations. The duplicate-screened sample, both label sets and the agreement computation are released as `kappa_human_results.csv`, alongside the full decision ledger in `fulltext_screening_decisions.csv`.

Data extraction. Reviewers extracted study characteristics using the pre-specified form (Supplementary Table S6), capturing method family, reported failure modes, evaluation paradigm, field strength, pulse sequence, reference standard and reproducibility indicators, with a verbatim supporting quotation for every assigned value. Method, failure-mode and evaluation categories were assigned from each record’s title, abstract and controlled vocabulary against the vocabulary fixed in advance (Supplementary Table S5); that deliberately conservative scope is why every prevalence count here is a lower bound rather than an exhaustive tally. Field strength and pulse sequence, typically stated only in acquisition subsections, were additionally taken from full text where available. We did not contact investigators: a variable not stated in the source was recorded as *not reported* and never inferred, which is why the 54 studies without a stated field strength divide into 50 whose full text does not state it and 4 assessed on abstract or partial text.

Table 2: PICOS elements and their application in this review.

Element	Definition	Application in this Review
Population	Target imaging modalities, clinical domains, and scanner physics	Human brain MRI acquisitions (acute ischemic stroke, neurovascular lesions, intracranial haemorrhages, tumours, and structural T1w, T2w, FLAIR, DWI scans across 1.5T, 3T, and 7T field strengths).
Intervention	AI-based and accelerated reconstruction methods, and their safety/fidelity evaluation	Deep learning and accelerated reconstruction methods (unrolled physics-based networks, GANs, diffusion priors, self-supervised, denoising, super-resolution), together with methods for evaluating their fidelity, artifacts, robustness, or failure modes.
Comparison	Baseline metrics, observers, and reference standards	Conventional pixel-wise metrics (PSNR, SSIM, NRMSE), perceptual/learned metrics (LPIPS, FID), model observers, fully sampled reference images, or radiologist reader assessments.
Outcome	Reconstruction quality, fidelity, and failure characterization	Image-quality/fidelity measures, artifact and failure-mode characterization, metric-radiologist agreement, uncertainty estimates, robustness to distribution shift, and diagnostic/observer outcomes.
Study Design	Methodological criteria and empirical validation settings	Primary empirical studies, algorithmic benchmarking, counterfactual interventional stress-testing, and human or model observer evaluations.

363 *Quality appraisal, and the reader/algorithmic partition.* Methodological
 364 quality was evaluated with three design-matched instruments: QUADAS-2
 365 (Whiting et al., 2011) across its four domains for diagnostic accuracy and
 366 radiologist reader studies; a reporting-transparency check against STARD prin-
 367 ciples (Bossuyt et al., 2015), reported descriptively rather than scored, since
 368 STARD yields no risk-of-bias rating; and, for algorithmic benchmarking studies,
 369 a reproducibility and robustness appraisal covering code and data availability,
 370 use of a named public benchmark, whether the reference standard was a fully
 371 sampled acquisition, and whether robustness to acquisition or distribution
 372 shift was reported. ROBINS-I (Sterne et al., 2016) was not applied, since few
 373 included studies are the non-randomised interventional designs it was built for.

374 The partition between the two instruments is not the one the extraction la-
 375 bels originally implied, and the correction matters for every reader-share statis-
 376 tic here. Extraction recorded 61 studies as observer or reader designs; three-

reviewer appraisal of their full text found four containing no human reader at all, each using the language of observer evaluation for an automated surrogate. Those four were reclassified as algorithmic, so **the verified reader subset is 57 and the algorithmic set is 206**, and $57 + 206 = 263$ closes the corpus exactly. All 57 were appraised by all three reviewers across two rounds, 39 in the first and 18 in the second, so every QUADAS-2 tally rests on 57 studies and $57 \times 7 = 399$ domain judgements. Every reader-derived count in this article uses 57, and the released per-study file encodes the same partition with the reclassification recorded per study. One further study (S001) carries an appraisal note that a human reader could not be *confirmed*, only abstract-level text being available; it is retained in the reader subset, since non-ascertainment is not a finding of absence, and it is flagged in the deposit.

Appraisal agreement. The three reviewers rated every applicable domain Low, Unclear or High independently, each judgement carrying a verbatim supporting quotation, before joint review to harmonize thresholds. Five first-round studies were later re-appraised on recovered full text (below), so their 35 judgements no longer rest on three independent ratings and are excluded from the coefficient. Over the remaining 52 studies and 364 judgements Fleiss' $\kappa = 0.788$ (95% bootstrap CI 0.739–0.833; 10,000 resamples, seed 20260821). That figure averages two rounds that behaved differently, and the difference is more informative than the average. The 34 first-round studies retaining three independent ratings were unanimous on all 238 of their judgements, so they contribute $\kappa = 1.00$ by construction and carry no information about latent reliability. The 18 of the second round were unanimous on 60 of 126 (47.6%), and across those 126 Fleiss' $\kappa = 0.390$ (95% CI 0.293–0.484) — fair on the Landis–Koch bands, with an interval reaching into moderate. The procedure was identical across rounds; the subsets were not, the second round comprising studies the first pass deferred as not straightforward. Two readings remain open — that those studies report their methods less completely and are genuinely harder to rate, and that the first round was less independent in practice than in design — and these data do not separate them.

408 The disagreement has a structure worth recording rather than smoothing.
409 Across the 126 second-round judgements one reviewer rated Low 33.3% of the
410 time and another 61.9%, a difference of threshold rather than of reading, which
411 is why κ sits at 0.390 while raw agreement is 63.5%. Six judgements split
412 three ways with no majority; all six were recorded as *Unclear*, on the ground
413 that a panel unable to converge is itself evidence the report did not support
414 a confident judgement, which is what *Unclear* encodes in QUADAS-2. The
415 rule was applied to those six and no other case, and the per-reviewer labels are
416 released so a reader can apply a different rule and recompute. Section 5.2 states
417 what this does and does not establish. Table 3 summarizes corpus-level quality
418 indicators and the QUADAS-2 domain judgments.

Table 3: Corpus-level provenance and methodological quality indicators across the 263 included studies, with QUADAS-2 domain-wise risk of bias and applicability concerns for the full appraised diagnostic/reader subset ($n = 57$, rated across seven domains; independently by all three reviewers except for the five studies re-appraised on recovered full text). The 39 studies of the first appraisal round and the 18 of the second are pooled here; the remaining 206 studies were appraised on reproducibility and robustness instead ($57 + 206 = 263$). Six of the 399 domain judgements split three ways with no majority and are recorded as *Unclear* (Section 3).

Indicator / Domain	Value / Risk of Bias Breakdown
Publication period	1995–2026; 229/263 (87.1%) published in 2020 or later, and 174/263 (66.2%) in 2023 or later, indicating a predominantly recent evidence base.
Basis of assessment, current	Full text available for 258/263 (98.1%); partial or truncated full text for 1/263 (0.4%); title, abstract and MeSH only for 4/263 (1.5%).
Basis of assessment, at appraisal	The first appraisal round was made on the text obtainable at the time, which for 18/263 studies was abstract and MeSH only (6 of them within the QUADAS-2 subset). A later library sweep recovered full text for 14 of those 18 (including 5 of the 6), and all 14 have since been re-appraised on the recovered full text, so appraisal basis now stands at 257/263 full text, 2/263 partial text and 4/263 abstract and MeSH only — within the QUADAS-2 subset, 56/57 on full text and 1/57 (S001) on abstract-level text. The two rows now differ by a single study (S200), appraised on truncated text whose full version has since become available.
Leading venues	The eight most frequent of the 87 distinct venue titles represented in the corpus, accounting for 128/263 (48.7%) of included studies: <i>Magnetic Resonance in Medicine</i> (40), <i>IEEE Transactions on Medical Imaging</i> (17), <i>AJNR</i> (15), <i>Magnetic Resonance Imaging</i> (14), <i>JMRI</i> (12), <i>Scientific Reports</i> (12), <i>NeuroImage</i> (9), <i>Medical Image Analysis</i> (9).

(Table 3 continued)

Predominant design	Algorithmic / methodological reconstruction studies ($n = 206$), alongside a verified observer/reader subset ($n = 57$) (Radmanesh et al., 2022; Lang et al., 2024). Extraction recorded 61 reader designs; appraisal of their full text found 4 with no human reader, and those 4 were reclassified as algorithmic ($202 + 4 = 206$; $61 - 4 = 57$). Of the 57: 34 are appraised and verified by all three reviewers; 18 were appraised by all three in a second round and await final adjudication of 6 split domain judgements; and 5 (S004, S020, S027, S034, S115) were re-appraised by the reviewers on full text recovered after the first round, with domain judgements verified against the source text. Every reader-derived statistic in this article uses 57.
Reference standard	Expert reader consensus (70), an unspecified stated ground truth (62), a fully sampled acquisition (69), a conventional or standard-of-care acquisition (15); not stated in 47/263 (17.9%).
Reproducibility	Code availability stated in 97/263 (36.9%); a named public dataset or benchmark used in 109/263 (41.4%), most commonly fastMRI, HCP, ADNI and IXI. Availability is variable and remains an area for improvement.
Sampling and setting	Sampling frame not stated in 95/263 (36.1%); recruitment centre not stated in 174/263 (66.2%). Ethics approval or informed consent reported in 159/263 (60.5%).

QUADAS-2 Risk of Bias ($n = 57$ Reader / Diagnostic Accuracy Studies)

Patient Selection	Low: 15/57 (26.3%); Unclear: 19/57 (33.3%); High: 23/57 (40.4%). The weakest domain: sampling is predominantly purposive or pathology-enriched, and most of the appraised studies state no enrolment rule.
Index Test (AI Reconstruction)	Low: 26/57 (45.6%); Unclear: 9/57 (15.8%); High: 22/57 (38.6%). Low risk falls below half once the second round is included. High ratings arise where readers scored accelerated reconstructions with the reference visible, or where test parameters were set using reference-derived labels.
Reference Standard	Low: 26/57 (45.6%); Unclear: 15/57 (26.3%); High: 16/57 (28.1%). High ratings arise where no artifact-free reference exists, or where lesion truth is defined by the comparator images themselves (incorporation bias). Reading the recovered full text moved four of the re-appraised studies from Low to High here, each because the comparator images, or software applied to them, supplied the truth against which they were scored.
Flow and Timing	Low: 45/57 (78.9%); Unclear: 7/57 (12.3%); High: 5/57 (8.8%). The strongest domain, since retrospective undersampling of the same raw k -space gives a zero interval between index test and reference.

QUADAS-2 Applicability Concerns ($n = 57$)

Patient Selection	Low: 36/57 (63.2%); Unclear: 4/57 (7.0%); High: 17/57 (29.8%), chiefly healthy-volunteer, lesion-free, or non-clinical cohorts.
-------------------	---

(Table 3 continued)

Index Test	Low: 30/57 (52.6%); Unclear: 4/57 (7.0%); High: 23/57 (40.4%), where the evaluated method is image-domain post-processing, cross-field synthesis, harmonization, or a quality index rather than reconstruction from under-sampled k -space.
Reference Standard	Low: 39/57 (68.4%); Unclear: 3/57 (5.3%); High: 15/57 (26.3%).

Across all 228 risk-of-bias judgements: 112 (49.1%) Low, 50 (21.9%) Unclear, 66 (28.9%) High. Across all 171 applicability judgements: 105 (61.4%) Low, 11 (6.4%) Unclear, 55 (32.2%) High. Over the 364 judgements that rest on three independent ratings Fleiss' $\kappa = 0.788$ (95% CI 0.739–0.833), but agreement is not uniform across the two appraisal rounds: the 238 retained first-round judgements were unanimous, the 126 of the second round reach $\kappa = 0.390$ (95% CI 0.293–0.484); the 35 judgements of the 5 re-appraised studies are excluded, having been conducted subsequently on recovered full text. Section 3 gives the decomposition and Section 5.2 what it bounds. The high proportion of Unclear ratings reflects unreported blinding and sampling procedures, compounded by the one first-round study for which only abstract-level text is available. Six second-round judgements split three ways with no majority and are recorded as Unclear under the rule stated in Section 3. Per-study ratings, each with the verbatim supporting quotation, are released as `included_characteristics_supplement.csv`.

419 *Derived venue classification.* The analysis in Section 4.5 groups studies by
420 publication community. That grouping is post-hoc, was not part of the fixed
421 protocol, and is stated in full in Supplementary Table S13 so it can be re-
422 produced or contested. Each study’s journal title, and nothing else about the
423 study, is matched case-insensitively against three ordered keyword lists — en-
424 gineering and methods, then general neuroscience and neuroimaging, then clin-
425 ical imaging — with anything unmatched left *unclassified* rather than forced
426 into a group. The ordering is load-bearing, since titles such as *Medical Image*
427 *Analysis* would otherwise match two lists, and it is applied identically to all
428 263 studies. The rule as executed and the resulting per-study assignment are
429 released as `5_Code/venue_classification.py` and `venue_communityclas-`
430 `sification.csv`.

431 *Synthesis.* Due to heterogeneity in acquisition protocols, reconstruction ar-
432 chitectures, datasets, and outcome measures, evidence was synthesised narra-
433 tively in accordance with SWiM guidelines (Campbell et al., 2020). Studies were
434 structured around the primary research questions (RQ1–RQ3) and grouped by
435 reconstruction method family, field strength, pulse sequence, and failure mode.
436 Quantitative values reflect named primary studies rather than pooled meta-
437 analytic estimates; no effect measure is pooled, and no meta-analytic model

438 is fitted. Percentages are reported to one decimal place and rounded half-up
439 throughout, so 9/80 prints as 11.3% and not 11.2%. Sensitivity analysis was
440 performed by recalculating headline prevalence counts on the 258 studies verified
441 against complete full text (Section 5.2).

442 4. Results

443 All counts and prevalence statements in this section describe the 1,068 re-
444 ports assessed in full and the 263 included studies; because 81 candidate reports
445 remain unretrievable, they are lower bounds on the literature rather than esti-
446 mates of it.

447 4.1. Selection and corpus

448 From 27,327 records identified across seven databases, multi-tier dedupli-
449 cation removed 403 duplicates, leaving 26,924 unique records screened at title
450 and abstract. Of these, 1,149 proceeded to full-text eligibility, and full texts
451 were obtained for 1,068 of them (93.0%; Section 3). Each of the 81 reports still
452 out of reach was checked against its publisher platform and assigned a recorded
453 reason: 53 sit behind title-level subscriptions our institution does not hold,
454 across nineteen publishers; 13 are Springer conference chapters sold only per
455 chapter; 4 are ISMRM proceedings abstracts issued without a PDF; 4 resolve
456 to a whole proceedings volume or to front matter rather than a single report;
457 and 7 carry no DOI and no locatable source. Supplementary Table S14 gives
458 the per-publisher breakdown. Access difficulty is itself informative, though it
459 bounds the missing evidence rather than characterising it. Of the 341 assessed
460 reports obtainable only through title-level subscriptions and equivalent barriers,
461 294 (86.2%) proved ineligible on the same PICOS criteria, 2 were retained as
462 background and 45 were eligible ($294 + 2 + 45 = 341$; the 45 resolve to 44 stud-
463 ies, one being a preprint adjudicated a duplicate). Raw non-retrieval counts
464 therefore overstate the missing evidence by roughly sevenfold. But those 45
465 were concentrated in subscription clinical journals, and Section 4.5 shows they

differ systematically from the open literature in exactly the variable this review measures, so 86.2% is an upper bound on what non-retrieval costs us, not a demonstration that it costs us nothing. Of the 1,068 assessed reports, 787 were excluded (785 on PICOS grounds, 2 as duplicate reports of included studies) and 18 recorded as background, leaving **263 included primary studies** spanning 1995–2026 (Fig. 1). Each included study was verified against the fullest text obtainable for it: complete full text for 258, partial text for 1, and title, abstract and controlled vocabulary for 4 (Section 3). Their aggregate characteristics are given in Table 4, and their thematic distribution across the research questions is summarized in Supplementary Table S7; all 263 are listed, with full references, in Supplementary Tables S1–S3.

Corpus growth is recent: 34 studies appeared in 1995–2019, 55 in 2020–2022 and 174 in 2023–2026.

4.2. RQ1: the metric–reader comparison is almost never made, and where readers appear they agree only partially

Across the 84 studies reporting objective fidelity metrics and the 57 verified reader studies, the comparisons that actually exist show agreement between fidelity metrics and radiologist assessment that is partial and condition-dependent (Table 5). The size of that evidence base should be stated in the text and not left to the table caption: only two of the four rows in Table 5 were collected on reconstructions of undersampled k -space. Tang et al. (2025) collected radiologist rankings on synthetically corrupted image pairs, and Sommer et al. (2020) on clinical scans corrupted by real patient motion rather than by undersampling. The table that carries RQ1 therefore rests on two on-target rows, which is the honest size of the direct evidence for the metric–reader comparison in this setting. In the fastMRI 2020 challenge’s reader phase, six radiologists ranked submissions for quality of pathology depiction; their concordance, by Kendall’s coefficient, was 0.457 in the $4\times$ track and 0.386 in the $8\times$ track, rising to 0.781 only in the transfer track, and the organisers report that concordance improved as SSIM scores diverged, which is to say the metric separated methods

Table 4: Characteristics of the 263 included studies.

Demographic / Method Dimension	Category Breakdown	(n)	Percentage / Key Methodological Notes
Publication Period	1995–2019	$n = 34$	12.9% (Classical CS/PI & early CNNs)
	2020–2022	$n = 55$	20.9% (Unrolled physics & GAN priors)
	2023–2026	$n = 174$	66.2% (Diffusion priors & 3D/4D foundation models)
Reconstruction Paradigm (multi-label; 339 labels recorded)	Compressed Sensing / Parallel Imaging	$n = 93$	Iterative sparsity & coil sensitivity encoding
	Generative (GAN/Diffusion/VAE)	$n = 80$	Adversarial perception & posterior sampling
	Motion Correction	$n = 57$	Intra-scan motion artifact suppression
	Physics-based Unrolled	$n = 37$	Explicit k -space data consistency layers
	Denoising	$n = 30$	Image-domain noise suppression
	Super-resolution	$n = 25$	Resolution enhancement post-acquisition
MRI Pulse Sequences (multi-label; 490 labels)	Self-supervised / Unsupervised	$n = 17$	Training without fully sampled reference scans
	T1-weighted (T1w)	$n = 163$	High-resolution anatomical cortical profiling
	T2-weighted (T2w)	$n = 122$	Ventricular & CSF pathology assessment
	FLAIR	$n = 85$	Hyperintense peri-ventricular lesion detection
	Diffusion-Weighted (DWI)	$n = 72$	Acute ischemic lesion (DWI core) evaluation
	SWI / T_2^* -weighted	$n = 48$	Punctate cerebral microbleed hypointensity voids
Magnetic Field Strength (single-label; predominant field where several are used)	3.0 Tesla (3T)	$n = 153$	58.2% (Primary clinical scanner benchmark)
	7.0 Tesla (7T)	$n = 35$	13.3% (Ultra-high-field research acquisitions)
	1.5 Tesla (1.5T)	$n = 21$	8.0% (Routine clinical throughput imaging)
	Not stated	$n = 50$	19.0% (Full text obtained; no field strength stated)
	Not ascertainable	$n = 4$	1.5% (Only abstract or partial text available)
Evaluation Paradigm (multi-label; 177 labels recorded across 144 studies)	Pixel-wise (PSNR/SSIM/NRMSE)	$n = 84$	47.5% of recorded (legacy fidelity reporting)
	Observer / Reader Studies	$n = 57$	32.2% of recorded (task-based reader assessment; verified subset)
	Uncertainty Estimation	$n = 18$	10.2% of recorded (posterior variance maps)
	Learned / Perceptual	$n = 18$	10.2% of recorded (rank-trained metrics)
Co-reporting. Of the 144 studies with a recorded evaluation paradigm, 28 report more than one. Critically, only 18 studies in the entire corpus (6.8% of 263; 12.5% of the 144 recorded) report <i>both</i> a pixel-wise fidelity metric <i>and</i> a reader assessment on the same data, which is the minimum precondition for computing metric-versus-reader agreement on which safety claims implicitly depend. See Section 4.			




Figure 1: PRISMA 2020 flow diagram for study selection across the two search arms. Multi-tier deduplication reduced 27,327 identified records to 26,924 unique records; title and abstract screening advanced 1,149 candidates. Full texts were obtained for 1,068 of them (93.0%); 81 could not be retrieved. Following PRISMA 2020, exclusion and background retention are recorded as distinct dispositions: of the 1,068 reports assessed in full, 787 were excluded (785 on explicit PICOS grounds and 2 as duplicate reports of included studies) and a further 18 were retained as background material rather than excluded, leaving 263 primary studies in the analytical set ($263 + 787 + 18 = 1,068$).

reliably only where their outputs already differed grossly (Muckley et al., 2021). When radiologist preference is instead used as the training signal, purpose-built quality networks predict expert rankings at $75.2\% \pm 1.3\%$ and $79.2\% \pm 1.7\%$, against radiologists’ own interobserver agreement of $61.5\% \pm 5.5\%$; the same study finds rank-based reading far more reliable than Likert scoring (70.4% versus 25% agreement between radiologists) (Tang et al., 2025). Reader tolerance of acceleration is real but bounded: in a retrospective evaluation spanning accelerations up to $100\times$, 94% of volumes (65 of 69) were rated sufficient for potential screening use at up to $14\times$ (interreader intraclass correlation 0.875), with the threshold for full diagnostic quality lower still (Radmanesh et al., 2022). The one direct clinical observation of the failure this review is about comes from an

507 emergency-department and inpatient comparison, and it is a single case. Across
 508 66 paired examinations in which a two-minute ultrafast protocol (T1w, T2/T2*,
 509 FLAIR and DWI, machine-learning-assisted reconstruction) was read against a
 510 ten-minute reference protocol, two blinded neuroradiologists agreed on the main
 511 clinical diagnosis in 65 of 66 cases (98.5%), rated overall diagnostic quality no
 512 different between protocols ($P > .05$), preferred the reference protocol for image
 513 noise and geometric distortion ($P < .05$) and preferred the accelerated protocol
 514 for reduced motion artifact ($P < .01$) (Lang et al., 2024). The single discrepant
 515 case is the one that matters here: a punctate focus of restricted diffusion, read
 516 as possible acute-to-subacute infarct on the reference DWI, was less conspicu-
 517 ous on the accelerated DWI, which the authors attribute to a combination of
 518 image-quality difference and section positioning, the lesion being smaller than
 519 the section thickness (4 mm reference against 5 mm accelerated) and so subject
 520 to partial-volume contamination. They draw the inference themselves, that the
 521 protocol’s sensitivity to small findings is an open question. One case in 66 is
 522 not a prevalence estimate and we do not offer it as one. It is a worked instance
 523 of the pattern this review predicts: an aggregate quality comparison that finds
 524 no difference, and inside it a focal finding that a reader saw and the aggregate
 525 did not. Where readers assess artifacts directly, blinded reading detects what
 526 reference-based metrics cannot express: a two-neuroradiologist blinded study of
 527 28 clinical scans with real patient motion found significant artifact-score reduc-
 528 tions after network correction ($p < .03$), a setting in which the motion-corrupted
 529 reference invalidates PSNR and SSIM by construction (Sommer et al., 2020).
 530 Motion also interacts with the sampling scheme itself, appearing as ghosting
 531 under Cartesian trajectories and as streaking or blurring under radial ones,
 532 so the visual form of the same physical corruption differs by protocol (Schau-
 533 man et al., 2026). The reference pipeline is not neutral either: conventional
 534 SWI homodyne phase filtering has been shown to create punctate artifactual
 535 microhaemorrhage-like foci, so processing can manufacture findings as well as
 536 remove them (Li et al., 2015).

537 No included study reports a rank correlation between a pixel-wise fidelity

metric and radiologist diagnostic assessment for deep learning-based brain reconstruction, and none evaluated a model observer. The second statement needs a boundary that the first does not. Our search vocabulary (Table 1, block B3) contains reader-study and metric terms but no model-observer vocabulary, so a study that used a channelised Hotelling or non-prewhitening observer and described it only in those terms could have been missed at retrieval. What we can say is that none of the 263 studies retrieved, screened and extracted here reports one, and that the term does not appear in any extracted evaluation label; whether the true field-wide count is zero or merely very small, it is not a paradigm this corpus can be said to use. That absence is itself the RQ1 finding: only 18 of 263 studies (6.8%, a screening-level lower bound; Section 3) record a fidelity metric and a reader assessment on the same data, so the field’s central proxy relationship, that high fidelity scores imply diagnostic adequacy, is largely unmeasured rather than merely disputed.

Table 6 reports how evaluation practice varies with the acceleration factor actually stated in each study. It is built only from what the corpus records: the acceleration factor was extracted post hoc from the full texts of all 263 included studies with a verbatim supporting quotation for every value, and 63 studies (24.0%) state one. Studies are grouped by the highest factor at which they evaluate their method. The table deliberately carries no fidelity threshold, reader-score difference or metric–reader rank statistic per acceleration band, because no included study reports one: that absence is the finding of Section 4, not an omission from the table.

4.3. RQ2: four mechanisms explain the divergence

Four mechanisms account for why the metrics miss what the readers see. Table 7 states each one formally alongside the clinical failure it permits, and Supplementary Table S8 gives the fuller derivations and evidence sources.

Spatial averaging. Conventional L_2 metrics average residual error uniformly over all $H \times W \times D$ voxels, so the total decomposes over lesion and background regions:

Table 5: What the reader studies in this corpus actually measured about metric–reader agreement (RQ1). Every value is quoted from the named study, with the reported coefficient and its definition stated. No included study reports a metric–reader rank correlation for deep learning-based brain reconstruction; the rows below are the closest measurements that exist, and the second column states for each row what the readers were actually shown, because two of the four rows were not collected on reconstructions of undersampled k -space.

Study	Readers and design	Comparison performed	Reported quantity (as defined in the study)
fastMRI 2020 challenge (Muckley et al., 2021)	6 radiologists ranking submissions for quality of pathology depiction	Radiologist concordance across accelerations; SSIM-based ranking vs reader ranking	Kendall's $W = 0.457$ ($4\times$), 0.386 ($8\times$), 0.781 (transfer track); concordance improved as SSIM scores diverged
Tang et al. (2025)	7 radiologists, rank-based reading; 2 also Likert scoring. Rankings were collected on synthetically corrupted image pairs, not on reconstructions of undersampled k-space	Quality networks trained on radiologist rankings vs MSE/SSIM as optimisation targets	Ranking accuracy $75.2\% \pm 1.3\%$ and $79.2\% \pm 1.7\%$ with reference image; radiologist interobserver agreement $61.5\% \pm 5.5\%$; ranking vs Likert agreement 70.4% vs 25%
Radmanesh et al. (2022)	Retrospective; neuroradiologist scoring of diagnostic quality at accelerations up to $100\times$	Reader-assigned acceleration thresholds for diagnostic and screening quality	$65/69$ volumes (94%) sufficient for potential screening at $\leq 14\times$; interreader ICC 0.875 ; mean absolute interreader difference in threshold acceleration factor $0.7\times$ (level 1) and $5.4\times$ (level 2)
Sommer et al. (2020)	2 neuroradiologists, blinded, 28 clinical cases with real motion	Artifact severity (5-point scale) before vs after network correction	Significant reduction in mean artifact scores ($p < .03$); reference-based metrics inapplicable, as the reference itself is motion-corrupted

$$\begin{aligned} \text{MSE}(\hat{X}, X) &= \frac{V_{\text{lesion}}}{V_{\text{total}}} \text{MSE}_{\text{lesion}} \\ &+ \left(1 - \frac{V_{\text{lesion}}}{V_{\text{total}}}\right) \text{MSE}_{\text{bg}}. \end{aligned} \quad (1)$$

Table 6: Evaluation practice by the highest acceleration factor (R) each study reports, among the 63 of 263 included studies that state one. Bands are half-open, so a study reporting exactly $R = 4$ falls in the second band and one reporting exactly $R = 8$ in the third. Counts are studies; percentages are within band. *Reader*, any observer or reader assessment recorded; *M+R*, a pixel-wise fidelity metric and a reader assessment recorded on the same data; *Halluc.*, hallucination or fidelity instability among the recorded failure modes; *Artifacts*, aliasing, ghosting or Gibbs artifacts among them. Acceleration factors were extracted from full text after the registered extraction round, each with a verbatim supporting quotation and excluding matches inside reference lists and citation contexts; the 407 retained quotations come from the 63 studies that state a factor, a mean of 6.5 quotations per such study, and are released as `acceleration_factors_evidence.csv`. The extraction is screening-level, so counts are lower bounds. No study at any acceleration factor reports a rank correlation between a fidelity metric and reader judgement.

Acc. (R)	Studies	Reader	M+R	Halluc.	Artifacts	Sequences most often evaluated
$2 \leq R < 4$	32	7 (21.9%)	1 (3.1%)	7 (21.9%)	25 (78.1%)	T1w 15, T2w 12, FLAIR 11
$4 \leq R < 8$	17	5 (29.4%)	1 (5.9%)	2 (11.8%)	12 (70.6%)	T1w 9, T2w 9, FLAIR 7
$R \geq 8$	14	5 (35.7%)	4 (28.6%)	2 (14.3%)	8 (57.1%)	T1w 9, T2w 8, FLAIR 8
All	63	17 (27.0%)	6 (9.5%)	11 (17.5%)	45 (71.4%)	T1w 33, T2w 29, FLAIR 26

The R -reporting subset is close to the whole corpus on reader assessment (27.0% against 21.7% for all 263), so it is not a systematically better-evaluated stratum. Reader assessment rises with acceleration (21.9% $\rightarrow$ 29.4% $\rightarrow$ 35.7%) and paired metric-reader reporting concentrates in the most accelerated band (4/14 against 1/32 in the $2 \leq R < 4$ band), but the counts are small and most studies in every band (including 9 of the 14 with $R \geq 8$) report no reader assessment at all. Artifact description is the most frequently recorded failure mode at every acceleration level, whereas hallucination or fidelity instability is recorded in 11 of 63.

568 With lesion volume fraction $f = V_{\text{lesion}}/V_{\text{total}}$ and local error ratio $r =$
569 $\text{MSE}_{\text{lesion}}/\text{MSE}_{\text{bg}}$, the global PSNR change caused by a localized edit follows
570 directly:

$$|\Delta\text{PSNR}| = 10 \log_{10} (1 + f(r - 1)) \quad (\text{dB}). \quad (2)$$

571 The consequence is quantitative rather than rhetorical, and it is worth anchoring
572 on a lesion a neuroradiologist would expect to resolve rather than on the smallest
573 arithmetically convenient one. An acute lacunar infarct of 100 mm^3 is a sphere
574 5.8 mm across: on a clinical DWI acquired at 4–5 mm section thickness (Lang
575 et al., 2024) that is several voxels wide in-plane and spans one to two sections, so
576 it is resolved, if only just, in the through-plane direction. It occupies a volume
577 fraction $f = 6.3\text{--}8.3 \times 10^{-5}$ of a 1.2–1.6 L brain. Even at a hundredfold local

Table 7: The four mechanisms by which a whole-volume fidelity score fails to register a focal diagnostic error (RQ2). The third column names a failure the mechanism permits, not one this review has observed at a stated frequency; the fourth states why the score does not move. Numerical entries are computed from Equation (2) under the stated error model, not measured in any primary study.

Mechanism	Formal statement	Clinical failure it permits	Why the score does not move
(a) Spatial error averaging	L_2 error is averaged uniformly over all $H \times W \times D$ voxels, so a lesion contributes with weight $f = V_{\text{lesion}}/V_{\text{total}}$	Erasure of an acute lacunar infarct on DWI or FLAIR, or of a punctate microhaemorrhage on SWI or T_2^*	By (2), erasing a 100 mm^3 infarct at $r = 100$ moves global PSNR by $0.027\text{--}0.036 \text{ dB}$; a 5 mm^3 microhaemorrhage by $1.3\text{--}1.8 \times 10^{-3} \text{ dB}$
(b) Reference dependence	Every reference-based score presupposes a registered, artifact-free reference image	A valid motion-corrected reconstruction is penalised against a corrupted reference, and genuine correction is scored as error	The comparison is no longer between reconstruction and truth; PSNR and SSIM become uninterpretable rather than merely noisy
(c) Loss symmetry	$(\hat{X}_i - X_i)^2 = (X_i - \hat{X}_i)^2$: the objective is indifferent to the sign of the error	A suppressed hyperintense lesion is scored identically to added background noise of the same magnitude	Diagnosis is asymmetric where the metric is not: a missed finding and a spurious brightening carry the same penalty
(d) Realism-reward trade-off	Adversarial and perceptual objectives reward outputs typical of the training distribution, $\log D_\phi(\hat{X})$, not outputs faithful to the measurement	Synthesis of plausible vessels, sulci or texture that was never acquired	Perceptual scores can rise as measurement fidelity falls, so the objective pays for the hallucination it should penalise

error ratio ($r = 100$), Equation (2) gives $|\Delta\text{PSNR}| = 0.027\text{--}0.036 \text{ dB}$. Erasing a resolvable acute infarct outright therefore moves global PSNR by about three hundredths of a decibel, which is the same order as the third decimal place of a score conventionally reported to two.

The limiting case is worth stating precisely, because it is easy to state loosely. A punctate cerebral microhaemorrhage of 5 mm^3 is 2.1 mm across. It is *not* resolved on clinical DWI, whose voxels are roughly $12\text{--}20 \text{ mm}^3$, so no reconstruction can be said to erase it there; it is not resolved in the first place. It is several voxels across on the high-resolution SWI or T_2^* -weighted acquisition on which microhaemorrhages are actually read, and on that sequence its era-




Figure 2: The PSNR cost of erasing a focal lesion, from Equation (2). Curves give $|\Delta\text{PSNR}|$ against erased lesion volume for three local error ratios r at a 1.4 L brain; the shaded band around $r = 100$ spans the 1.2–1.6 L range. Markers show the three clinical anchors used in the text. The horizontal line at 0.01 dB is the two-decimal precision at which PSNR is conventionally reported, and the shaded left-hand region is the volume below a single voxel on a clinical DWI acquisition, where erasure is not defined because the finding is not resolved. Every value is computed from the identity under the stated error model; nothing in this figure is measured.

sure costs $f = 3.1\text{--}4.2 \times 10^{-6}$ and $|\Delta\text{PSNR}| = 1.3\text{--}1.8 \times 10^{-3}$ dB, below the precision at which the score is reported at all. Scaling in the other direction, a 500 mm³ cortical infarct costs ≈ 0.15 dB (0.13–0.18 dB across the same brain-volume range). Fig. 2 plots the identity across this range, with the sub-voxel region marked. The argument does not need the extreme case: it holds by one to two orders of magnitude at a lesion size that is unambiguously resolved and unambiguously matters. These are computed bounds from the stated f and r , not measurements from any primary study, and the error model behind them, a uniform background error and a free local error ratio, is an assumption we state rather than a property of any reconstruction.

Loss-function symmetry. MSE and SSIM are indifferent to the sign of the error, since $(\hat{X}_i - X_i)^2 = (X_i - \hat{X}_i)^2$, whereas diagnosis is not: suppressing a true hyperintense lesion means the finding is missed, while adding background

noise costs nothing clinically. An evaluation respecting that asymmetry must penalise the two directions differently, for example

$$\begin{aligned}\mathcal{L}_{\text{asym}}(\hat{X}, X) = & \alpha \sum_{i \in \Omega_{\text{lesion}}} \max(0, X_i - \hat{X}_i)^2 \\ & + \beta \sum_{i \in \Omega_{\text{bg}}} (\hat{X}_i - X_i)^2, \quad (\alpha \gg \beta).\end{aligned}\tag{3}$$

Several included studies already build objectives of this shape: lesion-mask-weighted losses motivated explicitly by the dilution of lesion gradients among background voxels (Lee et al., 2026a), overlap-weighted reconstruction losses that up-weight structure against background (Jayakumar et al., 2023), and gradient-based structure-preservation terms (Zhang et al., 2026). Their existence is evidence that the symmetry problem is recognised in practice, though none of these studies pairs the correction with a reader assessment.

Realism–reward trade-off. Adversarial and perceptual objectives reward outputs that look like the training distribution, which is not the same as outputs faithful to the measurement, so a generative reconstruction can gain plausible texture that was never acquired (Zhou et al., 2024; Bhadra et al., 2021).

Reference dependence. All four metrics presuppose a registered, artifact-free reference. Intra-scan motion invalidates that assumption, and 69 of 263 studies rest their reference standard solely on a fully sampled acquisition, which is valid only where the subject did not move.

4.4. RQ3: what current evaluation paradigms can and cannot establish

Read against those mechanisms, the corpus supports a paradigm-by-paradigm account of what each evaluation approach can demonstrate about prior-induced diagnostic error. Table 8 states the account compactly; the argument behind it is given here, because two of its six rows are analytical claims we derive rather than findings we extracted, and an analysis presented only in a table has not been argued.

What the four empirical paradigms establish. Pixel-wise fidelity metrics establish aggregate agreement with a reference and nothing about focal change:

by (2) the penalty for erasing a resolvable acute infarct is of the order of the reporting precision, so a method can pass any threshold anyone would set while failing on the finding that matters. Perceptual and rank-trained metrics track reader preference better than MSE or SSIM as an optimisation target, but they inherit the variability of the readers they were trained on and, in the one study of this kind in the corpus, were calibrated on synthetically corrupted images rather than on reconstructions of undersampled k -space (Tang et al., 2025), so their transfer to the reconstruction setting is assumed rather than shown. Uncertainty estimates flag where a model is unsure, which is not the same question: a prior confident enough to remove a lesion is confident in the region where it removed it, so posterior variance can be low exactly where the diagnostic error is. Task-based observer assessment does settle the question directly, and it is the only paradigm here that does, which is why its scarcity in this corpus (57 of 263 studies, and 18 paired with a metric) is the review’s central finding rather than an incidental one.

Two limits that follow from the physics, not from the corpus. The first concerns the check that reconstruction papers most often use to argue that an output is faithful to the data: the measurement-space residual $\|\mathbf{A}\hat{X} - Y\|_2 / \|Y\|_2$ against a tolerance. This check cannot detect prior-induced edits at all, and the reason is structural rather than statistical. Write the reconstruction as the sum of its measured and unmeasured components, $\hat{X} = \mathbf{A}^+ \mathbf{A} \hat{X} + (\mathbf{I} - \mathbf{A}^+ \mathbf{A}) \hat{X}$. The second term lies in the null space $\mathcal{N}(\mathbf{A})$, and $\mathbf{A}$ annihilates it by definition, so any amount of structure written there, an erased lesion or an invented vessel alike, leaves the residual exactly unchanged. Undersampling is what creates that null space; the more aggressive the acceleration, the larger the subspace in which a learned prior can write freely without moving the number the field uses to certify data consistency. A residual check is therefore a two-sided measurement-fidelity band and not a safety test, and its apparent reassurance grows least trustworthy precisely where acceleration is highest.

The second limit is the constructive corollary. If prior-supplied content lives in $\mathcal{N}(\mathbf{A})$, then isolating $(\mathbf{I} - \mathbf{A}^+ \mathbf{A}) \hat{X}$ and inspecting it is the natural place to

look for erasure and hallucination, since that is where both must reside (Bhadra et al., 2021; Gottschling et al., 2025). This is the one paradigm in Table 8 that is not represented in the corpus at all, and its practical obstacle is the same one that limits third-party auditing generally: computing the projection requires the forward operator, hence the raw k -space and the coil sensitivities, which commercial inline pipelines do not export (Section 5). Causal and counterfactual auditing, which asks the most direct question of all, whether a specific finding survives a controlled change of acquisition, appears in a small part of the corpus and almost never for physics-constrained k -space reconstruction (Pawlowski et al., 2020; Güngör et al., 2023; Peng et al., 2026).

The answer to RQ3 is therefore not that current evaluation is weak but that it is mismatched: each paradigm in use answers a question adjacent to the safety question, and no paradigm now in use establishes reconstruction safety on its own. The two that could, task-based observer assessment and null-space inspection, are respectively rare and absent, and the requirements in Section 4.7 follow from that pairing.

4.5. Structure of the evaluation gap

The counts above establish that safety-relevant evaluation is scarce. Analysis of the same per-study records establishes where the scarcity comes from, and indicates a division of labour rather than uniform neglect.

The benchmarks cannot show the failure the field is worried about. One hundred and nine of the 263 studies (41.4%) name a public dataset, and across the whole corpus those references resolve to eleven datasets: fastMRI, the Human Connectome Project, ADNI, IXI, UK Biobank, BraTS, OASIS, Calgary–Campinas, M4Raw, the developing Human Connectome Project and TCGA. Dataset labels are multi-valued (136 cohort assignments across the 109 studies), and of the 109, the cohorts named are healthy volunteers or population samples (55, 50.5%), chronic neurodegenerative cohorts (31, 28.4%), brain tumour cohorts (16, 14.7%), or raw- k -space benchmarks not curated for pathology (34, 31.2%). These four shares are multi-label and sum to 124.8%, a study naming

Table 8: What each evaluation paradigm can and cannot establish about prior-induced diagnostic error (RQ3). Rows 1–4 summarize evaluation practice and performance as reported in the included studies; quoted values are representative results from the cited primary studies, not pooled estimates. Rows 5–6 state analytical limits derived by the authors from the acquisition physics (our own analysis; Section 4) and are labelled as such rather than attributed to any primary study.

Evaluation Paradigm	Core Mechanism	What It Can and Cannot Establish	Key References
1. Pixel-wise fidelity metrics	L_2 / SSIM whole-image voxel averaging	Aggregate agreement with a reference only. Cannot resolve focal change: by (2), erasing a resolvable 100 mm^3 lacunar infarct at $r = 100$ costs 0.03 dB. Requires a registered reference	Recorded in 84 of the 144 evaluation-labelled studies
2. Perceptual & learned metrics	Networks trained on radiologist rankings	Tracks reader preference better than MSE/SSIM as an optimisation target ($75.2\% \pm 1.3\%$ and $79.2\% \pm 1.7\%$ against readers' own $61.5\% \pm 5.5\%$). Inherits training-reader variability, and was calibrated on synthetically corrupted data rather than clinically undersampled acquisitions, so transfer is assumed (Section 4.4)	Tang et al. (2025) ($n = 18$ in corpus)
3. Uncertainty estimation	Posterior sampling / Bayesian variance $p(X Y)$	Flags out-of-distribution inputs and regions of low confidence. Does not attribute missing structure to prior vs sampling operator; high confidence in hallucinated regions is possible	Khawaled and Freiman (2024) ($n = 18$ in corpus)
4. Task-based observer assessment	Human reader performance on a diagnostic task; model observers are its automated form, unrepresented here (Barrett et al., 1993; Barrett and Myers, 2004; International Commission on Radiation Units and Measurements, 1996)	Establishes diagnostic sufficiency directly; the gold standard in assessment theory. Costly and reader-variable ($61.5\% \pm 5.5\%$, falling to 25% under absolute Likert scoring), and rarely paired with a metric: only 18/263 report both	Radmanesh et al. (2022), Lang et al. (2024) ($n = 57$ in corpus)
5. Measurement-space residual check	$\ \mathbf{A}\tilde{X} - Y\ _2 / \ Y\ _2$ against tolerance	Cannot detect prior-induced edits: the residual is invariant to anything in $\mathcal{N}(\mathbf{A})$. A two-sided measurement-fidelity band, not a safety test	Common diagnostic check in the reviewed studies
6. Null-space / hallucination mapping	Structure written into the null-space component $(\mathbf{I} - \mathbf{A}^+ \mathbf{A})\hat{X}$	Isolates prior-supplied content, where erasure and hallucination must reside. Requires $\mathbf{A}$, hence raw k -space and coil sensitivities	Bhadra et al. (2021), Gottschling et al. (2025); formalization for reconstruction auditing is author-derived

688 two cohorts being counted in both; every other percentage here is single-label.
 689 *Within this corpus, no study benchmarks against a dataset curated for acute is-*
 690 *chemic stroke or intracranial haemorrhage.* The sharper version of this finding
 691 sits in the annotation layer the field has already built. fastMRI+ supplies 7,570
 692 expert bounding-box annotations across thirty brain pathology categories (Zhao
 693 et al., 2022), and not one of the thirty is a haemorrhage category of any kind,
 694 microhaemorrhage included. The apparent counterexample should be stated
 695 plainly rather than left for a reader to find, because it is the category closest
 696 to this article’s own worked example: fastMRI+ *does* carry an image-level *La-*
 697 *cunar Infarct* label (113 annotations across 32 subjects), alongside study-level
 698 *Small Vessel Chronic White Matter Ischemic Change* and a *Global Ischemia*
 699 label recorded for a single subject. It does not rescue the benchmark, and the
 700 reasons are the ones that matter for reconstruction safety. No fastMRI+ cat-
 701 egory distinguishes acute from chronic infarction, so a lacunar label carries no
 702 acuity and cannot identify the acute presentation whose erasure is at issue. The
 703 limit is in the source images as much as in the label set: the annotated se-
 704 ries are axial T2-FLAIR, T1w or contrast-enhanced T1w only, so neither DWI,
 705 on which infarct acuity is established, nor SWI or T_2^* -weighted imaging, on
 706 which microhaemorrhage is read, was annotated at all. Acute vascular pathol-
 707 ogy is therefore absent even where annotation infrastructure exists, and it is
 708 absent for a reason that adding labels alone would not repair. Supplementary
 709 Table S10 sets out the full derivation, including the sub-selection of 1,001 of
 710 the 5,847 fastMRI brain studies, the category listing, and the single-annotator
 711 limitation that travels with these annotations. Stroke-lesion datasets do exist,
 712 ISLES (Hernandez Petzsche et al., 2022) and ATLAS (Liew et al., 2022) among
 713 them, but they are segmentation benchmarks distributed as processed images
 714 without raw k -space, so they cannot serve reconstruction evaluation as released.
 715 This matters mechanically: lesion erasure cannot be measured on cohorts that
 716 contain no acute lesions, so a model can lead every public leaderboard in the
 717 field while remaining untested against the failure mode that would matter most
 718 in an emergency department. Where a pathology-labelled dataset is used at all

719 it is BraTS, whose tumours are large, high-contrast mass lesions, close to the
720 easiest possible case for lesion preservation and the least informative about the
721 lacunar and punctate findings at issue here. The benchmark infrastructure and
722 the metric therefore fail in the same direction.

723 *Rigour is divided between two literatures.* Of the 263 studies, 84 (31.9%)
724 release code but report no reader assessment, 44 (16.7%) report a reader assess-
725 ment but release no code, 122 (46.4%) do neither, and only 13 (4.9%) do both.
726 The division tracks publication venue. Applying the post-hoc classification rule
727 stated in Section 3 to all 263 venues, clinical imaging journals ($n = 131$) carry
728 reader assessment at nearly three times the rate of engineering and methods
729 venues ($n = 78$), which in turn release code at nearly twice the clinical rate;
730 general neuroscience and neuroimaging venues ($n = 37$) sit between. Supple-
731 mentary Table S12 gives the full cross-tabulation. The rule leaves 17 of the 263
732 journal titles unmatched, and those are reported as a fourth *unclassified* group
733 rather than forced into one of the three; the four groups between them account
734 for all 57 reader studies and all 97 code releases. The tendency for the two
735 practices to fall in different studies rather than the same ones gives a Fisher
736 exact odds ratio of 0.43 ($p = 0.013$). The estimate is insensitive to the one
737 labelling choice that could move it: taking the reader subset as the 61 studies
738 originally labelled at extraction rather than the 57 verified at appraisal gives an
739 odds ratio of 0.43 and $p = 0.010$. We nonetheless report the association descrip-
740 tively rather than as a hypothesis test, because the venue grouping it rests on is
741 a post-hoc keyword rule that was not part of the fixed protocol, and a p -value
742 computed on a partition chosen after seeing the data does not carry its nominal
743 meaning. The direction and magnitude are what we rely on. The magnitude
744 is the part that does not depend on the choice: computational verifiability and
745 diagnostic verifiability are largely held by different communities, only one study
746 in twenty holds both, and a safety argument requires both.

747 *Awareness of a failure mode does not produce the evaluation that would detect*
748 *it.* Generative studies record hallucination or fidelity instability in 31/80 (38.8%)
749 of cases, the highest share among the four largest method families, against




Figure 3: Reader assessment against hallucination awareness, by method family (recorded extraction labels; lower bounds). Bars are the share of each family’s studies that record a hallucination or instability failure mode (orange) and the share that carry any reader assessment (blue), using the verified reader subset ($n = 57$). Bar labels are rounded to the nearest integer; the one-decimal values quoted in the text are the same quantities (for example 15 against 15.1%, 39 against 38.8%). Generated from the released per-study file (`included_characteristics.csv`) by `5_Code/make_figures.py`; every value is recomputable from the deposit.

750 15.1% for compressed-sensing and parallel-imaging work and 12.3% for motion
751 correction. Yet at 9/80 (11.3%) they are also among the two least reader-
752 evaluated of those families, while the small self-supervised group records the
753 corpus’s highest hallucination share (8/17, 47.1%) and no reader assessment at
754 all (0/17); per-family shares are in Table 9 and Fig. 3. The sub-fields that most
755 clearly recognise that their models can synthesise structure are the least likely
756 to run the test capable of detecting it.

757 *Evaluation is not broadening as the field grows.* Reader-assessment share
758 fell from 32.4% (1995–2019) to 18.2% (2020–2022), recovering only partially to




Figure 4: Evaluation practice across publication eras. Grey bars give the number of included studies per era (34, 55, 174); lines give the share with any reader assessment and the share co-reporting a fidelity metric and a reader assessment on the same data, using the verified reader subset ($n = 57$). What recovery there is in 2023–2026 comes from subscription-walled clinical journals (52.0% reader assessment in that stratum against 15.4% in the open remainder; Section 5.2). Generated from the released per-study file by `5_Code/make_figures.py`; counts are lower bounds from recorded extraction labels.

20.7% (2023–2026) while the corpus grew fivefold (Table 9 and Fig. 4). What recovery there is comes entirely from the subscription-walled clinical literature: among 2023–2026 studies obtainable only through title-level library subscriptions, reader assessment runs at 52.0% (13 of 25), against 15.4% (23 of 149) in the openly accessible remainder of that era. Reader evidence, where it exists, sits disproportionately behind paywalls, and a review that cannot retrieve those texts will understate it; Section 5.2 takes the direction of that bias seriously rather than dismissing it. Among the 144 studies recording any evaluation paradigm, four in five record exactly one, and the share recording two or more fell across the same periods (Supplementary Table S12). Of the 84 studies recording a pixel-wise metric, 60 (71.4%) record no reader, uncertainty or perceptual measure beside it.

Readers are not displaced away from the sequences where the stakes are highest. One expectation is not borne out. The comparison must be made

Table 9: Evaluation practice across publication eras and method families (recorded labels; lower bounds). Reader-assessment share fell as the corpus grew and recovers only partially in the latest era, and that recovery is carried by subscription-walled clinical journals. Among the four largest method families the generative group, the one most associated with hallucination, is one of the two least reader-evaluated and is indistinguishable from physics-based unrolled work; the small self-supervised family ($n = 17$) records no reader assessment at all. Reader shares use the verified reader subset ($n = 57$; Section 3), so counts sum to 57 across the era rows and are multi-label across the family rows.

Era	n	Reader	Uncert.	Co-report
1995–2019	34	32.4% (11/34)	5.9%	5.9%
2020–2022	55	18.2% (10/55)	5.5%	7.3%
2023–2026	174	20.7% (36/174)	7.5%	6.9%
Method family (multi-label)				
Family	n	Reader share		
Compressed sensing / parallel imaging	93	25.8% (24/93)		
Motion correction	57	21.1% (12/57)		
Super-resolution	25	20.0% (5/25)		
Denoising	30	20.0% (6/30)		
Generative (GAN/diffusion/VAE)	80	11.3% (9/80)		
Physics-based / unrolled	37	10.8% (4/37)		
Self-supervised / unsupervised	17	0.0% (0/17)		

on disjoint strata to mean anything, because pulse sequence is a multi-label field and the naive contrast double-counts the 107 studies that acquire both an acute-pathology and a structural sequence (Supplementary Table S11 gives the correction). That the two disjoint strata below sum to $23 + 84 = 107$, the same number, is a third coincidence of unrelated tallies rather than a transcription error, verified against the released per-study file as the two flagged in Section efsec:requirements are. Restricted to studies acquiring one group and not the other, reader assessment runs at 5/23 (21.7%) for DWI or FLAIR without T1w or T2w against 10/84 (11.9%) for T1w or T2w without DWI or FLAIR — a difference in the direction one would hope for, resting on 23 studies in the smaller stratum and not approaching significance (Fisher exact odds ratio 2.06, $p = 0.31$). We report it as descriptive and draw only the negative conclusion it supports: there is no evidence that reader assessment is *displaced away* from the sequences on which acute pathology is read. The deficit is one of scarcity and of distribution across research communities, not of misdirection within the clinical literature.

789 4.6. Methodological quality and risk of bias

790 Across all 57 QUADAS-2 appraised diagnostic and reader-
 791 ological quality was highest for Flow and Timing (45/57 low risk of bias, 78.9%),
 792 reflecting the standard retrospective undersampling paradigm where reconstruc-
 793 tion and reference share the same acquisition timepoint without temporal degra-
 794 dation. Reference Standard and Index Test both sit at 26/57 low risk (45.6%),
 795 each now below half, with high ratings arising where radiologists scored acceler-
 796 ated images with reference ground truth visible or where lesion truth was defined
 797 by the comparator images themselves (incorporation bias). Patient Selection re-
 798 mains the weakest domain (15/57 low risk, 26.3%; 19/57 unclear, 33.3%; 23/57
 799 high risk, 40.4%), as cohorts were predominantly retrospective, single-centre,
 800 and pathology-enriched without explicit prospective enrolment rules. Applica-
 801 bility concerns follow a similar distribution, with low concerns in Patient Selec-
 802 tion for 36/57 (63.2%), Index Test for 30/57 (52.6%), and Reference Standard
 803 for 39/57 (68.4%).

804 The two rounds do not place the weaknesses in quite the same order, and
 805 the pooled shares above differ from the first round’s in every one of the seven
 806 domains. Low-risk shares fell for Index Test (53.8% to 45.6%), applicability of
 807 the Index Test (59.0% to 52.6%), Flow and Timing (82.1% to 78.9%), Refer-
 808 ence Standard (46.2% to 45.6%) and applicability of Patient Selection (64.1% to
 809 63.2%); they rose for Patient Selection risk of bias (23.1% to 26.3%) and applica-
 810 bility of the Reference Standard (59.0% to 68.4%). The two largest movements
 811 concern the index test and the applicability of the reference standard, and the
 812 first of these runs against the appraised studies. The second-round studies are
 813 therefore not a favourable addition to the appraisal, which is why the pooled
 814 figures are the ones reported here. The first-round shares quoted here already
 815 incorporate the five studies re-appraised on recovered full text.

816 Corpus-wide reporting transparency evaluated against STARD principles re-
 817 vealed substantial omissions across the broader 263 studies: 109/263 (41.4%)
 818 provided no identifiable reference standard (47 stating none at all, Table 3,
 819 and a further 62 naming a ground truth they left unspecified), scanner field

strength was unreported in 54/263 (20.5%), and 166/263 (63.1%) did not make
 code openly available. Overall, 220/263 studies (83.7%) omitted at least one
 key methodological indicator (field strength, reference standard, or open code).
 Two directions must be kept apart here, and conflating them has caused con-
 fusion elsewhere in this literature. With respect to the *reconstruction methods*,
 this reporting and selection pattern is anti-conservative: retrospective single-
 centre evaluation on curated data flatters reported performance relative to what
 prospective clinical deployment would show. With respect to *this review’s own*
conclusion, the same pattern is conservative: if published performance is flat-
 tered and failure is under-reported, then the scarcity of safety-relevant evalua-
 tion we document is an understatement, and the true gap is at least as large
 as the one measured. The safety argument turns on the second direction and
 is not weakened by the first. Detailed domain-level judgments and supporting
 quotations are available in the open project repository.

4.7. Requirements for safety-oriented evaluation

Read together, the synthesis implies five requirements that any evaluation
 intended to establish reconstruction safety must meet, each stated with the cor-
 pus evidence that motivates it and none met by more than a small minority
 of included studies. **R1, focal task-based measurement:** every global fi-
 delity score must be accompanied by a region-of-interest or task-based observer
 measure, because spatial averaging is the first mechanism by which lesion-scale
 error disappears (60/84, or 71.4%, of pixel-metric studies record nothing beside
 the metric; this equals the artifact share of the *R*-reporting subset in Table 6
 (45/63) to the decimal, which is a coincidence of two unrelated tallies rather than
 a transcription error, verified against the released per-study file). **R2, paired**
metric–reader reporting: metric validity claims require a reader comparison
 on the same data (6.8% of studies provide one). **R3, reference-independent**
validity: evaluation must remain interpretable when no registered artifact-free
 reference exists (26.2% rest solely on a fully sampled acquisition). **R4, robust-**
ness under acquisition shift as a primary outcome, reported as a result

rather than a caveat; robustness or distribution shift is recorded as a failure-mode concern in 76/263 studies (28.9%), almost always in the discussion rather than as a pre-specified endpoint. **R5, hallucination-specific detection**, operating on the image side, since measurement-space residual checks are provably blind to null-space edits. To these the analysis in Section 4.5 adds a precondition: benchmarks containing the pathology whose erasure is at issue, without which R1, R2 and R5 cannot be exercised at all.

5. Discussion

Three conclusions follow. First, the comparison on which the field’s central proxy rests is almost never made: no included study reports a rank correlation between a pixel-wise fidelity metric and radiologist diagnostic assessment, and only 18 of 263 record the two on the same data at all. Where readers do appear, what is measured is mostly agreement between readers, and it is weak and condition-dependent. The finding therefore rests more on the absence of measurement than on the measurements that exist. Second, the reason is structural, through the four mechanisms of Section 4: background voxels average focal error away, the loss functions are sign-symmetric where diagnosis is not, adversarial terms pay for invented texture, and every one of these metrics assumes a registered reference that motion destroys. The mechanisms are mathematical properties of the metrics themselves; how often each matters clinically is the part the thin reader evidence leaves less certain. Third, no evaluation paradigm now in use settles the safety question by itself, a conclusion resting on the paradigm-by-paradigm synthesis of Section 4.

The analysis of where the gap comes from, which was not part of the fixed protocol and is reported as exploratory throughout, changes what would follow from these findings if it holds. The shortfall does not look like a uniform oversight; it co-occurs with two structural features. Evaluation practice is divided between a community that releases code and a community that runs readers, with only 4.9% of studies doing both (an association we interpret descriptively

879 in Section 4.5, its venue partition being post-hoc), and the benchmark datasets
880 named in this corpus contain no acute pathology, so the failure mode of great-
881 est clinical consequence is not merely under-measured but unmeasurable on the
882 datasets in common use. If that association is causal, which this design cannot
883 establish, recommendations addressed to individual authors will not close the
884 gap; what would be required is benchmark construction and reporting conven-
885 tions that force the two halves of the evidence together.

886 This matters because accelerated reconstruction has left the research setting.
887 Deep-learning acceleration, denoising and image enhancement ship inside com-
888 mercial MRI products, several of them cleared as software as a medical device:
889 one such tool, an FDA-cleared vendor-agnostic DICOM-domain enhancement
890 product, is evaluated in this corpus by a prospective multicentre multireader
891 trial (Bash et al., 2021), which is also a reminder that a cleared product can
892 be validated on one sequence family and deployed across others. Published
893 guidance in both jurisdictions increasingly asks manufacturers to demonstrate
894 performance for the device’s stated intended use and under conditions repre-
895 sentative of deployment rather than on curated development data: the FDA,
896 Health Canada and MHRA guiding principles on good machine learning practice
897 ask that datasets be representative of the intended patient population and that
898 testing be independent of training (U.S. Food and Drug Administration et al.,
899 2021), the EU AI Act imposes accuracy and robustness obligations on high-risk
900 AI systems (European Parliament and Council of the European Union, 2024),
901 and the Medical Device Regulation requires clinical evidence appropriate to the
902 device’s intended purpose (European Parliament and Council of the European
903 Union, 2017). Our reading of that guidance is that aggregate pixel-wise fidelity
904 cannot support a claim about small-lesion detection, because (2) shows the ag-
905 gregate does not move when the small lesion is removed. That reading is ours:
906 no regulator has reviewed or endorsed it, and none of the three instruments
907 names an image-quality metric.

908 A practical obstacle compounds it, and it is the reason the requirements in
909 Section 4.7 are addressed to infrastructure rather than to authors. Commer-

910 cial inline pipelines export DICOM images after undisclosed normalisation, coil
 911 combination and vendor post-processing, and the raw complex k -space typi-
 912 cally never leaves the scanner. Every method in Table 8 that could detect a
 913 prior-induced edit needs something the exported DICOM does not carry: the
 914 null-space projection needs the forward operator and the coil sensitivities, a
 915 counterfactual audit needs the ability to re-run the acquisition under a con-
 916 trolled change, and even a measurement-space residual needs Y . Independent
 917 third-party auditing of a deployed reconstruction is therefore not merely difficult
 918 but structurally impossible without an open raw-data interface, and no amount
 919 of reporting discipline by authors of research papers can supply one.

920 Three things follow that are worth stating as claims rather than as recom-
 921 mendations.

922 *The cheapest useful change costs no new data.* A reconstruction paper that
 923 claims its method preserves small lesions can state, from (2) and its own brain
 924 volume and assumed local error ratio, the change in global PSNR that erasing
 925 the smallest lesion it claims to preserve would produce. For a lacunar infarct
 926 that is about 0.03 dB, which is smaller than the difference the paper is likely to
 927 be reporting between its method and its baseline. This is arithmetic, not a new
 928 experiment, and it converts an implicit claim into an explicit one at the point
 929 where the claim is made. Nothing in this review suggests such a line would be
 930 difficult; the finding is that nobody currently reports it.

931 *The two-community split is self-reinforcing, which is why exhortation has*
 932 *not worked.* Each community’s evaluation is close to the best it can do with
 933 the data it can obtain. Engineering venues evaluate on the datasets that carry
 934 raw k -space, and those contain healthy volunteers, chronic cohorts and tumours
 935 rather than acute pathology, so a reader study run on them would not test lesion
 936 preservation even if it were run. Clinical venues have the pathology, the readers
 937 and the acute presentations, and no access to raw k -space or to the vendor’s
 938 reconstruction, so a reproducibility claim is not available to them. The 4.9% of
 939 studies doing both are not more diligent than the rest; they are the studies that
 940 happened to sit where both kinds of access coincide. An instruction to individual

941 authors to run more reader studies asks most of them to do something the data
942 available to them makes uninformative.

943 *The missing object is a dataset, and its requirements are now specifiable.*

944 The public resources discussed in this review supply each necessary compo-
945 nent separately, and the distinction between what the corpus uses and what
946 merely exists is part of the point: raw multi-coil k -space is available in three
947 of the eleven datasets the included studies actually name (fastMRI, Calgary-
948 Campinas, M4Raw); acute cerebrovascular pathology with expert segmentation
949 exists in ISLES and ATLAS, which no included study uses because they are dis-
950 tributed as processed images; and reader-level pathology annotation on brain
951 MRI exists in fastMRI+, which no included study uses either. None combines
952 them, and the combination is what every requirement in Section 4.7 presup-
953 poses: raw k -space so the forward operator is known, acute lesions confirmed
954 against a standard independent of the images being evaluated, lesion-level an-
955 notation so a focal figure of merit can be computed, and a licence permitting
956 redistribution so the benchmark can be used by people who did not collect it.
957 That object does not exist, and until it does the safety question is not merely
958 unanswered but unaskable on public data.

959 5.1. Priorities for future research

960 Three of the five requirements in Section 4.7 translate directly into research
961 priorities, and we do not restate them here: paired metric-reader reporting
962 (R2), robustness under acquisition shift reported as a primary endpoint rather
963 than a caveat (R4), and benchmarks carrying the pathology whose erasure is at
964 issue, which is the precondition on which R1, R2 and R5 all depend. Calibrated
965 voxel-level uncertainty that tracks local reconstruction error (Khawaled and
966 Freiman, 2024) belongs with R4 and is the form in which R4 is most likely to
967 be actionable for generative and self-supervised architectures.

968 A fourth priority does not follow from anything this review extracted, and we
969 separate it for that reason. Many current pipelines operate slice by slice and in-
970 troduce inter-slice discontinuities on isotropic 3D acquisitions, so reconstruction

971 and auditing frameworks should model 3D anatomical coherence directly (Peng
972 et al., 2026) while meeting scanner-side latency constraints. Slice-wise versus
973 volumetric processing was not an extraction field in this review. We state the
974 priority on mechanistic grounds and on the cited work, not on a corpus count.

975 5.2. Threats to validity

976 The strongest conclusions rest on a small subset of reader-based studies, and
977 we state that prominently rather than in passing: 57 of 263 studies carry any
978 reader assessment and 18 pair it with a metric. Four threat classes bound what
979 follows from that base.

980 *How to read the limitations below.* They are not all bounded in the same
981 way, and treating them as one undifferentiated list of caveats would misrepresent
982 what this review knows about itself. Three kinds appear. Some are *bounded*
983 *by released evidence*: the non-retrieval analysis, the paywall stratification and
984 the appraisal-agreement decomposition are each quantified from files a reader
985 can recompute. Some are *bounded by argument*: the direction of publication
986 bias and of the reporting asymmetry is reasoned rather than measured, and we
987 say which way each points. One is simply *unbounded*, and it is the sharpest:
988 extraction has no reliability estimate.

989 Three properties of the extraction layer follow from that and should be
990 weighed together. First, its depth is deliberately shallow. Method, failure-
991 mode and evaluation categories were assigned from title, abstract and controlled
992 vocabulary, so the prevalence counts are floors on what studies *report at that*
993 *level* rather than measurements of what they did. Second, the shallowness is
994 visible in the counts themselves: only 84 of 263 studies are recorded as reporting
995 a pixel-wise fidelity metric, and 39 of the 57 verified reader studies are recorded
996 as reporting none, which for this literature is implausible as a description of
997 practice and entirely plausible as a description of abstracts. A further 119 of
998 263 studies (45.2%) carry no evaluation-paradigm label at all, which is why
999 we report both the /263 and the /144 denominators wherever the distinction
1000 matters. Third, the headline 18/263 mixes two depths: the reader label was

1001 verified against full text and corrected at appraisal, while the metric label was
1002 not. That asymmetry biases the intersection downward by an unknown factor,
1003 and it is the single respect in which our headline figure is least defensible as a
1004 point estimate — though the direction is conservative, since correcting it could
1005 only widen the measured gap.

1006 What we offer in place of a reliability coefficient is auditability, and the
1007 substitution is deliberate rather than a shortfall we failed to notice. Every
1008 extracted value is released with the verbatim quotation that supports it, so a
1009 reader can check whether an assignment is *correct*, which is a stronger and more
1010 useful guarantee than a coefficient showing that two people were *consistent*. It
1011 is not a substitute for reliability in the statistical sense, and we do not claim
1012 it is: consistency and correctness bound different failure modes, and only the
1013 first is estimable from our records. A duplicate-extraction exercise on a random
1014 subsample, reported with a κ , is the obvious next step and we name it as an
1015 outstanding task rather than a completed one.

1016 *Construct validity.* We searched with broad, high-sensitivity terms for
1017 architectures, artifacts and evaluation paradigms, and the corpus spans the
1018 reconstruction-safety construct, from pixel-wise metrics through perceptual
1019 losses to reader studies. What the corpus cannot guarantee is that a metric
1020 means the same thing wherever it appears: studies differ in what they treat as
1021 the quantity of interest, so a rank correlation with reader scores in one paper
1022 is not straightforwardly the same measurement as in another, and grouping
1023 such values under one heading is an interpretation we make, not a property of
1024 the data. The venue classification underlying Section 4.5 is likewise a post-hoc
1025 keyword assignment, released with the data rather than extracted from the
1026 studies.

1027 *Internal validity.* A protocol fixed before screening began, a documented
1028 search across seven databases, and independent screening and appraisal by three
1029 reviewers provide robust protection, and every extracted value is released with
1030 the verbatim snippet supporting it. Five specific threats bound these findings.
1031 First, 81 of the 1,149 candidate reports remain without obtainable text through

1032 structural licensing barriers. Screening the nearest obtainable stratum — the
 1033 341 reports recovered only through title-level subscriptions — found 86.2% ineli-
 1034 gible on the same criteria, so at that rate the 81 would contribute on the order of
 1035 11 studies rather than 81. That bounds the missing evidence without establish-
 1036 ing that the bias is nil. Non-retrieval is demonstrably *not* random with respect
 1037 to the variable this review measures: within the same stratum, 2023–2026 stud-
 1038 ies carry reader assessment at 52.0% against 15.4% in the openly accessible
 1039 literature of that era. Both findings point the same way. Reader evidence is
 1040 over-represented behind paywalls, so a review that cannot retrieve those texts
 1041 understates how much reader assessment the field performs, and our headline
 1042 scarcity figures are conservative bounds rather than point estimates. Second,
 1043 four included studies were assessed on abstract and controlled vocabulary alone
 1044 and one on partial text; a sensitivity analysis recomputing all headline metrics
 1045 on the 258 studies verified against complete full text leaves every finding direc-
 1046 tionally stable (e.g., metric–reader co-reporting remains 7.0% vs 6.8%). Third,
 1047 every QUADAS-2 tally in this article rests on the whole verified reader subset
 1048 ($n = 57, 399$ domain judgements), so none is computed on a partial base. The
 1049 second round moved all seven domain distributions, and the movements that
 1050 most concern the index test run against the appraised studies: low risk of bias
 1051 for the Index Test falls from 53.8% to 45.6%, and high applicability concern
 1052 for the Index Test rises from 33.3% to 40.4% (Section 4). Four studies remain
 1053 abstract-only and are appraised on that basis. Fourth, reliability is measured
 1054 for screening and for appraisal but not for extraction, and extraction is the
 1055 sharpest methodological limitation of the review. The screening and appraisal
 1056 coefficients, and the decomposition that must be read with the appraisal figure,
 1057 are given in Section 3 and not restated here; the consequence for this section
 1058 is that 126 of the 399 QUADAS-2 judgements were reached at a level of agree-
 1059 ment a single rater could not be relied on to reproduce, so the domain shares
 1060 in Table 3 carry that uncertainty. It does not propagate to the prevalence find-
 1061 ings, which rest on extraction labels rather than QUADAS-2 ratings, but it
 1062 bounds how firmly the risk-of-bias picture can be stated. Neither screening nor

appraisal agreement transfers to extraction. The duplicate-screening sample is
 also smaller than the 300–400 records a referee might reasonably ask for, and
 the interval reflects that. Fifth, four studies recorded at extraction as observer
 or reader designs were found at appraisal to contain no human reader and were
 reclassified as algorithmic, moving the reader subset from 61 to 57 and the al-
 gorithmic set from 202 to 206. The effect on the headline numbers is small and
 in the direction of a larger gap: metric–reader co-reporting moves from 7.6%
 to 6.8%, the 2023–2026 reader share from 23.0% to 20.7%, the code-and-reader
 intersection from 5.3% to 4.9%, and the Fisher exact p from 0.010 to 0.013 at
 an odds ratio that moves only from 0.43 to 0.43. Those four corrections are
 also the source of the 4/61 (6.6%) error rate quoted above, and the design field
 is the only one that received three-reviewer full-text verification: a comparable
 rate elsewhere would not overturn any finding here, since the headline shares are
 separated from their comparators by more than that margin, but it is a source
 of variability measured in one place and assumed in the rest.

External validity. Among the 209 studies stating a field strength the corpus
 covers 1.5, 3 and 7 Tesla (21, 153 and 35 respectively), together with several
 vendors’ coil configurations and the major pulse sequences; the remaining 54
 (20.5%) state none, and nothing about field-strength coverage should be read
 into them. Tasks, protocols, architectures, datasets and metrics otherwise vary
 substantially, which is why we report the direction and consistency of find-
 ings rather than their magnitudes. Many algorithmic studies evaluate on retro-
 spectively undersampled single-centre data, which flatters performance relative
 to prospective multi-centre use; where the public data were themselves pre-
 processed before undersampling, that inflation is larger again (Shimron et al.,
 2022). Prospective multicentre reader evidence is thin, and the examples this re-
 view leans on each carry a boundary that must travel with them: the emergency-
 department comparison is single-centre and its one discrepant finding is a single
 case (Lang et al., 2024); the widest-ranging acceleration reader study is retro-
 spective and single-centre (Radmanesh et al., 2022); and the largest prospective
 multicentre multireader trial evaluated 3D T1-weighted acquisitions only, with

an image-domain DICOM enhancement tool rather than reconstruction from undersampled k -space, reported improved rather than degraded lesion conspicuity, and has several authors affiliated with the vendor of the evaluated tool (Bash et al., 2021) — we cite it for what it does establish, that such tools are cleared and deployed. Finally, the benchmark concentration of Section 4.5 — eleven public datasets named across the corpus, not one curated for acute ischaemic stroke or intracranial haemorrhage — caps how far any leaderboard result can be carried into emergency imaging.

Conclusion validity. Because pixel-wise metrics agree only loosely with radiologist judgement (RQ1), a study resting on PSNR and SSIM alone supports a safety claim less well than one with readers or uncertainty-aware observers, and we weight the evidence accordingly. Publication bias favouring successful reconstructions is a material limitation and works in the same direction, making the reported prevalence of failure evidence an understatement: a literature with little incentive to report failure is exactly the corpus a review about failure must draw on, and the absence of a common effect measure rules out funnel-plot and small-study asymmetry tests. We applied no certainty framework such as GRADE, since the review produces no pooled effect estimates for such a framework to rate. Prevalence counts derive from recorded extraction labels and understate true practice, and no range we quote should be read as an interval estimate rather than as values from named studies.

6. Conclusion

We reviewed 263 primary studies on the safety, artifacts and fidelity of deep learning-based brain MRI reconstruction. Wherever metrics and readers have been compared, agreement between conventional fidelity scores and radiologist judgement is weak, and at most 18 of the 263 studies record the paired metric–reader evaluation capable of exposing that divergence. Every prevalence figure here is a floor derived from screening-level labels rather than a measurement of practice, and although screening agreement was high, the extraction labels

1123 those floors rest on carry no reliability estimate of their own. The weakness
1124 is structural rather than incidental: the four mechanisms identified under RQ2
1125 explain why, and the safety question is not one that any current evaluation
1126 paradigm answers alone.

1127 The contribution we would emphasise is the account of why the gap persists,
1128 and we offer it as a hypothesis these data are consistent with rather than as
1129 an established structural fact: the venue grouping it rests on is post-hoc, so
1130 although the association between code release and reader assessment is strong
1131 and would be conventionally significant on a pre-specified partition, we do not
1132 read it as a test (Section 4.5). What does not depend on that classification
1133 is the magnitude. Only 4.9% of studies both release code and run readers, so
1134 computational verifiability and diagnostic verifiability are, as recorded, held by
1135 different studies roughly nineteen times in twenty; and no dataset named in this
1136 corpus is curated for acute stroke or intracranial haemorrhage. A reconstruction
1137 model can therefore satisfy every metric and every public benchmark in current
1138 use while the question of whether it preserves a small acute lesion has never been
1139 asked. Closing that gap requires benchmarks that contain the pathology at issue
1140 and reporting conventions under which no pixel-wise metric appears alone, and
1141 until both exist, claims that accelerated reconstruction is diagnostically safe will
1142 continue to rest on measurements that cannot see the failures that matter.

1143 One change is available immediately and needs no new data. Equation (2)
1144 converts any lesion volume into the PSNR movement its erasure would cause, so
1145 an author reporting a reconstruction at a stated acceleration can state, alongside
1146 the global score, the smallest focal change that score is capable of registering
1147 at their image size. For a 100 mm³ infarct that figure is around 0.03 dB, which
1148 is below the precision at which PSNR is conventionally reported. Publishing
1149 that bound next to the metric costs one line, requires no reader and no new
1150 acquisition, and makes explicit what the metric is silent about. It is the cheap-
1151 est useful step towards evaluation that could detect the failure this review is
1152 concerned with.

1153 **Funding Statement**

1154 This research was supported by institutional research computing infrastruc-
1155 ture and scholarship funding from RMIT University. The funder had no role in
1156 the design of the review, the search strategy, the eligibility criteria, the screening
1157 or appraisal of studies, the synthesis or interpretation of findings, the writing
1158 of this report, or the decision to submit it for publication. No external, com-
1159 mercial, or industry funding was received, and no vendor of MRI hardware or
1160 reconstruction software contributed financial or non-financial support to this
1161 work.

1162 **Ethics statement**

1163 This review synthesises secondary data from published, de-identified study
1164 reports. It involved no human participants, no animal experimentation and
1165 no collection of private health data, so institutional review board approval and
1166 patient informed consent were not required.

1167 **Data and code availability**

1168 The review protocol, executed search strategy, screening criteria, every
1169 title/abstract and full-text decision with its recorded reason, the completed
1170 PRISMA 2020 checklist, the per-study extraction and QUADAS-2 worksheets
1171 for all 263 included studies, the analysis code, and the duplicate-screening
1172 sample and per-reviewer appraisal worksheet with their agreement computa-
1173 tions (`kappa_worksheet_214_HUMAN*.xlsx`, `kappa_human_results.csv` and
1174 `quadas2_three_reviewer_worksheet_18.csv`) are openly available on the
1175 Open Science Framework at <https://doi.org/10.17605/OSF.IO/EUXA8> and
1176 mirrored at <https://github.com/dmai287/brain-mri-recon-safety-slr>,
1177 archived at Zenodo under DOI 10.5281/zenodo.21845858 (release v1.0.2,
1178 commit 219242a; concept DOI 10.5281/zenodo.21836966). Every count
1179 reported here is regenerable from `included_characteristics.csv` in that re-
1180 lease, and the deposit encodes the same reader/algorithmic partition (57/206)

1181 that this article reports, with the four reclassified studies flagged individually.
1182 The two figures generated from that file and the Δ PSNR identity figure are
1183 produced by `5_Code/make_figures.py` and `5_Code/make_dpsnr_figure.py`.
1184 The deposit is retrospective and is documented as such. Full-text PDFs of the
1185 included studies are not redistributed, as publisher licence terms do not permit
1186 it; the DOI list in the released bibliography (`included_studies_263.bib`)
1187 identifies the corpus in full.

1188 **CRediT authorship contribution statement**

1189 **Dat Tat Mai:** Conceptualization, Methodology, Software, Investigation,
1190 Data curation, Formal analysis, Visualization, Project administration, Writing
1191 – original draft, Writing – review & editing. **Thai Viet Pham:** Investiga-
1192 tion, Data curation, Validation, Resources, Writing – review & editing. **Thu**
1193 **Nguyen Thi Dang:** Investigation, Data curation, Validation, Writing – re-
1194 view & editing. **James Jin Kang:** Supervision, Validation, Writing – review
1195 & editing.

1196 **Declaration of competing interest**

1197 The authors declare that they have no known competing financial interests or
1198 personal relationships that could have appeared to influence the work reported
1199 in this paper.

1200 **Declaration of generative AI and AI-assisted technologies in the** 1201 **manuscript preparation process**

1202 During the preparation of this work the authors used generative AI technolo-
1203 gies solely to assist with language editing and proofreading of the manuscript
1204 text. After using these tools, the authors reviewed and edited the content as
1205 needed and take full responsibility for the content of the published article. All
1206 study screening, data extraction, and quality appraisals reported in this review
1207 were conducted manually by the authors.

1208 **Appendix A. The included studies and how to trace a claim to them**

1209 All 263 included primary studies appear in the reference list, so the evidence
1210 base is part of the article of record rather than of the supplement alone. Because
1211 the synthesis is narrative and most claims are grouped rather than attached to
1212 a single study, the reference list alone does not show which studies produced a
1213 given number. That mapping is carried by `included_characteristics.csv`,
1214 which holds one row per included study with its extracted method family, evalu-
1215 ation paradigm, failure modes, sequences, field strength, reference standard and
1216 code availability, each with the verbatim quotation supporting it, so any count,
1217 share or crosstab in Section 4 can be reproduced and the contributing stud-
1218 ies listed by filtering that file. Supplementary Tables S1–S3 give the same 263
1219 studies as a numbered index by publication year, and `included_character-`
1220 `istics_supplement.csv` adds the per-study QUADAS-2 and reproducibility
1221 ratings.

1222 The reference list therefore contains two kinds of entry, and a reader needs
1223 a rule to tell them apart. It prints 299: the 263 included studies, carried into
1224 the list under PRISMA 2020 item 17, and 36 narrative-only works cited in ar-
1225 gument. **An entry appearing in the numbered index of Supplementary**
1226 **Tables S1–S3 is a member of the reviewed corpus; an entry that does**
1227 **not is a narrative citation.** Equivalently, and checkable without the supple-
1228 ment, every narrative citation is cited at a specific point in the running text
1229 of Sections 1–6, whereas corpus entries not discussed by name appear without
1230 an in-text citation. The two lists were generated by independent mechanisms
1231 and agree exactly as sets. We would be glad to move the corpus to a separately
1232 labelled supplementary reference list if the editor prefers that production route.

1233 **Appendix B. Supplementary material**

1234 Supplementary material associated with this article (Supplementary Ta-
1235 bles S1–S12 and Supplementary Figure S1) is provided as a separate document.
1236 Tables S1–S3 list all 263 included primary studies with full bibliographic details;

Tables S4–S9 give the comparison against related secondary studies, the search vocabulary and its derivation, the extraction form, the thematic evidence architecture, the four metric-failure mechanisms with their evidence sources, and the provenance of every released data file. Tables S10–S12 carry the supporting detail for Section 4.5: the annotation-layer derivation for fastMRI+ (S10), the disjoint-strata correction for the sequence contrast (S11), and the venue cross-tabulation with evaluation-paradigm multiplicity by era (S12).

References

- Aali, A., Arvinte, M., Kumar, S., Arefeen, Y.I., Tamir, J.I., 2025. Robust multi-coil MRI reconstruction via self-supervised denoising. *Magnetic resonance in medicine* doi:10.1002/mrm.30591. PMID: 40457510.
- Abed, M.A., Demir, S., Grotklaes, A., Germani, E., Faghihroohi, S., Sabir, H., Albarqouni, S., 2026. Mriqt: Physics-aware diffusion model for image quality transfer in neonatal ultra-low-field MRI. 2026 IEEE 23rd International Symposium on Biomedical Imaging (ISBI) doi:10.1109/isbi61048.2026.11515815.
- Aggarwal, H.K., Mani, M.P., Jacob, M., 2018. MoDL: Model-based deep learning architecture for inverse problems. *IEEE Transactions on Medical Imaging* 38, 394–405.
- Al-Masni, M.A., Lee, S., Yi, J., Kim, S., Gho, S.M., Choi, Y.H., Kim, D.H., 2022. Stacked u-nets with self-assisted priors towards robust correction of rigid motion artifact in brain MRI. *NeuroImage* doi:10.1016/j.neuroimage.2022.119411.
- Alçalar, Y.U., Akçakaya, M., 2025. Sparsity-driven parallel imaging consistency for improved self-supervised MRI reconstruction. 2025 IEEE International Conference on Image Processing (ICIP) doi:10.1109/icip55913.2025.11084355.
- Altmann, S., Grauham, N.F., Brockstedt, L., Kondova, M., Schmidtman, I., Paul, R., Clifford, B., Feiweier, T., Hosseini, Z., Uphaus, T., Groppa, S., Brockmann, M.A., Othman, A.E., 2024. Ultrafast brain MRI with deep learning reconstruction for suspected acute ischemic stroke. *Radiology* doi:10.1148/radiol.231938. PMID: 38376403.
- Altmann, S., Grauham, N.F., Mercado, M.A.A., Almansour, H., Paul, R.H., Schöffling, V.I., Ottenhausen, M., Brockmann, M.A., Othman, A.E., 2026. Ultra-fast sub-minute brain MRI with deep-learning-reconstruction for anesthesia-free emergency imaging in children. *Clinical Neuroradiology* doi:10.1007/s00062-026-01627-7.
- Angella, P., Balbi, L., Ferrando, F., Traverso, P., Varriale, R., Pastore, V.P., Santacesaria, M., 2025a. Dima: Diffusing motion artifacts for unsupervised correction in brain MRI images. *IEEE Access* doi:10.1109/access.2025.3634749.
- Angella, P., Pastore, V.P., Santacesaria, M., 2025b. Assessing the use of diffusion models for motion artifact correction in brain MRI. 2025 IEEE 22nd International Symposium on Biomedical Imaging (ISBI) doi:10.1109/isbi60581.2025.10981000.

1276 Antun, V., Renna, F., Poon, C., Adcock, B., Hansen, A.C., 2020. On instabilities
1277 of deep learning in image reconstruction and the potential costs of AI. *Proceed-*
1278 *ings of the National Academy of Sciences* 117, 30088–30095. doi:10.1073/pnas
1279 .1907377117.

1280 Barrett, H.H., Myers, K.J., 2004. *Foundations of Image Science*. Wiley, Hoboken, NJ.

1281 Barrett, H.H., Yao, J., Rolland, J.P., Myers, K.J., 1993. Model observers for as-
1282 sessment of image quality. *Proceedings of the National Academy of Sciences* 90,
1283 9758–9765.

1284 Bash, S., Wang, L., Airriess, C., Zaharchuk, G., Gong, E., Shankaranarayanan, A.,
1285 Tanenbaum, L.N., 2021. Deep learning enables 60% accelerated volumetric brain
1286 MRI while preserving quantitative performance: A prospective, multicenter, multi-
1287 reader trial. *AJNR. American journal of neuroradiology* doi:10.3174/ajnr.a7358.
1288 pMID: 34824098.

1289 Bathla, G., Rai, P., Chan, V.E.Y., Nickel, M.D., Kollasch, P., Dueker, K., Fagan, A.J.,
1290 Benson, J.C., Huston, J., Messina, S.A., 2025. Deep-learning accelerated vessel wall
1291 imaging using T1-SPACE at ultra-high-field strength MRI. *American Journal of*
1292 *Neuroradiology* 47, 1424–1431. doi:10.3174/ajnr.a9097.

1293 Bauman, G., Afshari, R., Bieri, O., 2025. Ultra-high-resolution brain MRI at 0.55t:
1294 bstar and its application to magnetization transfer ratio imaging. *Zeitschrift fur*
1295 *medizinische Physik* doi:10.1016/j.zemedi.2024.12.001. pMID: 39824686.

1296 Baz, A.M., Bendella, Z., Katemann, C., Sprinkart, A.M., Weiss, K., Weber, O.M.,
1297 Peeters, J.M., Lehen, N.C., Clauberg, R., Luetkens, J.A., Radbruch, A., Wicht-
1298 mann, B.D., 2026. Deep learning improves image quality in motion-robust and
1299 sedation-free pediatric brain MRI. *European radiology* doi:10.1007/s00330-026-
1300 12482-y. pMID: 41925834.

1301 Begnoche, J.P., Schilling, K.G., Boyd, B.D., Cai, L.Y., Taylor, W.D., Landman, B.A.,
1302 2022. Epi susceptibility correction introduces significant differences far from local
1303 areas of high distortion. *Magnetic resonance imaging* doi:10.1016/j.mri.2022.05.
1304 016. pMID: 35644448.

1305 Behrendt, F., Bhattacharya, D., Mielsing, R., Maack, L., Krüger, J., Opfer, R., Schlae-
1306 fer, A., 2025. Guided reconstruction with conditioned diffusion models for un-
1307 supervised anomaly detection in brain mris. *Computers in biology and medicine*
1308 doi:10.1016/j.combiomed.2025.109660.

1309 Beljaards, L., Nagtegaal, M., Rao, C., Dong, Y., Osch, M.J.P.v., Pezzotti, N., Doneva,
1310 M., Staring, M., 2026. Deep-disorder: Motion correction in 3D MRI via segment re-
1311 construction and registration. *NMR in biomedicine* doi:10.1002/nbm.70286. pMID:
1312 41958180.

1313 Beljaards, L., Pezzotti, N., Rao, C., Doneva, M., Osch, M.J.P.v., Staring, M., 2024.
1314 AI-based motion artifact severity estimation in undersampled MRI allowing for
1315 selection of appropriate reconstruction models. *Medical physics* doi:10.1002/mp
1316 .16918. pMID: 38167996.

Benjamin, A.J.V., Bano, W., Mair, G., Thompson, G., Casado, A., Perri, C.D., Davies, M., Marshall, I., 2020. Diagnostic quality assessment of ir-prepared 3D magnetic resonance neuroimaging accelerated using compressed sensing and k-space sampling order optimization. *Magnetic resonance imaging* doi:10.1016/j.mri.2020.08.025. PMID: 32890675.

Bhadra, S., Kelkar, V.A., Brooks, F.J., Anastasio, M.A., 2021. On hallucinations in tomographic image reconstruction. *IEEE Transactions on Medical Imaging* 40, 3249–3260. doi:10.1109/TMI.2021.3077857.

Bian, W., Jang, A., Liu, F., 2024. Improving quantitative MRI using self-supervised deep learning with model reinforcement: Demonstration for rapid t1 mapping. *Magnetic resonance in medicine* doi:10.1002/mrm.30045. PMID: 38342980.

Bjørkeli, E.B., Geitung, J.T., Esmaceli, M., 2025. Artificial intelligence-powered four-fold upscaling of human brain synthetic metabolite maps. *The Journal of international medical research* doi:10.1177/03000605251330578. PMID: 40257058.

Bollmann, S., Kristensen, M.H., Larsen, M.S., Olsen, M.V., Pedersen, M.J., Østergaard, L.R., O'Brien, K., Langkammer, C., Fazlollahi, A., Barth, M., 2019. Sharqnet - sophisticated harmonic artifact reduction in quantitative susceptibility mapping using a deep convolutional neural network. *Zeitschrift für medizinische Physik* doi:10.1016/j.zemedi.2019.01.001. PMID: 30773331.

Bossuyt, P.M., Reitsma, J.B., Bruns, D.E., Gatsonis, C.A., Glasziou, P.P., Irwig, L., Lijmer, J.G., Moher, D., Rennie, D., de Vet, H.C.W., 2015. STARD 2015: An updated list of essential items for reporting diagnostic accuracy studies. *BMJ* 351, h5527.

Brunátová, J., Löcke, M., Uribe, S., Bertoglio, C., 2025. Denoising of dual-venC pc-MRI with large high/low venC ratios. *Magnetic resonance in medicine* doi:10.1002/mrm.30278.

Byanju, R., Klein, S., Cristobal-Huerta, A., Hernandez-Tamames, J.A., Poot, D.H.J., 2025. Myelin water imaging from accelerated 3D-grase acquisitions using subspace constrained reconstruction. *Magma (New York, N.Y.)* doi:10.1007/s10334-025-01276-w. PMID: 40679773.

Campbell, M., McKenzie, J.E., Sowden, A., Katikireddi, S.V., Brennan, S.E., Ellis, S., Hartmann-Boyce, J., Ryan, R., Shepperd, S., Thomas, J., Welch, V., Thomson, H., 2020. Synthesis without meta-analysis (SWiM) in systematic reviews: reporting guideline. *BMJ* 368, l6890.

Cao, Y., Cao, J., Li, B., 2026. Loddpm: High-fidelity pseudo-healthy brain MRI synthesis via longitudinal differential guidance. *Biomedical Signal Processing and Control* 123, 110600. doi:10.1016/j.bspc.2026.110600.

Cao, Y., Levin, D., 1995. Using an image database to constrain the acquisition and reconstruction of MR images of the human head. *IEEE Transactions on Medical Imaging* doi:10.1109/42.387716.

Castro, D.C., Walker, I., Glocker, B., 2020. Causality matters in medical imaging. *Nature Communications* 11, 3673.

1359 Cerjanic, A.M., Diano, A.M., Johnson, C.L., 2025. Elastographic magnetization pre-
1360 pared imaging with rapid encoding. *Magnetic resonance in medicine* doi:10.1002/
1361 mrm.30482. pMID: 40065646.

1362 Chan, K.L., Ziegs, T., Henning, A., 2022. Improved signal-to-noise performance of
1363 multinet GRAPPA 1h FID MRSI reconstruction with semi-synthetic calibration
1364 data. *Magnetic Resonance in Medicine* 88, 1500–1515. doi:10.1002/mrm.29314.

1365 Chen, G., Xie, H., Rao, X., Liu, X., Otikovs, M., Frydman, L., Sun, P., Zhang, Z., Pan,
1366 F., Yang, L., Zhou, X., Liu, M., Bao, Q., Liu, C., 2025a. MRI motion correction
1367 through disentangled CycleGAN based on multi-mask k-space subsampling. *IEEE*
1368 *Transactions on Medical Imaging* 44, 1907–1921. doi:10.1109/tmi.2024.3523949.

1369 Chen, L., Lin, J., Fu, G., Yang, P., Liu, M., 2025b. Pw-GAN: Pseudo-warping field
1370 guided GAN for unsupervised denoising of fetal brain MRI images. *bioRxiv* doi:10.
1371 1101/2025.04.30.651598.

1372 Chen, X., Peng, Z., Valadez, G.H., 2023. Dd-cisenet: Dual-domain cross-iteration
1373 squeeze and excitation network for accelerated MRI reconstruction. *arXiv.org*
1374 doi:10.48550/arxiv.2305.00088.

1375 Choi, J.W., Cho, Y.J., Lee, S.B., Lee, S., Hwang, J.Y., Choi, Y.H., Cheon, J.E.,
1376 Lee, J., 2025a. Accelerated brain magnetic resonance imaging with deep learning
1377 reconstruction: a comparative study on image quality in pediatric neuroimaging.
1378 *Pediatric Radiology* 55, 1903–1914. doi:10.1007/s00247-025-06314-2.

1379 Choi, K.S., Park, C., Lee, J.Y., Lee, K.H., Jeon, Y.H., Hwang, I., Yoo, R.E., Yun, T.J.,
1380 Lee, M.J., Jung, K.H., Kang, K.M., 2025b. Prospective evaluation of accelerated
1381 brain MRI using deep learning-based reconstruction: Simultaneous application to
1382 2D spin-echo and 3D gradient-echo sequences. *Korean journal of radiology* doi:10.
1383 3348/kjr.2024.0653. pMID: 39780631.

1384 Chow, L.S., Rajagopal, H., 2017. Modified-BRISQUE as no reference image quality
1385 assessment for structural MR images. *Magnetic Resonance Imaging* 43, 74–87.
1386 doi:10.1016/j.mri.2017.07.016.

1387 Chung, K.J., Souza, R., Frayne, R., 2020. Restoration of lossy JPEG-compressed brain
1388 MR images using cross-domain neural networks. *IEEE Signal Processing Letters* 27,
1389 141–145. doi:10.1109/lsp.2019.2961072.

1390 Ciceri, T., Squarcina, L., Pigoni, A., Ferro, A., Montano, F., Bertoldo, A., Persico,
1391 N., Boito, S., Triulzi, F.M., Conte, G., Brambilla, P., Peruzzo, D., 2023. Geomet-
1392 ric reliability of super-resolution reconstructed images from clinical fetal MRI in
1393 the second trimester. *Neuroinformatics* doi:10.1007/s12021-023-09635-5. pMID:
1394 37284977.

1395 Conklin, J., Longo, M.G.F., Tabari, A., Filho, A.L.G., Liu, W., Splitthoff, D.N.,
1396 Lo, W.C., Cauley, S.F., Setsompop, K., Schaefer, P.W., Kirsch, J.E., Rapalino,
1397 O., Huang, S.Y., 2022. Clinical validation of wave-caipi susceptibility-weighted
1398 imaging for routine brain MRI at 1.5 t. *European radiology* doi:10.1007/s00330-
1399 022-08871-8. pMID: 35925387.

1400 Cordero-Grande, L., Ortuño-Fisac, J.E., del Hoyo, A.A., Uus, A., Deprez, M., Santos,
1401 A., Hajnal, J.V., Ledesma-Carbayo, M.J., 2023. Fetal MRI by robust deep gen-
1402 erative prior reconstruction and diffeomorphic registration. *IEEE Transactions on*
1403 *Medical Imaging* doi:10.1109/tmi.2022.3217725.

1404 Dabrowski, O., Falcone, J.L., Klauser, A., Songeon, J., Kocher, M., Chopard, B.,
1405 Lazeyras, F., Courvoisier, S., 2025. Sismik for brain MRI: Deep-learning-based
1406 motion estimation and model-based motion correction in k-space. *IEEE transactions*
1407 *on medical imaging* doi:10.1109/tmi.2024.3446450. PMID: 39159019.

1408 Delattre, B.M.A., Boudabbous, S., Hansen, C., Neroladaki, A., Hachulla, A.L., Vargas,
1409 M.I., 2019. Compressed sensing MRI of different organs: ready for clinical daily
1410 practice? *European Radiology* 30, 308–319. doi:10.1007/s00330-019-06319-0.

1411 Desalegn, L., Jifara, W., 2024. Hara-GAN: Hybrid attention and relative average
1412 discriminator based generative adversarial network for MR image reconstruction.
1413 *IEEE Access* doi:10.1109/access.2024.3364699.

1414 Dewey, B.E., Remedios, S.W., Sanjayan, M., Rjeily, N.B., Lee, A.Z., Wyche, C.,
1415 Duncan, S., Prince, J.L., Calabresi, P.A., Fitzgerald, K.C., Mowry, E.M., 2025.
1416 Super-resolution in clinically available spinal cord mris enables automated atrop-
1417 phy analysis. *AJNR. American journal of neuroradiology* doi:10.3174/ajnr.a8526.
1418 PMID: 39366765.

1419 Dieckmeyer, M., Roy, A.G., Senapati, J., Wachinger, C., Grundl, L., Döpfert, J.,
1420 Bertran, P.F., Lemke, A., Zimmer, C., Kirschke, J.S., Hedderich, D.M., 2021. Effect
1421 of MRI acquisition acceleration via compressed sensing and parallel imaging on brain
1422 volumetry. *Magma (New York, N.Y.)* doi:10.1007/s10334-020-00906-9. PMID:
1423 33502667.

1424 Ding, J., Duan, Y., Zhuo, Z., Yuan, Y., Zhang, G., Song, Q., Gao, B., Zhang, B.,
1425 Wang, M., Yang, L., Hou, Y., Yuan, J., Feng, C., Wang, J., Lin, L., Liu, Y., 2021.
1426 Acceleration of brain tof-mra with compressed sensitivity encoding: A multicenter
1427 clinical study. *AJNR. American journal of neuroradiology* doi:10.3174/ajnr.a7091.
1428 PMID: 33858820.

1429 Dolui, S., Wang, Z., Wolf, R.L., Nabavizadeh, A., Xie, L., Tosun, D., Nasrallah, I.M.,
1430 Wolk, D.A., Detre, J.A., Initiative, A.D.N., 2024. Automated quality evaluation
1431 index for arterial spin labeling derived cerebral blood flow maps. *Journal of magnetic*
1432 *resonance imaging : JMRI* doi:10.1002/jmri.29308. PMID: 38400805.

1433 Dong, Z., Reese, T.G., Lee, H.H., Huang, S.Y., Polimeni, J.R., Wald, L.L., Wang,
1434 F., 2025. Romer-epti: Rotating-view motion-robust super-resolution epti for snr-
1435 efficient distortion-free in-vivo mesoscale diffusion MRI and microstructure imaging.
1436 *Magnetic resonance in medicine* doi:10.1002/mrm.30365. PMID: 39552568.

1437 Du, J., Wang, L., Liu, Y., Zhou, Z., He, Z., Jia, Y., 2020. Brain MRI super-resolution
1438 using 3D dilated convolutional encoder–decoder network. *IEEE Access* doi:10.1109/
1439 *access*.2020.2968395.

1440 Duffy, B.A., Zhao, L., Sepehrband, F., Min, J., Wang, D.J., Shi, Y., Toga, A.W., Kim,
1441 H., Initiative, A.D.N., 2021. Retrospective motion artifact correction of structural
1442 MRI images using deep learning improves the quality of cortical surface reconstruc-
1443 tions. *NeuroImage* doi:10.1016/j.neuroimage.2021.117756. PMID: 33460797.

1444 Duong, S.T.M., Phung, S.L., Bouzerdoum, A., Ang, S.P., Schira, M.M., 2021. Correct-
1445 ing susceptibility artifacts of MRI sensors in brain scanning: A 3D anatomy-guided
1446 deep learning approach. *Sensors* (Basel, Switzerland) doi:10.3390/s21072314.
1447 pMID: 33810289.

1448 Eeden Risager, K., Gholamalizadeh, T., Mehdipour Ghazi, M., 2025. Non-reference
1449 quality assessment for medical imaging: Application to synthetic brain mris. *Deep*
1450 *Generative Models* doi:10.1007/978-3-031-72744-3_19.

1451 Eggl, M.F., D. Santis, S., 2026. Simulation-based inference at the theoretical limit
1452 for fast, robust microstructural MRI with minimal diffusion data. *Communications*
1453 *Medicine* doi:10.1038/s43856-026-01614-6.

1454 Elhabian, S., Vachet, C., Piven, J., Styner, M., Gerig, G., 2016. Compressive sensing
1455 based q-space resampling for handling fast bulk motion in hardi acquisitions. 2016
1456 IEEE 13th International Symposium on Biomedical Imaging (ISBI) doi:10.1109/is
1457 bi.2016.7493412.

1458 Elsaid, N.M.H., Wu, Y.C., 2019. Super-resolution diffusion tensor imaging using sr-
1459 cnn: A feasibility study. *Annual International Conference of the IEEE Engineer-*
1460 *ing in Medicine and Biology Society. IEEE Engineering in Medicine and Biology*
1461 *Society. Annual International Conference* doi:10.1109/embc.2019.8857125. pMID:
1462 31946482.

1463 Ernst, P., Chatterjee, S., Rose, G., Speck, O., Nürnberger, A., 2023. Sinogram upsam-
1464 pling using primal-dual unet for undersampled CT and radial MRI reconstruction.
1465 *Neural networks : the official journal of the International Neural Network Society*
1466 doi:10.1016/j.neunet.2023.08.004.

1467 Esteban, O., Birman, D., Schaer, M., Kiar, G., Gorgolewski, K.J., Poldrack, R.A.,
1468 2017. MRIQC: Advancing the automatic prediction of image quality in MRI from
1469 unseen sites. *PLoS ONE* 12, e0184661. doi:10.1371/journal.pone.0184661.

1470 European Parliament and Council of the European Union, 2017. Regulation (EU)
1471 2017/745 on medical devices (Medical Device Regulation). *Official Journal of the*
1472 *European Union*, L 117, 5 May 2017.

1473 European Parliament and Council of the European Union, 2024. Regulation (EU)
1474 2024/1689 laying down harmonised rules on artificial intelligence (Artificial Intelli-
1475 gence Act). *Official Journal of the European Union*, 12 July 2024.

1476 Fadeev, A., Eltonsy, N., Tourassi, G., Elmaghraby, A., 2005. Quality evaluation of
1477 fast morphing interpolation model for 3D volume reconstruction. *Proceedings of*
1478 *the Fifth IEEE International Symposium on Signal Processing and Information*
1479 *Technology*, 2005. doi:10.1109/isspit.2005.1577201.

1480 Fantini, I., Rittner, L., Yasuda, C., Lotufo, R., 2018. Automatic detection of mo-
1481 tion artifacts on MRI using deep CNN. 2018 International Workshop on Pattern
1482 Recognition in Neuroimaging (PRNI) doi:10.1109/prni.2018.8423948.

1483 Feizollah, S., Tardif, C.L., 2025. 3D mermaid: 3D multi-shot enhanced recovery motion
1484 artifact insensitive diffusion for submillimeter, multi-shell, and snr-efficient diffusion
1485 imaging. *Magnetic resonance in medicine* doi:10.1002/mrm.30436. pMID: 40035173.

1486 Fujita, S., Polak, D., Nickel, D., Splitthoff, D.N., Huang, Y., Gil, N., Buathong,
1487 S., Chiang, C.H., Lo, W.C., Clifford, B., Cauley, S.F., Conklin, J., Huang, S.Y.,
1488 2025. Motion-informed 3d deep learning reconstruction in patients with cognitive
1489 impairment. *American Journal of Neuroradiology* 47, 446–453. doi:10.3174/ajnr
1490 .a8977.

1491 Gao, Y., Zhu, X., Moffat, B.A., Glarin, R., Wilman, A.H., Pike, G.B., Crozier, S.,
1492 Liu, F., Sun, H., 2021. xqsm: quantitative susceptibility mapping with octave
1493 convolutional and noise-regularized neural networks. *NMR in biomedicine* doi:10.
1494 1002/nbm.4461. PMID: 33368705.

1495 Gardner, E.A., Ellis, J.H., Hyde, R.J., Aisen, A.M., Quint, D.J., Carson, P.L., 1995.
1496 Detection of degradation of magnetic resonance (MR) images: comparison of an
1497 automated MR image-quality analysis system with trained human observers. *Aca-*
1498 *demic radiology* doi:10.1016/s1076-6332(05)80184-9. PMID: 9419562.

1499 Garma-Oehmichen, A., Acuna-Luna, K.P., Santos-Diaz, A., 2021. Evaluation of echo
1500 planar imaging (EPI) distortion correction using synb0-disco and reversed phase
1501 encoding acquisition, in: 2021 43rd Annual International Conference of the IEEE
1502 Engineering in Medicine & Biology Society (EMBC), pp. 4023–4026. doi:10.1109/
1503 embc46164.2021.9630596.

1504 Giannakopoulos, I.I., Muckley, M.J., Kim, J., Breen, M., Johnson, P.M., Lui, Y.W.,
1505 Lattanzi, R., 2024. Accelerated MRI reconstructions via variational network and
1506 feature domain learning. *Scientific Reports* doi:10.1038/s41598-024-59705-0.

1507 Gil, N., Tabari, A., Lo, W.C., Clifford, B., Lang, M., Awan, K., Gaudet, K., Split-
1508 thoff, D.N., Polak, D., Cauley, S., Huang, S.Y., 2024. Quantitative evaluation of
1509 scout accelerated motion estimation and reduction (samer) mprage for morpho-
1510 metric analysis of brain tissue in patients undergoing evaluation for memory loss.
1511 *NeuroImage* doi:10.1016/j.neuroimage.2024.120865. PMID: 39349147.

1512 Glessgen, C., Gallichan, D., Moor, M., Hainc, N., Federau, C., 2019. Evaluation
1513 of 3D fat-navigator based retrospective motion correction in the clinical setting
1514 of patients with brain tumors. *Neuroradiology* doi:10.1007/s00234-019-02160-w.
1515 PMID: 30671618.

1516 Goffi, F., Bianchi, A.M., Schiena, G., Brambilla, P., Maggioni, E., 2025. Multi-metric
1517 approach for the comparison of denoising techniques for resting-state fMRI. *Human*
1518 *brain mapping* doi:10.1002/hbm.70080. PMID: 40309965.

1519 Gong, E., Pauly, J.M., Wintermark, M., Zaharchuk, G., 2018. Deep learning enables
1520 reduced gadolinium dose for contrast-enhanced brain MRI. *Journal of Magnetic*
1521 *Resonance Imaging* 48, 330–340. doi:10.1002/jmri.25970.

1522 Gottschling, N.M., Antun, V., Hansen, A.C., Adcock, B., 2025. The troublesome
1523 kernel: On hallucinations, no free lunches, and the accuracy–stability trade-off in
1524 inverse problems. *SIAM Review* 67, 73–104. doi:10.1137/23M1568739.

1525 Güngör, A., Dar, S.U., Öztürk, c., Korkmaz, Y., Bedel, H.A., Elmas, G., Ozbey,
1526 M., Çukur, T., 2023. Adaptive diffusion priors for accelerated MRI reconstruction.
1527 *Medical image analysis* doi:10.1016/j.media.2023.102872. PMID: 37384951.

1528 Gurbani, S.S., Schreibmann, E., Maudsley, A.A., Cordova, J.S., Soher, B.J., Poptani,
1529 H., Verma, G., Barker, P.B., Shim, H., Cooper, L.A.D., 2018. A convolutional neural
1530 network to filter artifacts in spectroscopic MRI. *Magnetic resonance in medicine*
1531 doi:10.1002/mrm.27166. PMID: 29520831.

1532 Hagiwara, A., Otsuka, Y., Hori, M., Tachibana, Y., Yokoyama, K., Fujita, S., Andica,
1533 C., Kamagata, K., Irie, R., Koshino, S., Maekawa, T., Chougar, L., Wada, A.,
1534 Takemura, M.Y., Hattori, N., Aoki, S., 2019. Improving the quality of synthetic
1535 FLAIR images with deep learning using a conditional generative adversarial network
1536 for pixel-by-pixel image translation. *AJNR. American journal of neuroradiology*
1537 doi:10.3174/ajnr.a5927. PMID: 30630834.

1538 Hales, P.W., Pfeuffer, J., Clark, C.A., 2020. Combined denoising and suppression
1539 of transient artifacts in arterial spin labeling MRI using deep learning. *Journal of*
1540 *magnetic resonance imaging : JMRI* doi:10.1002/jmri.27255.

1541 Hammernik, K., Klatzer, T., Kobler, E., Recht, M.P., Sodickson, D.K., Pock, T.,
1542 Knoll, F., 2018. Learning a variational network for reconstruction of accelerated
1543 MRI data. *Magnetic Resonance in Medicine* 79, 3055–3071.

1544 Hampe, N., Katscher, U., Berg, C.A.T.v.d., Tha, K.K., Mandija, S., 2020. Investigat-
1545 ing the challenges and generalizability of deep learning brain conductivity mapping.
1546 *Physics in medicine and biology* doi:10.1088/1361-6560/ab9356. PMID: 32408291.

1547 Harada, S., Takatsu, Y., Murayama, K., Sano, Y., Ikedo, M., 2025. Verification of reso-
1548 lution and imaging time for high-resolution deep learning reconstruction techniques.
1549 *Magnetic Resonance Imaging* 123, 110463. doi:10.1016/j.mri.2025.110463.

1550 Harper, J.R., Cherukuri, V., O'Reilly, T., Yu, M., Mbabazi-Kabachelor, E., Mu-
1551 lando, R., Sheth, K.N., Webb, A.G., Warf, B.C., Kulkarni, A.V., Monga, V.,
1552 Schiff, S.J., 2021. Assessing the utility of low resolution brain imaging: treatment
1553 of infant hydrocephalus. *NeuroImage. Clinical* doi:10.1016/j.nicl.2021.102896.
1554 PMID: 34911199.

1555 Haskell, M.W., Cauley, S.F., Bilgic, B., Hossbach, J., Splittthoff, D.N., Pfeuffer, J., Set-
1556 sompop, K., Wald, L.L., 2019. Network accelerated motion estimation and reduction
1557 (namer): Convolutional neural network guided retrospective motion correction using
1558 a separable motion model. *Magnetic resonance in medicine* doi:10.1002/mrm.27771.
1559 PMID: 31045278.

1560 Heo, H.Y., Singh, M., Yedavalli, V., Jiang, S., Zhou, J., 2024. Cest and nuclear over-
1561 hauser enhancement imaging with deep learning-extrapolated semisolid magnetiza-
1562 tion transfer reference: Scan-rescan reproducibility and reliability studies. *Magnetic*
1563 *resonance in medicine* doi:10.1002/mrm.29937. PMID: 38009996.

1564 Hernandez Petzsche, M.R., de la Rosa, E., Hanning, U., Wiest, R., Valenzuela, W.,
1565 Reyes, M., Meyer, M., Liew, S.L., Kofler, F., Ezhov, I., et al., 2022. ISLES 2022: A
1566 multi-center magnetic resonance imaging stroke lesion segmentation dataset. *Scien-
1567 tific Data* 9, 762. doi:10.1038/s41597-022-01875-5.

1568 Hewlett, M., Petrov, I., Johnson, P.M., Drangova, M., 2024. Deep-learning-based
1569 motion correction using multichannel MRI data: a study using simulated artifacts
1570 in the fastmri dataset. *NMR in biomedicine* doi:10.1002/nbm.5179.

1571 Ho, B.C., Kim, D., Kumar, A., Weiss, S., Vossler, H., Mormino, E., Zaharchuk, G.,
1572 Initiative, A.D.N., 2026. Evaluation of image-level harmonization methods for multi-
1573 center MR neuroimaging. *Journal of magnetic resonance imaging : JMRI* doi:10.
1574 1002/jmri.70221. PMID: 41489091.

1575 Hossbach, J., Splitthoff, D.N., Cauley, S., Clifford, B., Polak, D., Lo, W.C., Meyer,
1576 H., Maier, A., 2023. Deep learning-based motion quantification from k-space for
1577 fast model-based magnetic resonance imaging motion correction. *Medical physics*
1578 doi:10.1002/mp.16119. PMID: 36433748.

1579 Hosseini, M.S., Aghamiri, S.M.R., Ardekani, A.F., BagheriMofidi, S.M., Safari, M.,
1580 2025. Autocornn: An unsupervised physics-aware deep learning model for geo-
1581 metric distortion correction of brain MRI images towards MR-only stereotactic ra-
1582 diosurgery. *Journal of imaging informatics in medicine* doi:10.1007/s10278-024-
1583 01171-1. PMID: 39080159.

1584 Hu, Y., Xu, Y., Tian, Q., Chen, F., Shi, X., Moran, C.J., Daniel, B.L., Hargreaves,
1585 B.A., 2021. Run-up: Accelerated multishot diffusion-weighted MRI reconstruction
1586 using an unrolled network with u-net as priors. *Magnetic resonance in medicine*
1587 doi:10.1002/mrm.28446. PMID: 32783339.

1588 Hu, Y., Zhang, X., Chen, D., Yan, Z., Shen, X., Yan, G., Ou-Yang, L., Lin, J.,
1589 Dong, J., Qu, X., 2022. Spatiotemporal flexible sparse reconstruction for rapid
1590 dynamic contrast-enhanced MRI. *IEEE transactions on bio-medical engineering*
1591 doi:10.1109/tbme.2021.3091881. PMID: 34166181.

1592 Huang, J., Zhou, G., Yu, Z., Hu, W., 2022. [deep parallel MRI reconstruction based
1593 on a complex-valued loss function]. *Nan fang yi ke da xue xue bao = Journal of*
1594 *Southern Medical University* doi:10.12122/j.issn.1673-4254.2022.12.02. PMID:
1595 36651242.

1596 Huang, Y., Wang, Z., Zhang, X., Cao, J., Tu, Z., Lin, M., Li, L., Jiang, X., Guo, D.,
1597 Qu, X., 2025. Improve robustness to mismatched sampling rate: An alternating
1598 deep low-rank approach for exponential function reconstruction and its biomed-
1599 ical magnetic resonance applications. *Journal of Magnetic Resonance* 376, 107898.
1600 doi:10.1016/j.jmr.2025.107898.

1601 Huber, L.R., Stirnberg, R., Morgan, A.T., Feinberg, D.A., Ehse, P., Knudsen, L.,
1602 Gulban, O.F., Koiso, K., Gephart, I., Swegle, S., Wardle, S.G., Persichetti, A.S.,
1603 Beckett, A.J.S., Stöcker, T., Boulant, N., Poser, B.A., Bandettini, P.A., 2025. Short-
1604 term gradient imperfections in high-resolution epi lead to fuzzy ripple artifacts.
1605 *Magnetic resonance in medicine* doi:10.1002/mrm.30489. PMID: 40173320.

1606 Huijben, E.M.C., Amirrajab, S., Pluim, J.P.W., 2025. Enhancing reconstruction-based
1607 out-of-distribution detection in brain MRI with model and metric ensembles. *Com-
1608 puter methods and programs in biomedicine* doi:10.1016/j.cmpb.2025.109045.

1609 International Commission on Radiation Units and Measurements, 1996. *Medical Imag-
1610 ing — The Assessment of Image Quality*. Technical Report ICRU Report 54. ICRU.
1611 Bethesda, MD.

Islam, K.T., Dayarathna, S., Zhong, S., Zakavi, P., Kavnoudias, H., Farquharson, S., Durbridge, G., Sun, H., Bacchi, S., Egan, G.F., Barth, M., Dwyer, A., McMahon, K.L., Parizel, P.M., Law, M., Chen, Z., 2026. Synpoc: a high-quality generative diffusion model for transforming ultra-low-field point-of-care MRI using high-field MRI representations. *Scientific reports* doi:10.1038/s41598-025-33162-9. pMID: 41580445.

Ismail, A.R., Azhary, M.Z.R., Azwan, N.A.Z.N., Ismail, A., Alsaiani, N.A., 2024. Performance evaluation of medical image denoising using convolutional autoencoders, in: 2024 3rd International Conference on Creative Communication and Innovative Technology (ICCIT), pp. 1–6. doi:10.1109/iccit62134.2024.10701250.

Jayadeepthi, K., Yakshitha, K., Reddy, K.P.K., Gowtham, K., Mahesh, K., 2026. Deep learning based denoising framework for high-fidelity brain MRI reconstruction, in: 2026 4th International Conference on Inventive Computing and Informatics (ICICI), pp. 696–702. doi:10.1109/icici68773.2026.11580971.

Jayakumar, N., Hossain, T., Zhang, M., 2023. SADIR: Shape-aware diffusion models for 3D image reconstruction. *Shape in Medical Imaging* doi:10.1007/978-3-031-46914-5_23.

Jeong, S., 2026. Evaluating reliability of automated quantitative brain morphometry from fetal t2-weighted MRI. *Frontiers in Neuroscience* doi:10.3389/fnins.2026.1817732.

Ji, Y., Wu, W., Buck, M.H.S.d., Okell, T., Jezzard, P., 2023. Highly accelerated intracranial time-of-flight magnetic resonance angiography using wave-encoding. *Magnetic resonance in medicine* doi:10.1002/mrm.29647. pMID: 37010811.

Jiang, Y., Zhang, S., Chi, J., 2023. Multi-modal brain tumor data completion based on reconstruction consistency loss. *Journal of digital imaging* doi:10.1007/s10278-022-00697-6. pMID: 36856903.

Jochmann, T., Salman, F., Dwyer, M.G., Bergsland, N., Zivadinov, R., Haueisen, J., Schweser, F., 2025. Quantitative susceptibility mapping in magnetically inhomogeneous tissues. *Magnetic resonance in medicine* doi:10.1002/mrm.30537. pMID: 40312865.

Joshi, D., Gautam, A., Borotikar, B., 2025. Arterial spin labeling-based CBF map generation using a vector quantized variational autoencoder, in: 2025 IEEE International Conference on Electronics, Computing and Communication Technologies (CONECCT), pp. 1–5. doi:10.1109/conecct65861.2025.11306510.

Jun, Y., Arefeen, Y., Cho, J., Fujita, S., Wang, X., Grant, P.E., Gagoski, B., Jaimes, C., Gee, M.S., Bilgic, B., 2024. Zero-deepsub: Zero-shot deep subspace reconstruction for rapid multiparametric quantitative MRI using 3D-qalas. *Magnetic resonance in medicine* doi:10.1002/mrm.30018. pMID: 38282270.

Jung, W., Jeong, G., Kim, S., Hwang, I., Choi, S.H., Jeon, Y.H., Choi, K.S., Lee, J.Y., Yoo, R.E., Yun, T.J., Kang, K.M., 2025. Reliability of brain volume measures of accelerated 3D t1-weighted images with deep learning-based reconstruction. *Neuroradiology* doi:10.1007/s00234-024-03461-5. pMID: 39316090.

1654 Kamesh Iyer, S., Moon, B.F., Josselyn, N., Kurtz, R.M., Song, J.W., Ware, J.B.,
1655 Nabavizadeh, S.A., Witschey, W.R., 2022. Quantitative susceptibility mapping
1656 using plug-and-play alternating direction method of multipliers. *Scientific Reports*
1657 doi:10.1038/s41598-022-22778-w.

1658 Kanemaru, N., Takao, H., Amemiya, S., Abe, O., 2022. The effect of a post-scan
1659 processing denoising system on image quality and morphometric analysis. *Journal of*
1660 *neuroradiology = Journal de neuroradiologie* doi:10.1016/j.neurad.2021.11.007.

1661 Kang, S.H., Chung, J.Y., Lee, Y., Initiative, A.D.N., 2026. Simulation data-based dual
1662 domain network (sim-ddnet) for motion artifact reduction in MR images. *MAGNE-*
1663 *TOCHEMISTRY* doi:10.3390/magnetochemistry12010014.

1664 Kang, S.K., Shin, S.A., Seo, S., Byun, M.S., Lee, D.Y., Kim, Y.K., Lee, D.S., Lee, J.S.,
1665 2021. Deep learning-based 3D inpainting of brain MR images. *Scientific Reports*
1666 doi:10.1038/s41598-020-80930-w.

1667 Karimi, D., Jaimes, C., Machado-Rivas, F., Vasung, L., Khan, S., Warfield,
1668 S.K., Gholipour, A., 2021a. Deep learning-based parameter estimation in fe-
1669 tal diffusion-weighted MRI. *NeuroImage* doi:10.1016/j.neuroimage.2021.118482.
1670 pMID: 34455242.

1671 Karimi, D., Vasung, L., Machado-Rivas, F., Jaimes, C., Khan, S., Gholipour, A.,
1672 2021b. Accurate parameter estimation in fetal diffusion-weighted MRI, in: *MICCAI*
1673 2021, LNCS 12907. doi:10.1007/978-3-030-87234-2/46.

1674 Karthik, A., Aggarwal, K., Kapoor, A., Singh, D., Hu, L., Gandhamal, A., Kumar, D.,
1675 2024. Comprehensive assessment of imaging quality of artificial intelligence-assisted
1676 compressed sensing-based MR images in routine clinical settings. *BMC medical*
1677 *imaging* doi:10.1186/s12880-024-01463-6. pMID: 39434010.

1678 Karuna, Y., Syamala, N., Ravikumar, C.V., Thakur, P., Saladi, S., 2025. Modified
1679 energy-based GAN for intensity in homogeneity correction in brain MR images.
1680 *Scientific reports* doi:10.1038/s41598-025-08552-8. pMID: 40691452.

1681 Kayvanrad, M., Lin, A., Joshi, R., Chiu, J., Peters, T., 2016. Diagnostic quality
1682 assessment of compressed sensing accelerated magnetic resonance neuroimaging.
1683 *Journal of Magnetic Resonance Imaging* 44, 433–444. doi:10.1002/jmri.25149.

1684 Kazerouni, A., Aghdam, E.K., Heidari, M., Azad, R., Fayyaz, M., Hacıhaliloglu, I.,
1685 Merhof, D., 2023. Diffusion models in medical imaging: A comprehensive survey.
1686 *Medical Image Analysis* 88, 102846.

1687 Kharaji, M., Canton, G., Guo, Y., Mosi, M.H., Zhou, Z., Balu, N., Mossa-Basha, M.,
1688 2025. Dante-caipi accelerated contrast-enhanced 3D t1: Deep learning-based image
1689 quality improvement for vessel wall MRI. *AJNR. American journal of neuroradiol-*
1690 *ogy* doi:10.3174/ajnr.a8424. pMID: 39038956.

1691 Khawaled, S., Freiman, M., 2024. NPB-rec: A non-parametric bayesian deep-learning
1692 approach for undersampled MRI reconstruction with uncertainty estimation. *Artifi-*
1693 *cial intelligence in medicine* doi:10.1016/j.artmed.2024.102798. pMID: 38462289.

1694 Kiani Kalejahi, B., 2026. GAN-based cross-modality brain MRI synthesis: Paired
1695 versus unpaired training and comparison with diffusion and transformer models.
1696 Biomimetics doi:10.3390/biomimetics11030175.

1697 Kim, K.H., Park, S.H., 2017. Artificial neural network for suppression of banding
1698 artifacts in balanced steady-state free precession MRI. Magnetic Resonance Imaging
1699 37, 139–146. doi:10.1016/j.mri.2016.11.020.

1700 Kim, S., Park, S., 2021. Simultaneous multislice brain MRI t1 mapping with im-
1701 proved low-rank modeling. Tomography (Ann Arbor, Mich.) doi:10.3390/tomogr
1702 aphy7040047. pMID: 34698294.

1703 Kim, S.H., Choi, Y.H., Lee, J.S., Lee, S.B., Cho, Y.J., Lee, S.H., Shin, S.M., Cheon,
1704 J.E., 2022. Deep learning reconstruction in pediatric brain MRI: comparison of
1705 image quality with conventional T2-weighted MRI. Neuroradiology 65, 207–214.
1706 doi:10.1007/s00234-022-03053-1.

1707 Kim, T., Ji, S., Min, K., Kim, M., Youn, J., Oh, C., Kim, J., Lee, J., 2026. Vessel
1708 segmentation for χ -separation in quantitative susceptibility mapping. Magnetic
1709 resonance in medicine doi:10.1002/mrm.70054. pMID: 40891385.

1710 Knoll, F., Hammernik, K., Zhang, C., Moeller, S., Pock, T., Sodickson, D.K.,
1711 Akçakaya, M., 2020a. Deep-learning methods for parallel magnetic resonance imag-
1712 ing reconstruction: A survey of the current approaches, trends, and issues. IEEE
1713 Signal Processing Magazine 37, 128–140. doi:10.1109/MSP.2019.2950640.

1714 Knoll, F., Zbontar, J., Sriram, A., Muckley, M.J., Bruno, M., Defazio, A., Parente,
1715 M., Geras, K.J., Katsnelson, J., Chandarana, H., Zhang, Z., Drozdal, M., Romero,
1716 A., Rabbat, M., Vincent, P., Pinkerton, J., Wang, D., Yakubova, N., Owens, E.,
1717 Zitnick, C.L., Recht, M.P., Sodickson, D.K., Lui, Y.W., 2020b. fastmri: A publicly
1718 available raw k -space and DICOM dataset of knee images for accelerated MR image
1719 reconstruction using machine learning. Radiology: Artificial Intelligence 2, e190007.

1720 Koffler, A., Kerkerling, K.M., Göschel, L., Fillmer, A., Kolbitsch, C., 2024. Quanti-
1721 tative MR image reconstruction using parameter-specific dictionary learning with
1722 adaptive dictionary-size and sparsity-level choice. IEEE Transactions on Biomedical
1723 Engineering doi:10.1109/tbme.2023.3300090.

1724 Kulkarni, P.H., Merchant, S.N., Awate, S.P., 2021. Dictionary+wavelet model with
1725 nested-minorized vb-em for sms-caipi r-fMRI reconstruction. IEEE Open Journal
1726 of Signal Processing doi:10.1109/ojasp.2021.3100751.

1727 Kumar, V., Sharma, M.K., Jehadeesan, R., Venkatraman, B., Suman, G., Patra, A.,
1728 Goenka, A.H., Sheet, D., 2021. Learning to generate missing pulse sequence in MRI
1729 using deep convolution neural network trained with visual turing test, in: 2021 43rd
1730 Annual International Conference of the IEEE Engineering in Medicine & Biology
1731 Society (EMBC), pp. 3419–3422. doi:10.1109/embc46164.2021.9630435.

1732 Lacerda, L.M., Sperl, J.I., Menzel, M.I., Sprenger, T., Barker, G.J., Dell'Acqua, F.,
1733 2016. Diffusion in realistic biophysical systems can lead to aliasing effects in diffusion
1734 spectrum imaging. Magnetic resonance in medicine doi:10.1002/mrm.26080. pMID:
1735 26714794.

1736 Lan, H., Varghese, B.A., Sheikh-Bahaei, N., Sepehrband, F., Toga, A.W., Choupan, J.,
1737 Initiative, A.D.N., 2025. Diffusion based multi-domain neuroimaging harmonization
1738 method with preservation of anatomical details. *NeuroImage* doi:10.1016/j.neur
1739 oimage.2025.121297.

1740 Lang, M., Clifford, B., Lo, W.C., Applewhite, B.P., Tabari, A., Goncalves Filho,
1741 A.L.M., Hosseini, Z., Figueiro Longo, M.G., Cauley, S.F., Setsompop, K., Bilgic,
1742 B., Feiweier, T., Lev, M.H., Schaefer, P.W., Rapalino, O., Huang, S.Y., Conklin, J.,
1743 2024. Clinical evaluation of a 2-minute ultrafast brain MR protocol for evaluation of
1744 acute pathology in the emergency and inpatient settings. *AJNR. American Journal*
1745 *of Neuroradiology* 45, 379–385. doi:10.3174/ajnr.A8143.

1746 Lang, M., Tabari, A., Polak, D., Ford, J., Clifford, B., Lo, W., Manzoor, K., Split-
1747 thoff, D.N., Wald, L.L., Rapalino, O., Schaefer, P., Conklin, J., Cauley, S., Huang,
1748 S.Y., 2023. Clinical evaluation of scout accelerated motion estimation and reduc-
1749 tion technique for 3D MR imaging in the inpatient and emergency department set-
1750 tings. *AJNR. American journal of neuroradiology* doi:10.3174/ajnr.a7777. pMID:
1751 36702502.

1752 Lau, V., Xiao, L., Zhao, Y., Su, S., Ding, Y., Man, C., Wang, X., Tsang, A., Cao, P.,
1753 Lau, G.K.K., Leung, G.K.K., Leong, A.T.L., Wu, E.X., 2023. Pushing the limits of
1754 low-cost ultra-low-field MRI by dual-acquisition deep learning 3D superresolution.
1755 *Magnetic resonance in medicine* doi:10.1002/mrm.29642.

1756 Lee, J., Jung, W., Yang, S., Park, J.H., Hwang, I., Chung, J.W., Choi, S.H., Choi,
1757 K.S., 2024. Deep learning-based super-resolution and denoising algorithm improves
1758 reliability of dynamic contrast-enhanced MRI in diffuse glioma. *Scientific reports*
1759 doi:10.1038/s41598-024-76592-7. pMID: 39455814.

1760 Lee, J., Kim, H., Chung, H., Eom, H., Jang, J., Sohn, C.H., Choi, K.S., 2026a. Lesion-
1761 aware post-training of latent diffusion models for synthesizing diffusion MRI from
1762 CT perfusion, in: *Medical Image Computing and Computer Assisted Intervention*
1763 – MICCAI 2025, Springer. doi:10.1007/978-3-032-04937-7_27.

1764 Lee, N.G., Cui, S.X., Nayak, K.S., 2026b. Cartesian maxgirf: Model-based epi recon-
1765 struction incorporating gradient nonlinearity and concomitant field effects. *Magnetic*
1766 *resonance in medicine* doi:10.1002/mrm.70113. pMID: 41039908.

1767 Li, H., Chiew, M., Dragonu, I., Jezzard, P., Okell, T.W., 2026. Few-shot learning
1768 for highly accelerated 3D time-of-flight mra reconstruction. *Magnetic resonance in*
1769 *medicine* doi:10.1002/mrm.70072. pMID: 40931540.

1770 Li, M., Chen, C., Xiong, Z., Liu, Y., Rong, P., Shan, S., Liu, F., Sun, H.,
1771 Gao, Y., 2025a. Quantitative susceptibility mapping via deep neural networks
1772 with iterative reverse concatenations and recurrent modules. *Medical physics*
1773 doi:10.1002/mp.17747.

1774 Li, N., Wang, W.T., Pham, D.L., Butman, J.A., 2015. Artifactual microhemorrhage
1775 generated by susceptibility weighted image processing. *Journal of magnetic reso-*
1776 *nance imaging : JMRI* doi:10.1002/jmri.24728. pMID: 25143262.

1777 Li, Y., Liao, Y.P., Wang, J., Lu, W., Zhang, Y., 2025b. Patient-specific MRI super-
1778 resolution via implicit neural representations and knowledge transfer. *Physics in*
1779 *medicine and biology* doi:10.1088/1361-6560/adbed4. pMID: 40064110.

Li, Z., Fan, Q., Bilgic, B., Wang, G., Wu, W., Polimeni, J.R., Miller, K.L., Huang, S.Y., Tian, Q., 2023. Diffusion MRI data analysis assisted by deep learning synthesized anatomical images (deepanat). *Medical image analysis* doi:10.1016/j.media.2023.102744. PMID: 36867912.

Liang, D., Liu, Y., Wang, H., Zhu, Y., 2020. Deep learning for magnetic resonance image reconstruction: An overview. *IEEE Transactions on Radiation and Plasma Medical Sciences* 4, 641–655.

Liang, Z., Guo, X., Xu, W., Ibrahim, Y., Voets, N., Pretorius, P.M., Initiative, A.D.N., Noble, J.A., Kamnitsas, K., 2026. Itermask3d: Unsupervised anomaly detection and segmentation with test-time iterative mask refinement in 3D brain MRI. *Medical image analysis* doi:10.1016/j.media.2025.103763.

Liebrand, L.C., Karkalousos, D., Poirion, É., Emmer, B.J., Roosendaal, S.D., Marquering, H.A., Majoie, C.B.L.M., Savatovsky, J., Caan, M.W.A., 2025. Deep learning for efficient reconstruction of highly accelerated 3D FLAIR MRI in neurological deficits. *Magma (New York, N.Y.)* doi:10.1007/s10334-024-01200-8. PMID: 39212832.

Liew, S.L., Tavenner, B.P., Donnelly, M.R., Zavaliangos-Petropulu, A., Jeong, J.N., Barisano, G., Hutton, A., Simon, J.P., Juliano, J.M., Suri, A., et al., 2022. A large, curated, open-source stroke neuroimaging dataset to improve lesion segmentation algorithms. *Scientific Data* 9, 320. doi:10.1038/s41597-022-01401-7.

Liu, G., Ueguchi, T., Ogawa, S., 2026a. Block-interleaved segmented echo-planar imaging for improved activity detection in submillimeter high-resolution functional MRI at 7 t. *Magnetic resonance in medicine* doi:10.1002/mrm.30569. PMID: 40432253.

Liu, H., Chen, Z., Law, L.H., Liu, Y., Wang, Z., Wang, J., Zhang, Y., Shen, D., Huang, J., Chan, K.W.Y., 2026b. Accelerating cest MRI through complementary under-sampling and multi-offset transformer reconstruction. *Communications Engineering* doi:10.1038/s44172-025-00580-6.

Liu, H., Versteeg, E., Fuderer, M., Heide, O.v.d., Schilder, M.B., Berg, C.A.T.v.d., Sbrizzi, A., 2025a. Time-efficient, high-resolution 3t whole-brain relaxometry using cartesian 3D MR spin tomography in time-domain (MR-stat) with cerebrospinal fluid suppression. *Magnetic resonance in medicine* doi:10.1002/mrm.30384. PMID: 39607873.

Liu, J., Kocak, M., Supanich, M., Deng, J., 2020. Motion artifacts reduction in brain MRI by means of a deep residual network with densely connected multi-resolution blocks (DRN-DCMB). *Magnetic Resonance Imaging* 71, 69–79. doi:10.1016/j.mri.2020.05.002.

Liu, J., Lin, Q., Xiong, Z., Shan, S., Liu, C., Li, M., Liu, F., Pike, G.B., Sun, H., Gao, Y., 2026c. Highly undersampled MRI reconstruction via a single posterior sampling of diffusion models. *IEEE Transactions on Medical Imaging* doi:10.1109/tmi.2026.3654585.

Liu, Z., Patel, V., Zhou, X., Tao, S., Yu, T., Ma, J., Nickel, D., Liebig, P., Westerland, E.M., Mojahed, H., Gupta, V., Middlebrooks, E.H., 2025b. Deep learning reconstruction for 7t MP2RAGE and SPACE MRI: Improving image quality

1823 at high acceleration factors. *American Journal of Neuroradiology* 46, 2446–2454.
1824 doi:10.3174/ajnr.a8841.

1825 Liu, Z., Zhou, X., Tao, S., Ma, J., Nickel, D., Liebig, P., Mostapha, M., Patel, V.,
1826 Westerhold, E.M., Mojahed, H., Gupta, V., Middlebrooks, E.H., 2025c. Application
1827 of deep learning accelerated image reconstruction in t2-weighted turbo spin-echo
1828 imaging of the brain at 7t. *AJNR. American journal of neuroradiology* doi:10.
1829 3174/ajnr.a8662. pMID: 39832954.

1830 Lustig, M., Donoho, D., Pauly, J.M., 2007. Sparse MRI: The application of compressed
1831 sensing for rapid MR imaging. *Magnetic Resonance in Medicine* 58, 1182–1195.

1832 Ma, Q., Lai, Z., Wang, Z., Qiu, Y., Zhang, H., Qu, X., 2024. MRI reconstruction with
1833 enhanced self-similarity using graph convolutional network. *BMC medical imaging*
1834 doi:10.1186/s12880-024-01297-2. pMID: 38760778.

1835 Ma, X., Uğurbil, K., Wu, X., 2022. Mitigating transmit-b1 artifacts by predicting
1836 parallel transmission images with deep learning: A feasibility study using high-
1837 resolution whole-brain diffusion at 7 tesla. *Magnetic resonance in medicine* doi:10.
1838 1002/mrm.29238. pMID: 35403237.

1839 Mallardi, G., Calefato, F., Lanubile, F., Logroscino, G., Tafuri, B., 2025. Diffusion
1840 models for neuroimaging data augmentation: Assessing realism and clinical rele-
1841 vance. *Journal of Medical Systems* doi:10.1007/s10916-025-02300-1.

1842 Marchetto, E., Eichhorn, H., Gallichan, D., Schnabel, J.A., Ganz, M., 2025. Agree-
1843 ment of image quality metrics with radiological evaluation in the presence of mo-
1844 tion artifacts. *Magma (New York, N.Y.)* doi:10.1007/s10334-025-01266-y. pMID:
1845 40493331.

1846 Marchetto, E., Flassbeck, S., Mao, A., Assländer, J., 2026. Contrast-optimized ba-
1847 sis functions for self-navigated motion correction in quantitative MRI. *Magnetic*
1848 *resonance in medicine* doi:10.1002/mrm.70090. pMID: 40961378.

1849 Marie, H.S., Draz, M.M., Said, W., Elbaz, M., 2026. Bald uakari algorithm: a novel
1850 nature-inspired metaheuristic for missing pixel imputation in medical images. *Clus-*
1851 *ter Computing* doi:10.1007/s10586-026-06180-z.

1852 Mason, A., Rioux, J., Clarke, S.E., Costa, A., Schmidt, M., Keough, V., Huynh, T.,
1853 Beyea, S., 2020. Comparison of objective image quality metrics to expert radiolo-
1854 gists' scoring of diagnostic quality of MR images. *IEEE Transactions on Medical*
1855 *Imaging* 39, 1064–1072.

1856 Matsuo, K., Nakaura, T., Morita, K., Uetani, H., Nagayama, Y., Kidoh, M., Hoka-
1857 mura, M., Yamashita, Y., Shinoda, K., Ueda, M., Mukasa, A., Hirai, T., 2023.
1858 Feasibility study of super-resolution deep learning-based reconstruction using k-
1859 space data in brain diffusion-weighted images. *Neuroradiology* 65, 1619–1629.
1860 doi:10.1007/s00234-023-03212-y.

1861 McIlvain, G., 2022. Quantitative effects of off-resonance related distortion on brain
1862 mechanical property estimation with magnetic resonance elastography. *NMR in*
1863 *Biomedicine* doi:10.1002/nbm.4616.

1864 Mello, V.B.B., McKinley, R., Wiest, R., Rummel, C., 2026. Fast surface reconstruction
1865 of human brain MRI: benchmarking deep-learning based morphometry tools.
1866 Scientific reports doi:10.1038/s41598-026-55397-w. pMID: 42286078.

1867 Meng, Y., Allen, J.W., Sharghi, V.K., Qiu, D., 2025. Motion and temporal b0-shift cor-
1868 rections for qsm and r 2 * mapping using dual-echo spiral navigators and conjugate-
1869 phase reconstruction. Magnetic resonance in medicine doi:10.1002/mrm.30266.

1870 Meng, Z., Guo, R., Li, Y., Guan, Y., Wang, T., Zhao, Y., Sutton, B., Li, Y., Liang, Z.,
1871 2020. Accelerating t2mapping of the brain by integrating deep learning priors with
1872 low?rank and sparse modeling. Magnetic Resonance in Medicine 85, 1455–1467.
1873 doi:10.1002/mrm.28526.

1874 Mercier, C., Faisan, S., Pron, A., Girard, N., Auzias, G., Chonavel, T., Rousseau, F.,
1875 2025. Intersection-based slice motion estimation for fetal brain imaging. Computers
1876 in biology and medicine doi:10.1016/j.combiomed.2025.110005. pMID: 40112563.

1877 Miao, J., Huang, F., Narayan, S., Wilson, D.L., 2013. A new perceptual difference
1878 model for diagnostically relevant quantitative image quality evaluation: a prelimi-
1879 nary study. Magnetic resonance imaging doi:10.1016/j.mri.2012.09.009. pMID:
1880 23218792.

1881 Miao, J., Huang, F., Wilson, D.L., 2010. Comprehensive quantitative image quality
1882 evaluation of compressed sensing MRI reconstructions using a weighted perceptual
1883 difference model (case-PDM): selective evaluation, disturbance calibration, and ag-
1884 gregative evaluation of noise, blur, aliasing, and oil-painting artifacts, in: SPIE
1885 Proceedings, p. 762709. doi:10.1117/12.845507.

1886 Milovic, C., Tejos, C., Silva, J., Shmueli, K., Irarrazaval, P., 2024. XSIM : A structural
1887 similarity index measure optimized for MRI QSM. Magnetic Resonance in Medicine
1888 93, 411–421. doi:10.1002/mrm.30271.

1889 Mirza, M.U., Dalmaz, O., Bedel, H.A., Elmas, G., Korkmaz, Y., Gungor, A., Dar,
1890 S.U.H., Oguz, K.K., Çukur, T., 2026. Learning fourier-constrained diffusion bridges
1891 for MRI reconstruction. IEEE Transactions on Medical Imaging doi:10.1109/tmi.
1892 2026.3673956.

1893 Mohsen, M., Tamer, R., Hossam, H., Ahmed, O., Sherif, S., 2025. Deep super-
1894 resolution and age prediction for brain MRI using SRCNN: A lightweight frame-
1895 work for high-fidelity reconstruction and regression, in: 2025 Twelfth International
1896 Conference on Intelligent Computing and Information Systems (ICICIS), pp. 38–45.
1897 doi:10.1109/icicis66182.2025.11313176.

1898 Montalt-Tordera, J., Muthurangu, V., Hauptmann, A., Steeden, J.A., 2021. Machine
1899 learning in magnetic resonance imaging: Image reconstruction. Physica Medica 83,
1900 79–87.

1901 Moshe, Y.H., Buchsweiler, Y., Teicher, M., Artzi, M., 2024. Handling missing MRI
1902 data in brain tumors classification tasks: Usage of synthetic images vs. duplicate
1903 images and empty images. Journal of magnetic resonance imaging : JMRI doi:10.
1904 1002/jmri.29072.

1905 Muckley, M.J., Riemenschneider, B., Radmanesh, A., Kim, S., Jeong, G., Ko, J., Jun,
1906 Y., Shin, H., Hwang, D., Mostapha, M., Arberet, S., Nickel, D., Ramzi, Z., Ciuciu,
1907 P., Starck, J.L., Teuwen, J., Karkalousos, D., Zhang, C., Sriram, A., Huang, Z.,
1908 Yakubova, N., Lui, Y.W., Knoll, F., 2021. Results of the 2020 fastMRI challenge
1909 for machine learning MR image reconstruction. *IEEE Transactions on Medical*
1910 *Imaging* 40, 2306–2317. doi:10.1109/tmi.2021.3075856.

1911 Müller, S.J., Khadhraoui, E., Voit, D., Riedel, C.H., Frahm, J., Romero, J.M., Ernst,
1912 M., 2023. Comparison of epi DWI and steam DWI in early postoperative MRI con-
1913 trols after resection of tumors of the central nervous system. *Clinical Neuroradiology*
1914 doi:10.1007/s00062-023-01261-7.

1915 Muro, I., Isoiwa, T., Shibukawa, S., Usui, K., Otsuka, Y., 2025. Using deep learning
1916 to simultaneously reduce noise and motion artifacts in brain MR imaging. *Magnetic*
1917 *resonance in medical sciences : MRMS : an official journal of Japan Society of Mag-*
1918 *netic Resonance in Medicine* doi:10.2463/mrms.mp.2024-0098. pMID: 39938896.

1919 Muro, I., Shimizu, S., Tsukamoto, H., 2022. [improvement of motion artifacts in brain
1920 MRI using deep learning by simulation training data]. *Nihon Hoshasen Gijutsu*
1921 *Gakkai zasshi* doi:10.6009/jjrt.780108. pMID: 35046218.

1922 Nagaraj, U.D., Greer, J.S., Debnath, P., Tkach, J.A., Kline-Fath, B.M., 2025a. Slice-
1923 to-volume reconstruction of fetal brain MR imaging in clinical practice. *AJNR.*
1924 *American journal of neuroradiology* doi:10.3174/ajnr.a8728.

1925 Nagaraj, U.D., Meineke, J., Sriwastwa, A., Tkach, J.A., Leach, J.L., Doneva, M.,
1926 2025b. Evaluation of synthetic images derived from a neural network in pedi-
1927 atric brain magnetic resonance imaging. *Magnetic Resonance Imaging* 121, 110427.
1928 doi:10.1016/j.mri.2025.110427.

1929 Nimje, S., Artières, T., Guye, M., de Rochefort, L., 2025. Insights on scan-specific
1930 deep-learning strategies for brain MRI parallel imaging reconstruction. *NMR in*
1931 *Biomedicine* 38, e70079. doi:10.1002/nbm.70079. pMID: 40545734.

1932 Ning, L., Laun, F., Gur, Y., DiBella, E.V.R., Deslauriers-Gauthier, S., Megherbi,
1933 T., Ghosh, A., Zucchelli, M., Menegaz, G., Fick, R., St-Jean, S., Paquette, M.,
1934 Aranda, R., Descoteaux, M., Deriche, R., O'Donnell, L., Rath, Y., 2015. Sparse
1935 reconstruction challenge for diffusion MRI: Validation on a physical phantom to
1936 determine which acquisition scheme and analysis method to use? *Medical image*
1937 *analysis* doi:10.1016/j.media.2015.10.012. pMID: 26606457.

1938 Nishioka, N., Shimizu, Y., Kaneko, Y., Shirai, T., Suzuki, A., Amemiya, T., Ochi,
1939 H., Bito, Y., Takizawa, M., Ikebe, Y., Kameda, H., Harada, T., Fujima, N., Kudo,
1940 K., 2025. Accelerating FLAIR imaging via deep learning reconstruction: potential
1941 for evaluating white matter hyperintensities. *Japanese journal of radiology* doi:10.
1942 1007/s11604-024-01666-5. pMID: 39316286.

1943 Nonninger, J.N., Kienast, P., Pogledic, I., Mallouhi, A., Barkhof, F., Trattnig, S.,
1944 Robinson, S.D., Kasprian, G., Haider, L., 2025. Assessment of AI-accelerated t2-
1945 weighted brain MRI, based on clinical ratings and image quality evaluation. *Euro-*
1946 *pean journal of radiology* doi:10.1016/j.ejrad.2025.112123. pMID: 40315626.

Okell, T.W., Woods, J.G., Chiew, M., 2026. Combined angiography and perfusion using radial imaging and arterial spin labeling with structural contrast. *Magnetic resonance in medicine* doi:10.1002/mrm.70073. PMID: 40955094.

Olsson, H., Millward, J.M., Starke, L., Gladysz, T., Klein, T., Fehr, J., Lai, W.C., Lippert, C., Niendorf, T., Waiczies, S., 2024. Simulating rigid head motion artifacts on brain magnitude MRI data-outcome on image quality and segmentation of the cerebral cortex. *PLoS one* doi:10.1371/journal.pone.0301132. PMID: 38626138.

Ortiz-Gonzalez, A., Kobler, E., Schletter, L., Effland, A., 2026. Motiondps: Motion-compensated 3D brain MRI reconstruction. *IEEE Transactions on Medical Imaging* doi:10.1109/tmi.2026.3705847.

Osadebey, M.E., Pedersen, M., Arnold, D.L., Wendel-Mitoraj, K.E., 2018. Blind blur assessment of MRI images using parallel multiscale difference of gaussian filters. *Biomedical engineering online* doi:10.1186/s12938-018-0514-4. PMID: 29898715.

Page, M.J., McKenzie, J.E., Bossuyt, P.M., Boutron, I., Hoffmann, T.C., Mulrow, C.D., et al., 2021. The PRISMA 2020 statement: An updated guideline for reporting systematic reviews. *BMJ* 372, n71.

Park, S., Beckett, A., Häkkinen, S., Walker, E., Ma, S.J., Kim, S., Kim, H., Feinberg, D.A., 2025a. Higher spatial resolution and sensitivity in whole brain functional MRI at 7t using 3D epi accelerated with variable density caipi sampling and temporal random walk. *Magnetic resonance in medicine* doi:10.1002/mrm.30512.

Park, Y.W., Yoo, R.E., Shin, I., Jeon, Y.H., Singh, K.P., Lee, M.D., Kim, S., Yang, K., Jeong, G., Ryu, L., Han, K., Ahn, S.S., Lee, S.K., Jain, R., Choi, S.H., 2025b. A multinational study of deep learning-based image enhancement for multiparametric glioma MRI. *Scientific Reports* doi:10.1038/s41598-025-17993-0.

Pawar, K., Chen, Z., Shah, N.J., Egan, G.F., 2022. Suppressing motion artefacts in MRI using an inception-resnet network with motion simulation augmentation. *NMR in biomedicine* doi:10.1002/nbm.4225. PMID: 31865624.

Pawlowski, N., Castro, D.C., Glocker, B., 2020. Deep structural causal models for medical imaging. *Advances in Neural Information Processing Systems (NeurIPS)* 33, 15061–15073.

Peng, W., Xia, T., Ribeiro, F.D.S., Bosschieter, T., Adeli, E., Zhao, Q., Glocker, B., Pohl, K.M., 2026. Latent causal modeling for 3D brain MRI counterfactuals. *Deep Generative Models* doi:10.1007/978-3-032-05472-2_19.

Pirkil, C.M., Cencini, M., Kurzwski, J.W., Waldmannstetter, D., Li, H., Sekuboyina, A., Endt, S., Peretti, L., Donatelli, G., Pasquariello, R., Costagli, M., Buonincontri, G., Tosetti, M., Menzel, M.I., Menze, B.H., 2022. Learning residual motion correction for fast and robust 3D multiparametric MRI. *Medical image analysis* doi:10.1016/j.media.2022.102387. PMID: 35180675.

Pizarro, R., Assemlal, H.E., Jegathambal, S.K.B., Jubault, T., Antel, S., Arnold, D., Shmuel, A., 2023. Deep learning, data ramping, and uncertainty estimation for detecting artifacts in large, imbalanced databases of MRI images. *Medical image analysis* doi:10.1016/j.media.2023.102942. PMID: 37797482.

1989 Pruessmann, K.P., Weiger, M., Scheidegger, M.B., Boesiger, P., 1999. SENSE: Sensi-
1990 tivity encoding for fast MRI. *Magnetic Resonance in Medicine* 42, 952–962.

1991 Qiao, Y., Shi, Y., 2022. Unsupervised deep learning for fod-based susceptibility
1992 distortion correction in diffusion MRI. *IEEE transactions on medical imaging*
1993 doi:10.1109/tmi.2021.3134496. PMID: 34882551.

1994 Qiu, Z., Hu, S., Zhao, W., Sakaie, K., Sun, J.E.P., Griswold, M.A., Jones, D.K.,
1995 Ma, D., 2024. Self-calibrated subspace reconstruction for multidimensional MR
1996 fingerprinting for simultaneous relaxation and diffusion quantification. *Magnetic*
1997 *resonance in medicine* doi:10.1002/mrm.29969. PMID: 38102776.

1998 Radmanesh, A., Muckley, M.J., Murrell, T., Lindsey, E., Sriram, A., Knoll, F., Sod-
1999 ickson, D.K., Lui, Y.W., 2022. Exploring the acceleration limits of deep learning
2000 variational network-based two-dimensional brain MRI. *Radiology. Artificial intelli-*
2001 *gence* doi:10.1148/ryai.210313. PMID: 36523647.

2002 Ramos-Llordén, G., Lobos, R.A., Kim, T.H., Tian, Q., Witzel, T., Lee, H.H., Scholz,
2003 A., Keil, B., Yendiki, A., Bilgiç, B., Haldar, J.P., Huang, S.Y., 2023. High-fidelity,
2004 high-spatial-resolution diffusion magnetic resonance imaging of ex vivo whole human
2005 brain at ultra-high gradient strength with structured low-rank echo-planar imaging
2006 ghost correction. *NMR in biomedicine* doi:10.1002/nbm.4831. PMID: 36106429.

2007 Rassmann, S., Kugler, D., Ewert, C., Reuter, M., 2026. Regression is all you need for
2008 medical image translation. *IEEE transactions on medical imaging* doi:10.1109/tm
2009 i.2025.3650412. PMID: 41477795.

2010 Rauland, A., Meisler, S.L., Alexander-Bloch, A.F., Bagautdinova, J., Baller, E.B.,
2011 Gur, R.E., Gur, R.C., Luo, A.C., Moore, T.M., Popovych, O.V., Reetz, K., Roalf,
2012 D.R., Shinohara, R.T., Sotardi, S., Sydnor, V.J., Vossough, A., Eickhoff, S.B.,
2013 Cieslak, M., Satterthwaite, T.D., 2025. White matter bundle reconstruction from
2014 single-shell diffusion magnetic resonance imaging: Test-retest reliability and pre-
2015 dictive capability across orientation distribution function reconstruction methods.
2016 *Human brain mapping* doi:10.1002/hbm.70429. PMID: 41376618.

2017 Raymond, C., Yao, J., Clifford, B., Feiweier, T., Oshima, S., Telesca, D., Zhong, X.,
2018 Meyer, H., Everson, R.G., Salamon, N., Cloughesy, T.F., Ellingson, B.M., 2025.
2019 Leveraging physics-based synthetic MR images and deep transfer learning for arti-
2020 fact reduction in echo-planar imaging. *AJNR. American journal of neuroradiology*
2021 doi:10.3174/ajnr.a8566. PMID: 39947682.

2022 Rezvani, S., Moazami, S., Azevedo, C.J., Oberai, A.A., 2025. Probabilistic brain
2023 MR image transformation using generative models. *Scientific Reports* doi:10.1038/
2024 s41598-025-31958-3.

2025 Risk, B.B., Kociuba, M.C., Rowe, D.B., 2018. Impacts of simultaneous multislice
2026 acquisition on sensitivity and specificity in fMRI. *NeuroImage* doi:10.1016/j.neur
2027 oimage.2018.01.078. PMID: 29408461.

2028 Rizzuti, G., Schakel, T., Huttinga, N.R.F., Dankbaar, J.W., Leeuwen, T.v., Sbrizzi,
2029 A., 2024. Towards retrospective motion correction and reconstruction for clinical 3D
2030 brain MRI protocols with a reference contrast. *Magma (New York, N.Y.)* doi:10.
2031 1007/s10334-024-01161-y. PMID: 38758490.

2032 Ryu, K., Nam, Y., Gho, S.M., Jang, J., Lee, H.J., Cha, J., Baek, H.J., Park, J., Kim,
2033 D.H., 2019. Data-driven synthetic MRI FLAIR artifact correction via deep neural
2034 network. *Journal of magnetic resonance imaging : JMRI* doi:10.1002/jmri.26712.

2035 Saeed, S., Nickel, D., Weiland, E., Frohwein, L.J., Niehoff, J.H., Reiman, G., Schreck,
2036 J., Haag, N., Boriesosdick, J., Schoenbeck, D., Katz, M., Shahzadi, I., Woeltjen,
2037 M.M., Kroeger, J.R., Borggreffe, J., Moenninghoff, C., 2026. Improved image quality
2038 and reduced acquisition time in brain MRI using deep learning-based reconstruction:
2039 A quantitative and subjective assessment compared to standard mprage in 0.55 t
2040 MRI. *Magnetic resonance imaging* doi:10.1016/j.mri.2026.110656.

2041 Safari, M., Wang, S., Eidex, Z., Qiu, R.L.J., Chang, C.W., Yu, D.S., Yang, X., 2025a.
2042 A physics-informed deep learning model for MRI brain motion correction. *Medical*
2043 *physics* doi:10.1002/mp.70197. PMID: 41395855.

2044 Safari, M., Wang, S., Li, Q., Eidex, Z., Qiu, R.L.J., Chang, C.W., Mao, H., Yang, X.,
2045 2025b. Res-mocodiff: residual-guided diffusion models for motion artifact correction
2046 in brain MRI. *Physics in medicine and biology* doi:10.1088/1361-6560/ae1110.
2047 PMID: 41061726.

2048 Sagawa, H., Fushimi, Y., Nakajima, S., Fujimoto, K., Miyake, K.K., Numamoto,
2049 H., Koizumi, K., Nambu, M., Kataoka, H., Nakamoto, Y., Saga, T., 2021. Deep
2050 learning-based noise reduction for fast volume diffusion tensor imaging: Assessing
2051 the noise reduction effect and reliability of diffusion metrics. *Magnetic resonance*
2052 *in medical sciences : MRMS : an official journal of Japan Society of Magnetic*
2053 *Resonance in Medicine* doi:10.2463/mrms.tn.2020-0061. PMID: 32963184.

2054 Sakitis, C.J., Rowe, D.B., 2025. Bayesian merged utilization of GRAPPA and SENSE
2055 (bmugs) for in-plane accelerated reconstruction increases fMRI detection power.
2056 *Magnetic resonance imaging* doi:10.1016/j.mri.2024.110252.

2057 Sanchez, P., Kascenas, A., Liu, X., O'Neil, A.Q., Tsaftaris, S.A., 2022. What is
2058 healthy? generative counterfactual diffusion for lesion localization, in: *Deep Gener-*
2059 *ative Models (DGM4MICCAI), MICCAI Workshop, Springer.* pp. 34–44.

2060 Sanchez, T., Zalevskyi, V., Mihailov, A., Martí Juan, G., Eixarch, E., Jakab, A.,
2061 Dunet, V., Koob, M., Auzias, G., Bach Cuadra, M., 2026. Automatic quality
2062 control in multi-centric fetal brain MRI super-resolution reconstruction. *Perinatal,*
2063 *Preterm and Paediatric Image Analysis* doi:10.1007/978-3-032-05997-0_1.

2064 Sannanjanja, B., Zhu, C., Colip, C.G., Somasundaram, A., Ibrahim, M., Khrisat, T.,
2065 Mossa-Basha, M., 2022. Image-quality assessment of 3D intracranial vessel wall MRI
2066 using dante or dante-caipi for blood suppression and imaging acceleration. *AJNR.*
2067 *American journal of neuroradiology* doi:10.3174/ajnr.a7531. PMID: 35618420.

2068 Schauman, S., Niekerk, A.v., Norbeck, O., Rydén, H., Avventi, E., Skare, S., 2026. An
2069 exploration of motion-sampling interactions in 3D MRI for neuroimaging. *Magnetic*
2070 *resonance in medicine* doi:10.1002/mrm.70132. PMID: 41204061.

2071 Schlemper, J., Caballero, J., Hajnal, J.V., Price, A.N., Rueckert, D., 2017. A deep
2072 cascade of convolutional neural networks for dynamic MR image reconstruction.
2073 *IEEE Transactions on Medical Imaging* 37, 491–503.

2074 Shan, S., Gao, Y., Liu, P.Z.Y., Whelan, B., Sun, H., Dong, B., Liu, F., Wadding-
2075 ton, D.E.J., 2023. Distortion-corrected image reconstruction with deep learning
2076 on an MRI-linac. *Magnetic resonance in medicine* doi:10.1002/mrm.29684. pMID:
2077 37125656.

2078 Shan, S., Gao, Y., Waddington, D., Chen, H., Whelan, B., Liu, P., Wang, Y., Liu,
2079 C., Gan, H., Gao, M., Liu, F., 2024. Image reconstruction with b² inhomogeneity
2080 using a deep unrolled network on an open-bore MRI-linac. *IEEE Transactions on*
2081 *Instrumentation and Measurement* 73, 1–9. doi:10.1109/tim.2024.3481545.

2082 Sheng, H., Wang, X., Jiang, M., Zhang, Z., 2022. Deep learning-based diffusion-
2083 weighted magnetic resonance imaging in the diagnosis of ischemic penumbra in
2084 early cerebral infarction. *Contrast media & molecular imaging* doi:10.1155/2022/
2085 6270700. pMID: 35291425.

2086 Shimron, E., Tamir, J.I., Wang, K., Lustig, M., 2022. Implicit data crimes: Ma-
2087 chine learning bias arising from misuse of public data. *Proceedings of the National*
2088 *Academy of Sciences* 119, e2117203119. doi:10.1073/pnas.2117203119.

2089 Simkó, A., Ruiter, S., Löfstedt, T., Garpebring, A., Nyholm, T., Bylund, M., Jonsson,
2090 J., 2023. Improving MR image quality with a multi-task model, using convolutional
2091 losses. *BMC medical imaging* doi:10.1186/s12880-023-01109-z. pMID: 37784039.

2092 Singh, M., Mahmud, S.Z., Yedavalli, V., Zhou, J., Kamson, D.O., Zijl, P.v., Heo,
2093 H.Y., 2025. Learning-based motion artifact correction in the z-spectral domain
2094 for chemical exchange saturation transfer MRI. *Magnetic resonance in medicine*
2095 doi:10.1002/mrm.30440. pMID: 39834104.

2096 Solak, M., Tören, M., Asan, B., Kaba, E., Beyazal, M., Çeliker, F.B., 2026. Braingan
2097 framework for generating synthetic contrast-enhanced multisequence MRI in brain
2098 metastases. *Journal of Imaging Informatics in Medicine* doi:10.1007/s10278-026-
2099 02059-y.

2100 Sommer, K., Saalbach, A., Brosch, T., Hall, C., Cross, N.M., Andre, J.B., 2020.
2101 Correction of motion artifacts using a multiscale fully convolutional neural net-
2102 work. *AJNR. American journal of neuroradiology* doi:10.3174/ajnr.a6436. pMID:
2103 32054615.

2104 Song, W., Zeng, C., Zhang, X., Wang, Z., Huang, Y., Lin, J., Wei, W., Qu, X.,
2105 2022. Jointly estimating bias field and reconstructing uniform MRI image by deep
2106 learning. *Journal of Magnetic Resonance* 343, 107301. doi:10.1016/j.jmr.2022.
2107 107301.

2108 Song, Y., Sohl-Dickstein, J., Kingma, D.P., Kumar, A., Ermon, S., Poole, B., 2021.
2109 Score-based generative modeling through stochastic differential equations, in: *In-*
2110 *ternational Conference on Learning Representations (ICLR)*.

2111 Starck, S., Sideri-Lampretsa, V., Kainz, B., Menten, M.J., Mueller, T.T., Rueckert,
2112 D., 2026. Diff-def: Diffusion-generated deformation fields for conditional atlases.
2113 *IEEE transactions on medical imaging* doi:10.1109/tmi.2025.3595421.

2114 Sterne, J.A.C., Hernán, M.A., Reeves, B.C., Savović, J., Berkman, N.D., Viswanathan,
2115 M., Henry, D., White, I.R., Higgins, J.P.T., 2016. Robins-i: A tool for assessing
2116 risk of bias in non-randomised studies of interventions. *BMJ* 355, i4919.

2117 Sui, Y., Afacan, O., Jaimes, C., Gholipour, A., Warfield, S.K., 2022. Scan-specific gen-
2118 erative neural network for MRI super-resolution reconstruction. *IEEE Transactions*
2119 *on Medical Imaging* doi:10.1109/tmi.2022.3142610.

2120 Sujit, S.J., Coronado, I., Kamali, A., Narayana, P.A., Gabr, R.E., 2019. Automated
2121 image quality evaluation of structural brain MRI using an ensemble of deep learning
2122 networks. *Journal of magnetic resonance imaging : JMRI* doi:10.1002/jmri.26693.
2123 pMID: 30811739.

2124 Tabari, A., Lang, M., Awan, K., Liu, W., Clifford, B., Lo, W.C., Splittthoff, D.N.,
2125 Cauley, S., Rapalino, O., Schaefer, P., Huang, S.Y., Conklin, J., 2023. Optimized
2126 flow compensation for contrast-enhanced t1-weighted wave-caipi 3D mprage imaging
2127 of the brain. *European radiology experimental* doi:10.1186/s41747-023-00351-y.
2128 pMID: 37394534.

2129 Tabari, A., Vejdani-Jahromi, M., Lang, M., Nickel, D., Lo, W.C., Clifford, B., Con-
2130 klin, J., Huang, S.Y., 2026. Clinical validation of deep learning-accelerated versus
2131 wave-CAIPI postcontrast 3d T1-MPRAGE for evaluation of intracranial enhancing
2132 lesions. *American Journal of Neuroradiology* 47, 394–400. doi:10.3174/ajnr.a8992.

2133 Takada, S., Nakaura, T., Yoshida, N., Uetani, H., Shiraishi, K., Kobayashi, N., Mat-
2134 suo, K., Morita, K., Nagayama, Y., Kidoh, M., Yamashita, Y., Takayanagi, R.,
2135 Hirai, T., 2025. Impact of super-resolution deep learning-based reconstruction for
2136 hippocampal MRI: A volunteer and phantom study. *European Journal of Radiology*
2137 191, 112289. doi:10.1016/j.ejrad.2025.112289.

2138 Tang, C., Eisenmenger, L.B., Rivera-Rivera, L., Huo, E., Junn, J.C., Kuner, A.D.,
2139 Oechtering, T.H., Peret, A., Starekova, J., Johnson, K.M., 2025. Incorporating
2140 radiologist knowledge into MRI quality metrics for machine learning using rank-
2141 based ratings. *Journal of magnetic resonance imaging : JMRI* doi:10.1002/jmri
2142 .29672. pMID: 39690114.

2143 Tang, H., Hu, N., Yuan, Y., Xia, C., Liu, X., Zuo, P., Stalder, A.F., Schmidt, M., Zhou,
2144 X., Song, B., Sun, J., 2019. Accelerated time-of-flight magnetic resonance angiog-
2145 raphy with sparse undersampling and iterative reconstruction for the evaluation
2146 of intracranial arteries. *Korean journal of radiology* doi:10.3348/kjr.2017.0634.
2147 pMID: 30672166.

2148 Tanno, R., Worrall, D.E., Kaden, E., Ghosh, A., Grussu, F., Bizzi, A., Sotiropoulos,
2149 S.N., Criminisi, A., Alexander, D.C., 2021. Uncertainty modelling in deep learning
2150 for safer neuroimage enhancement: Demonstration in diffusion MRI. *NeuroImage*
2151 doi:10.1016/j.neuroimage.2020.117366.

2152 Taymourtash, A., Schwartz, E., Nenning, K.H., Licandro, R., Kienast, P., Hiele,
2153 V., Prayer, D., Kasprian, G., Langs, G., 2025. Measuring the effects of motion
2154 corruption in fetal fMRI. *Human brain mapping* doi:10.1002/hbm.26806. pMID:
2155 39846325.

2156 Tessema, A.W., Jin, S., Cho, H., Han, S., Cho, H., 2026. Technical feasibility of
2157 an efficient deep learning-based super-resolution method for 3D MR angiography.
2158 *IEEE Access* doi:10.1109/access.2026.3712275.

2159 Thaler, C., Sedlacik, J., Forkert, N.D., Stellmann, J.P., Schön, G., Fiehler, J., Gel-
2160 lißen, S., 2023. Effect of geometric distortion correction on thickness and volume
2161 measurements of cortical parcellations in 3D t1w gradient echo sequences. PloS one
2162 doi:10.1371/journal.pone.0284440. pMID: 37058493.

2163 Tisdall, M.D., Atkins, M.S., 2006. Using human and model performance to com-
2164 pare MRI reconstructions. IEEE transactions on medical imaging doi:10.1109/tm
2165 i.2006.881374. pMID: 17117780.

2166 Trivedi, V.K., Sharma, B., Balamurugan, P., 2024. Mcddpm: Multichannel conditional
2167 denoising diffusion model for unsupervised anomaly detection in brain MRI. 2024
2168 17th International Congress on Image and Signal Processing, BioMedical Engineer-
2169 ing and Informatics (CISP-BMEI) doi:10.1109/cisp-bmei64163.2024.10906217.

2170 Uher, D., Drenthen, G.S., Poser, B.A., Hofman, P.A.M., Wagner, L.G., Lanen,
2171 R.H.G.J.v., Hoeberigs, C.M., Colon, A.J., Schijns, O.E.M.G., Jansen, J.F.A.,
2172 Backes, W.H., 2024. Deepflair: A neural network approach to mitigate signal and
2173 contrast loss in temporal lobes at 7 tesla FLAIR images. Magnetic resonance imag-
2174 ing doi:10.1016/j.mri.2024.04.013. pMID: 38621552.

2175 U.S. Food and Drug Administration, Health Canada, Medicines and Healthcare
2176 products Regulatory Agency, 2021. Good machine learning practice for medical
2177 device development: Guiding principles. Joint guiding principles, October 2021.
2178 URL: [https://www.fda.gov/medical-devices/software-medical-device-samd](https://www.fda.gov/medical-devices/software-medical-device-samd/good-machine-learning-practice-medical-device-development-guiding-principles)
2179 [/good-machine-learning-practice-medical-device-development-guiding-pri](https://www.fda.gov/medical-devices/software-medical-device-samd/good-machine-learning-practice-medical-device-development-guiding-principles)
2180 [nciples](https://www.fda.gov/medical-devices/software-medical-device-samd/good-machine-learning-practice-medical-device-development-guiding-principles).

2181 Usui, K., Muro, I., Shibukawa, S., Goto, M., Ogawa, K., Sakano, Y., Kyogoku, S.,
2182 Daida, H., 2023. Evaluation of motion artefact reduction depending on the artefacts'
2183 directions in head MRI using conditional generative adversarial networks. Scientific
2184 reports doi:10.1038/s41598-023-35794-1. pMID: 37237139.

2185 Vakli, P., Weiss, B., Szalma, J., Barsi, P., Gyuricza, I., Kemenczky, P., Somogyi, E.,
2186 Ádám Nárai, Gál, V., Hermann, P., Vidnyánszky, Z., 2023. Automatic brain MRI
2187 motion artifact detection based on end-to-end deep learning is similarly effective
2188 as traditional machine learning trained on image quality metrics. Medical image
2189 analysis doi:10.1016/j.media.2023.102850.

2190 Vanderhasselt, T., Naeyaert, M., Watté, N., Allemeersch, G., Raeymaeckers, S.,
2191 Dudink, J., Mey, J.d., Raeymaekers, H., 2020. Synthetic MRI of preterm in-
2192 fants at term-equivalent age: Evaluation of diagnostic image quality and auto-
2193 mated brain volume segmentation. AJNR. American journal of neuroradiology
2194 doi:10.3174/ajnr.a6533. pMID: 32299803.

2195 Verclytte, S., Beaugrard, G., Nickel, D., Muñoz-Ramírez, V., Norberciak, L., Bois-
2196 sel, M., Chaton, V., Kwiatkowski, A., 2026. Deep learning-accelerated 3d FLAIR
2197 enables reliable MS lesion detection. American Journal of Neuroradiology 47, 1560–
2198 1568. doi:10.3174/ajnr.a9140.

2199 Verdicchio, M., Grassi, F., Esposito, G., Cavaliere, C., Alfano, V., Cerrone, M., Pap-
2200 patà, S., Salvatore, E., Cecca, A.D., Salvatore, M., Aiello, M., 2025. Reliability

assessment of accelerated MR sequences for the estimation of structural and vascular markers in dementia. Scientific reports doi:10.1038/s41598-025-87224-z. pMID: 40240759.

Vranic, J.E., Cross, N.M., Wang, Y., Hippe, D.S., Weerd, E.d., Mossa-Basha, M., 2019. Compressed sensing-sensitivity encoding (cs-SENSE) accelerated brain imaging: Reduced scan time without reduced image quality. AJNR. American journal of neuroradiology doi:10.3174/ajnr.a5905. pMID: 30523142.

Wada, A., Akashi, T., Hagiwara, A., Nishizawa, M., Shimoji, K., Kikuta, J., Maekawa, T., Sano, K., Kamagata, K., Nakanishi, A., Aoki, S., 2024. Deep learning-driven transformation: A novel approach for mitigating batch effects in diffusion MRI beyond traditional harmonization. Journal of magnetic resonance imaging : JMRI doi:10.1002/jmri.29088.

Wang, C.J., Rost, N.S., Golland, P., 2023. Spatial-intensity transforms for medical image-to-image translation. IEEE transactions on medical imaging doi:10.1109/tmi.2023.3283948. pMID: 37285247.

Wang, G., Gong, E., Banerjee, S., Martin, D., Tong, E., Choi, J., Chen, H., Wintermark, M., Pauly, J.M., Zaharchuk, G., 2020. Synthesize high-quality multi-contrast magnetic resonance imaging from multi-echo acquisition using multi-task deep generative model. IEEE Transactions on Medical Imaging 39, 3089–3099. doi:10.1109/tmi.2020.2987026.

Wang, N., Liao, C., Cao, X., Nishimura, M., Brackner, Y.W.E., Yurt, M., Gao, M., Abraham, D., Alkan, C., Iyer, S.S., Zhou, Z., Jeong, H., Kerr, A., Haldar, J.P., Setsompop, K., 2025a. Spherical echo-planar time-resolved imaging (sepi) for rapid 3D quantitative t_2^* and susceptibility imaging. Magnetic resonance in medicine doi:10.1002/mrm.30255. pMID: 39250435.

Wang, X., Cabezas, M., Tang, Z., D'Souza, A., Liu, D., Wang, C., Cai, W., 2025b. Estimating uncertainty on deep learning-driven fibre orientation distribution enhancement, in: ICONIP 2024, CCIS 2285. doi:10.1007/978-981-96-6957-8/25.

Wang, X., Fan, H., Tan, Z., Vasylechko, S., Yang, E., Didier, R., Afacan, O., Uecker, M., Warfield, S.K., Gholipour, A., 2025c. Rapid, high-resolution and distortion-free R_2^* mapping of fetal brain using multi-echo radial flash and model-based reconstruction. Magnetic resonance in medicine doi:10.1002/mrm.30604. pMID: 40534142.

Wang, Z., Bovik, A.C., Sheikh, H.R., Simoncelli, E.P., 2004. Image quality assessment: from error visibility to structural similarity. IEEE Transactions on Image Processing 13, 600–612. doi:10.1109/TIP.2003.819861.

Wang, Z., Ramasawmy, R., Javed, A., Mugler, J.P., Meyer, C.H., Campbell-Washburn, A.E., 2025d. Variable-flip-angle 3D spiral-in-out turbo spin-echo imaging using concomitant gradient compensation and echo reordering at 0.55 t. Magnetic resonance in medicine doi:10.1002/mrm.30380. pMID: 39567364.

Wen, Y., Ma, H., Xiang, S., Feng, Z., Guan, C., Li, X., 2025. Accelerating brain t_2 -weighted imaging using artificial intelligence-assisted compressed sensing combined with deep learning-based reconstruction: a feasibility study at 5.0t MRI. BMC medical imaging doi:10.1186/s12880-025-01763-5. pMID: 40596927.

2244 Whiting, P.F., Rutjes, A.W.S., Westwood, M.E., Mallett, S., Deeks, J.J., Reitsma,
2245 J.B., Leeflang, M.M.G., Sterne, J.A.C., Bossuyt, P.M.M., 2011. QUADAS-2: A
2246 revised tool for the quality assessment of diagnostic accuracy studies. *Annals of*
2247 *Internal Medicine* 155, 529–536.

2248 Wichtmann, B.D., Katemann, C., Kadrija, M., Layer, Y.C., Bischoff, L.M., Scheuver,
2249 Y., Mezger, M., Weber, O.M., Luetkens, J.A., Attenberger, U.I., Radbruch, A.,
2250 Paech, D., 2025. Rapid and motion-robust pediatric brain imaging: T2-weighted
2251 turbo-spin-echo propeller acquisition with compressed sensing. *Pediatric radiology*
2252 doi:10.1007/s00247-024-06088-z. pMID: 39589491.

2253 Wismüller, A., Vosoughi, M.A., 2022. Self-supervised learning in medical imaging:
2254 anomaly detection in MRI using autoencoders, in: *Real-Time Image Processing*
2255 *and Deep Learning 2022*, p. 9. doi:10.1117/12.2622233.

2256 Wu, F., Luo, H., Wang, X., Yang, Q., Zhuang, Y., Lin, L., Dong, Y., Tulupov, A.,
2257 Zhang, Y., Cai, S., Chen, Z., Cai, C., Bao, J., Cheng, J., 2025a. Application of
2258 anti-motion ultra-fast quantitative MRI in neurological disorder imaging: Insights
2259 from huntington's disease. *Journal of magnetic resonance imaging : JMRI* doi:10.
2260 1002/jmri.29682.

2261 Wu, J., Chen, L., Li, Z., Li, X., Sun, T., Wang, L., Wang, R., Wei, H., Zhang,
2262 Y., 2025b. 3d isotropic high-resolution fetal brain MRI reconstruction from motion
2263 corrupted thick data based on physical-informed unsupervised learning. *IEEE Jour-*
2264 *nal of Biomedical and Health Informatics* 29, 8256–8267. doi:10.1109/jbhi.2025.
2265 3586049.

2266 Xia, L., Qi, X., Liu, J., Yuan, Y., Wang, Y., Xu, N., Zhang, J., 2026. Accurate multi-b-
2267 value DWI generation using two-stage deep learning: multicenter study. *European*
2268 *Journal of Radiology* 194, 112497. doi:10.1016/j.ejrad.2025.112497.

2269 Xie, K., Gao, L., Zhang, H., Zhang, S., Xi, Q., Zhang, F., Sun, J., Lin, T., Sui, J., Ni,
2270 X., 2024. GAN-based metal artifacts region inpainting in brain MRI imaging with
2271 reflective registration. *Medical physics* doi:10.1002/mp.16724.

2272 Xiong, Z., Gao, Y., Liu, F., Emery, D., Butcher, K., Wilman, A.H., Sun, H., 2026.
2273 Advancing quantitative susceptibility mapping with 2.5d diffusion models for rapid
2274 intracranial hemorrhage quantification. *Magnetic resonance in medicine* doi:10.
2275 1002/mrm.70358. pMID: 41873533.

2276 Xu, J., Zu, T., Hsu, Y.C., Wang, X., Chan, K.W.Y., Zhang, Y., 2025. DEISM: Deep
2277 reconstruction framework with self-calibration mechanisms for accelerated chemical
2278 exchange saturation transfer imaging. *IEEE Transactions on Biomedical Engineer-*
2279 *ing* 72, 2413–2424. doi:10.1109/tbme.2025.3543403.

2280 Yang, G., Yu, S., Dong, H., Slabaugh, G., Dragotti, P.L., Ye, X., Liu, F., Arridge, S.,
2281 Keegan, J., Guo, Y., Firmin, D., 2018. DAGAN: Deep de-aliasing generative adver-
2282 sarial networks for fast compressed sensing MRI reconstruction. *IEEE Transactions*
2283 *on Medical Imaging* 37, 1310–1321.

2284 Yang, Q., Lin, Y., Wang, J., Bao, J., Wang, X., Ma, L., Zhou, Z., Yang, Q., Cai,
2285 S., He, H., Cai, C., Dong, J., Cheng, J., Chen, Z., Zhong, J., 2022. Model-based
2286 synthetic data-driven learning (most-dl): Application in single-shot t2 mapping

2287 with severe head motion using overlapping-echo acquisition. IEEE transactions on
2288 medical imaging doi:10.1109/tmi.2022.3179981. pMID: 35657830.

2289 Yarach, U., Chatnuntawech, I., Liao, C., Teerapittayanon, S., Iyer, S.S., Kim, T.H.,
2290 Haldar, J., Cho, J., Bilgic, B., Hu, Y., Hargreaves, B., Setsompop, K., 2025.
2291 Blip-up blip-down circular epi (buda-cepi) for distortion-free dmri with rapid un-
2292 rolled deep learning reconstruction. Magnetic resonance imaging doi:10.1016/j.mri.
2293 2024.110277. pMID: 39566835.

2294 Yarach, U., Chatnuntawech, I., Setsompop, K., Suwannasak, A., Angkurawaranon,
2295 S., Madla, C., Hanprasertpong, C., Sangpin, P., 2024. Improved reconstruction
2296 for highly accelerated propeller diffusion 1.5t clinical MRI. Magnetic Resonance
2297 Materials in Physics, Biology and Medicine 37, 283–294. doi:10.1007/s10334-023-
2298 01142-7.

2299 Yasaka, K., Akai, H., Kato, S., Tajima, T., Yoshioka, N., Furuta, T., Kageyama, H.,
2300 Toda, Y., Akahane, M., Ohtomo, K., Abe, O., Kiryu, S., 2024. Iterative motion
2301 correction technique with deep learning reconstruction for brain MRI: A volunteer
2302 and patient study. Journal of imaging informatics in medicine doi:10.1007/s10278-
2303 024-01184-w. pMID: 38942939.

2304 Yasaka, K., Akai, H., Sugawara, H., Tajima, T., Akahane, M., Yoshioka, N., Kabasawa,
2305 H., Miyo, R., Ohtomo, K., Abe, O., Kiryu, S., 2022. Impact of deep learning recon-
2306 struction on intracranial 1.5 t magnetic resonance angiography. Japanese journal of
2307 radiology doi:10.1007/s11604-021-01225-2. pMID: 34851499.

2308 Yilmaz, K., Çukur, T., Patel, V.M., 2023. Self-supervised MRI reconstruction with
2309 unrolled diffusion models. Medical Image Computing and Computer Assisted Inter-
2310 vention – MICCAI 2023 doi:10.1007/978-3-031-43999-5_47.

2311 Yoo, H., Moon, H.E., Kim, S., Kim, D.H., Choi, Y.H., Cheon, J.E., Lee, J.S., Lee,
2312 S., 2025. Evaluation of image quality and scan time efficiency in accelerated 3D
2313 t1-weighted pediatric brain MRI using deep learning-based reconstruction. Korean
2314 journal of radiology doi:10.3348/kjr.2024.0701. pMID: 39898398.

2315 Yoshida, N., Kageyama, H., Akai, H., Sasaki, K., Sakurai, N., Koori, N., Yamamoto,
2316 S., Kodama, N., 2026. Deep learning approach to super-resolution correction of
2317 brain MRI motion artifacts for accurate hippocampal volumetry. Scientific reports
2318 doi:10.1038/s41598-026-44834-5. pMID: 41866419.

2319 Yoshida, N., Kageyama, H., Akai, H., Yasaka, K., Sugawara, H., Okada, Y., Kuni-
2320 matsu, A., 2022. Motion correction in MR image for analysis of vsrad using gen-
2321 erative adversarial network. PloS one doi:10.1371/journal.pone.0274576. pMID:
2322 36103561.

2323 You, S., Gondova, A., Izaguirre, C.S.A., Milo, G.T., Jeong, S., Lee, H.J., Tarui, T.,
2324 Rollins, C.K., Yun, H.J., Grant, P.E., Im, K., 2025. Conditional deep generative
2325 normative modeling for structural and developmental anomaly detection in the fetal
2326 brain. NeuroImage doi:10.1016/j.neuroimage.2025.121442. pMID: 40912664.

2327 Yu, S., Dai, G., Wang, Z., Li, L., Wei, X., Xie, Y., 2018. A consistency evaluation of
2328 signal-to-noise ratio in the quality assessment of human brain magnetic resonance
2329 images. BMC medical imaging doi:10.1186/s12880-018-0256-6. pMID: 29769079.

2330 Yuan, S., Liu, Y., Wei, R., Zhu, J., Men, K., Dai, J., 2024. A novel loss function to
2331 reproduce texture features for deep learning-based MRI-to-CT synthesis. *Medical*
2332 *physics* doi:10.1002/mp.16850.

2333 Zahneisen, B., Keating, B., Singh, A., Herbst, M., Ernst, T., 2016. Reverse retrospec-
2334 tive motion correction. *Magnetic resonance in medicine* doi:10.1002/mrm.25830.
2335 PMID: 26140504.

2336 Zaitsev, M., Maclaren, J., Herbst, M., 2015. Motion artifacts in MRI: A review. *NMR*
2337 *in Biomedicine* 28, 911–935.

2338 Zhang, C., Karkalousos, D., Bazin, P.L., Coolen, B.F., Vrenken, H., Sonke, J.J.,
2339 Forstmann, B.U., Poot, D.H.J., Caan, M.W.A., 2022. A unified model for recon-
2340 struction and $r2^*$ mapping of accelerated 7t data using the quantitative recurrent
2341 inference machine. *NeuroImage* doi:10.1016/j.neuroimage.2022.119680. PMID:
2342 36240989.

2343 Zhang, C., Klein, S., Cristobal-Huerta, A., Hernandez-Tamames, J.A., Poot, D.H.J.,
2344 2021. Apir4emc: Autocalibrated parallel imaging reconstruction for extended multi-
2345 contrast imaging. *Magnetic resonance imaging* doi:10.1016/j.mri.2021.02.002.
2346 PMID: 33592248.

2347 Zhang, J., Wu, J., Chen, S., Zhang, Z., Cai, S., Cai, C., Chen, Z., 2019a. Robust
2348 single-shot T2 mapping via multiple overlapping-echo acquisition and deep neural
2349 network. *IEEE Transactions on Medical Imaging* 38, 1801–1811. doi:10.1109/tmi.
2350 2019.2896085.

2351 Zhang, L., Wu, M., Wang, L., Steffens, D.C., Potter, G.G., Liu, M., 2026. Iterative
2352 learning for joint image denoising and motion artifact correction of 3D brain MRI.
2353 *Medical image analysis* doi:10.1016/j.media.2025.103872. PMID: 41274083.

2354 Zhang, Q., Ruan, G., Yang, W., Liu, Y., Zhao, K., Feng, Q., Chen, W., Wu, E.X.,
2355 Feng, Y., 2019b. MRI gibbs?ringing artifact reduction by means of machine learning
2356 using convolutional neural networks. *Magnetic Resonance in Medicine* 82, 2133–
2357 2145. doi:10.1002/mrm.27894.

2358 Zhao, C., Shao, M., Carass, A., Li, H., Dewey, B.E., Ellingsen, L.M., Woo, J.,
2359 Guttman, M.A., Blitz, A.M., Stone, M., Calabresi, P.A., Halperin, H., Prince,
2360 J.L., 2019. Applications of a deep learning method for anti-aliasing and super-
2361 resolution in MRI. *Magnetic resonance imaging* doi:10.1016/j.mri.2019.05.038.
2362 PMID: 31247254.

2363 Zhao, R., Yaman, B., Zhang, Y., Stewart, R., Dixon, A., Knoll, F., Huang, Z., Lui,
2364 Y.W., Hansen, M.S., Lungren, M.P., 2022. fastMRI+, clinical pathology annota-
2365 tions for knee and brain fully sampled magnetic resonance imaging data. *Scientific*
2366 *Data* 9, 152. doi:10.1038/s41597-022-01255-z.

2367 Zhou, L., Ni, X., Kong, Y., Zeng, H., Xu, M., Zhou, J., Wang, Q., Liu, C., 2023. Miti-
2368 gating misalignment in MRI-to-CT synthesis for improved synthetic CT generation:
2369 an iterative refinement and knowledge distillation approach. *Physics in medicine*
2370 *and biology* doi:10.1088/1361-6560/ad0ddc.

2371 Zhou, X., Balachandra, A.R., Romano, M.F., Chin, S.P., Au, R., Kolachalama, V.B.,
2372 2024. Adversarial learning for MRI reconstruction and classification of cognitively
2373 impaired individuals. *IEEE Access* doi:10.1109/access.2024.3408840.

2374 Zhu, D., Xu, F., Liu, D., Qin, Q., 2025a. A straightforward approach for 3D single-
2375 shot arterial spin labeling-based brain perfusion imaging: Preventing artifacts due to
2376 signal fluctuations. *Magnetic resonance in medicine* doi:10.1002/mrm.30439. PMID:
2377 39887515.

2378 Zhu, P., Liu, C., Fu, Y., Chen, N., Qiu, A., 2025b. Cycle-conditional diffusion model
2379 for noise correction of diffusion-weighted images using unpaired data. *Medical image*
2380 *analysis* doi:10.1016/j.media.2025.103579. PMID: 40273728.

2381 Zhu, S., Jin, Y., Spears, T., Zawar, I., Fletcher, P.T., 2026. Realdeal: Enhancing
2382 realism and details in brain image generation via image-to-image diffusion models.
2383 *Deep Generative Models* doi:10.1007/978-3-032-05472-2_15.

2384 Zibetti, M.V.W., Knoll, F., Regatte, R.R., 2022. Alternating learning approach for
2385 variational networks and undersampling pattern in parallel MRI applications. *IEEE*
2386 *Transactions on Computational Imaging* doi:10.1109/tci.2022.3176129.